

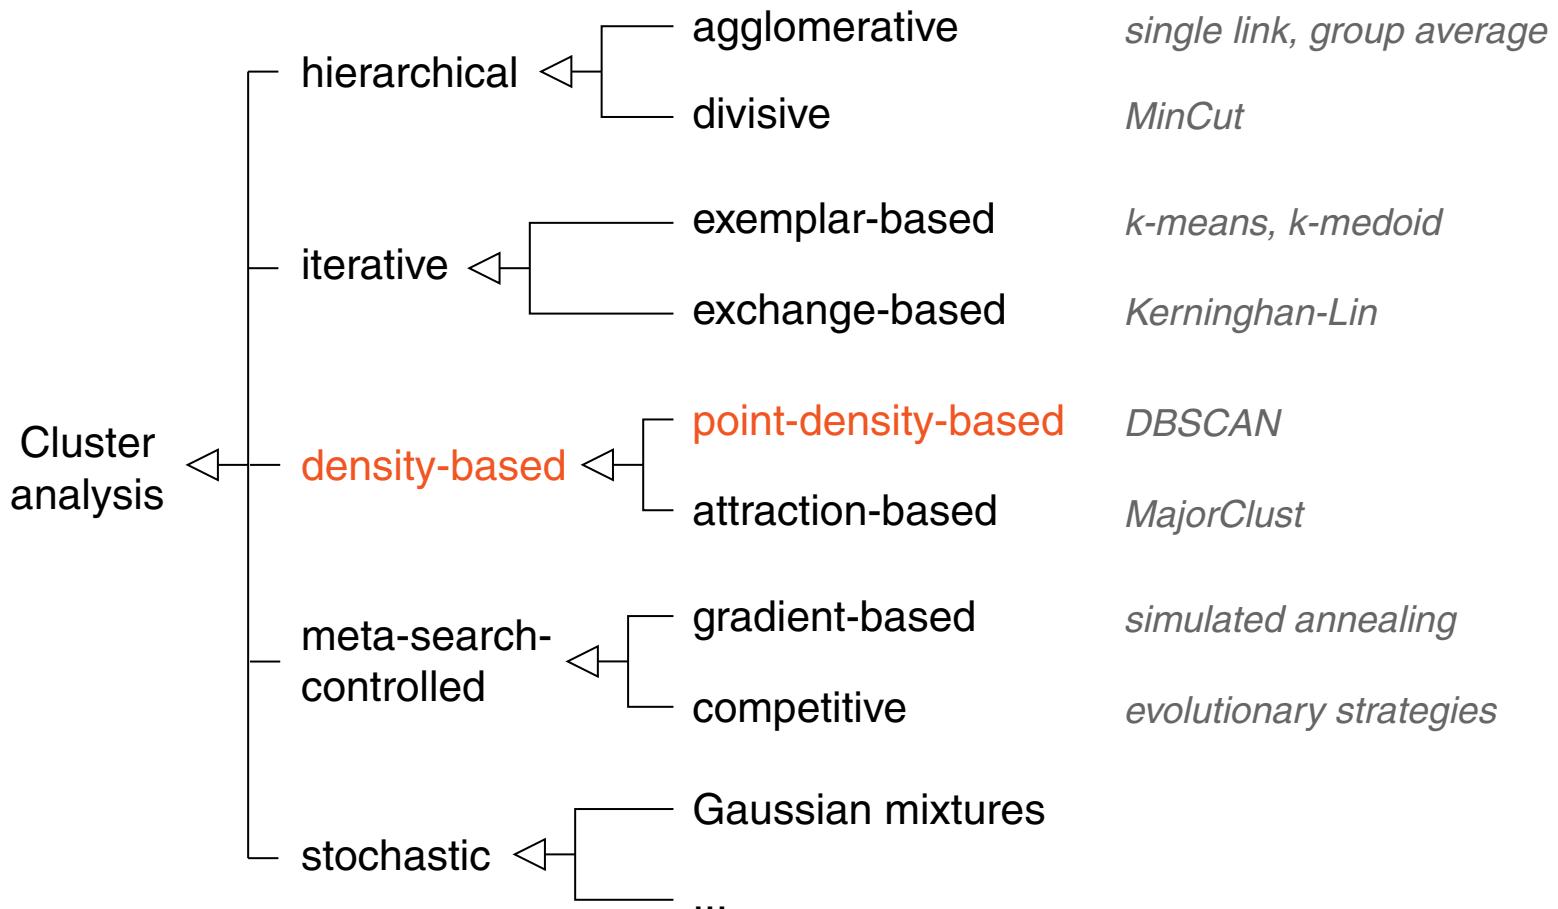
Chapter DM:II (continued)

II. Cluster Analysis

- Cluster Analysis Basics
- Hierarchical Cluster Analysis
- Iterative Cluster Analysis
- Density-Based Cluster Analysis
- Cluster Evaluation
- Constrained Cluster Analysis

Density-Based Cluster Analysis

Merging Principles



Density-Based Cluster Analysis

Density-based algorithms strive to partition the graph $G = \langle V, E, w \rangle$, better: the set of points V , into regions of similar density.

Approaches to density estimation:

- parameter-based: the type of the underlying data distribution is known
- parameterless: construction of histograms, superposition of kernel density estimators

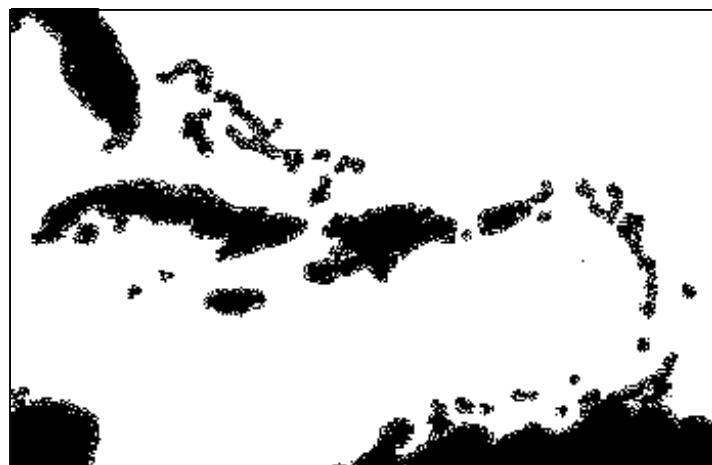
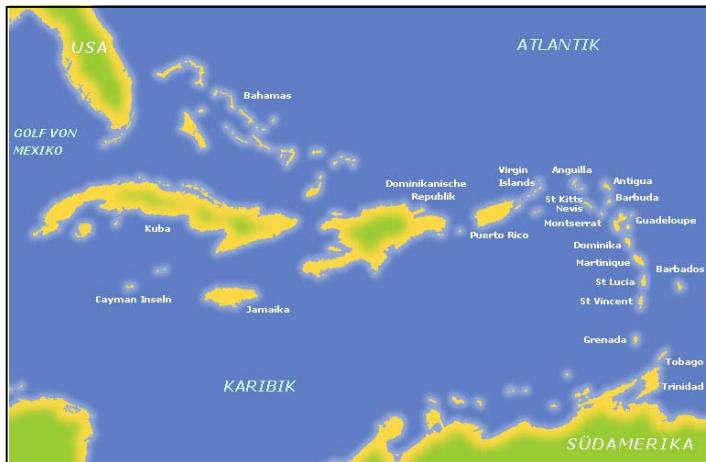
Density-Based Cluster Analysis

Density-based algorithms strive to partition the graph $G = \langle V, E, w \rangle$, better: the set of points V , into regions of similar density.

Approaches to density estimation:

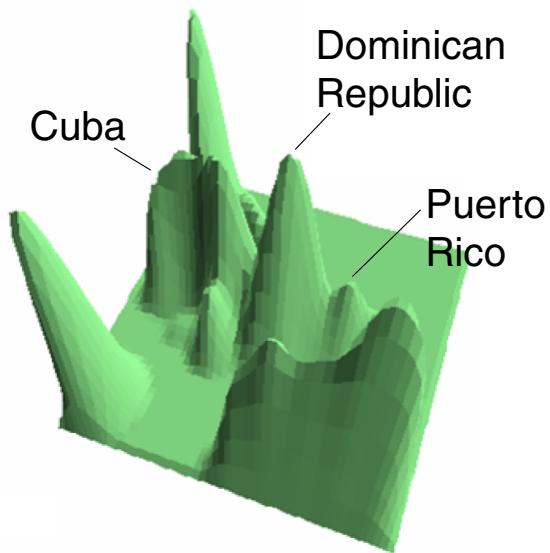
- parameter-based: the type of the underlying data distribution is known
- parameterless: construction of histograms, superposition of kernel density estimators

Example (Caribbean Islands) :




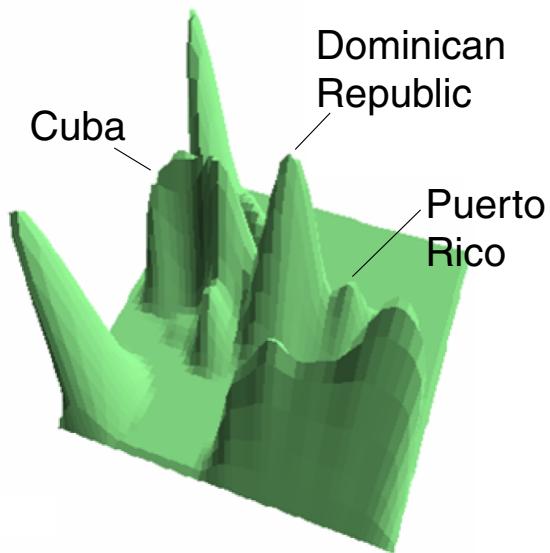
Density-Based Cluster Analysis

Density Estimation with Gaussian Kernel for the Example



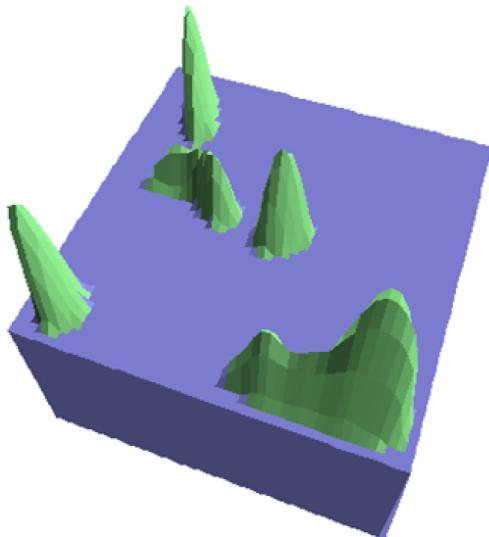
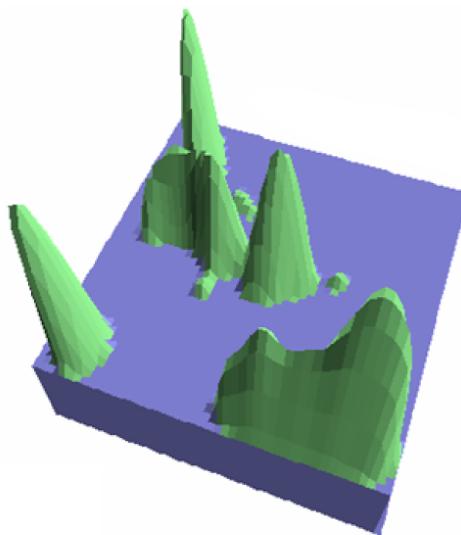
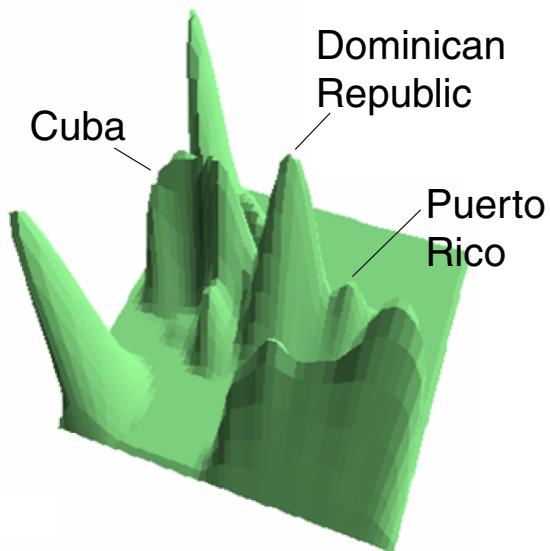
Density-Based Cluster Analysis

Density Estimation with Gaussian Kernel for the Example



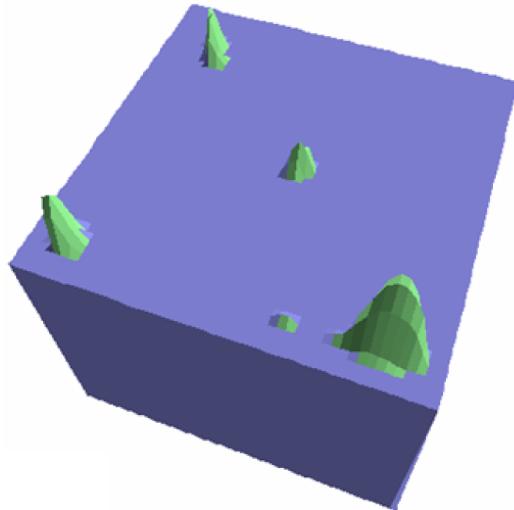
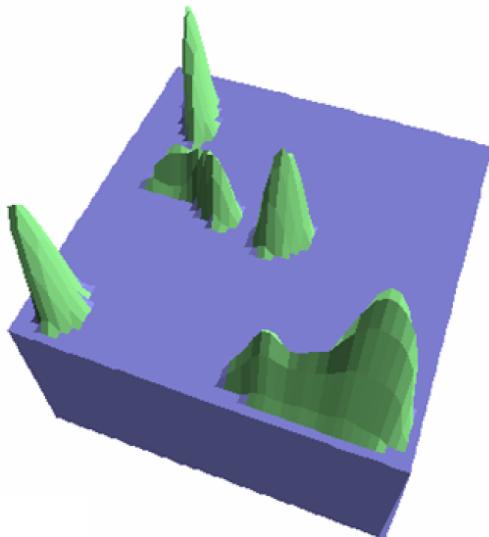
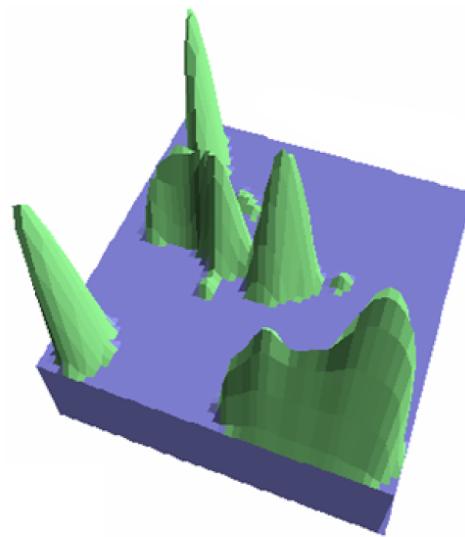
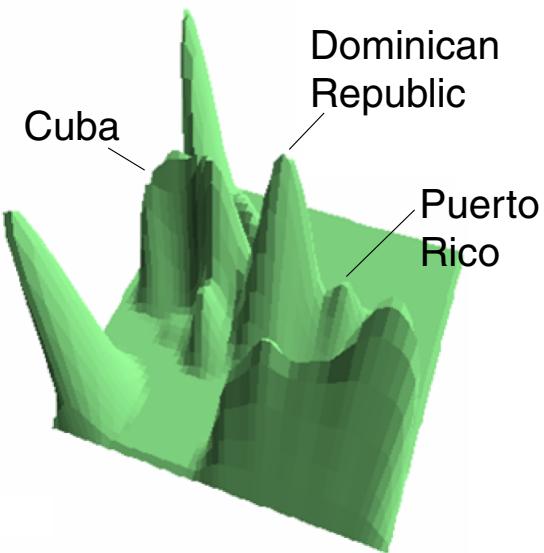
Density-Based Cluster Analysis

Density Estimation with Gaussian Kernel for the Example



Density-Based Cluster Analysis

Density Estimation with Gaussian Kernel for the Example



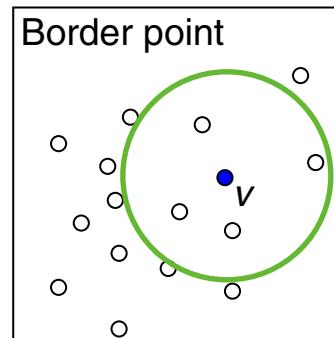
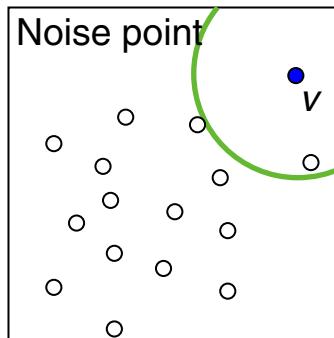
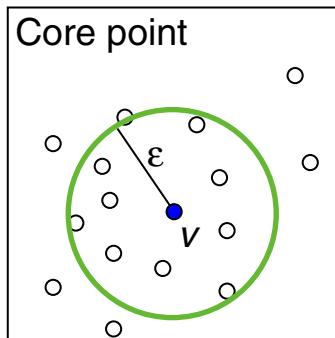
Remarks:

- The green three-dimensional landscape is the result of associating each point of the rasterized map (right-hand side) with a three-dimensional Gaussian kernel and superimposing them.
- Raising the “water level” in the three-dimensional landscape (\sim clipping at a certain contour line) corresponds to splitting the dendrogram and reveals possible clusters. Observe that no single water level (contour line) can be chosen such that all clusters can be identified.

Density-Based Cluster Analysis

DBSCAN: Density Estimation Principle [Ester et al. 1996]

Let $N_\varepsilon(v)$ denote the ε -neighborhood of some point $v \in V$. Distinguish between three kinds of points:



1. v is a core point $\Leftrightarrow |N_\varepsilon(v)| \geq \text{MinPts}$
2. v is a noise point \Leftrightarrow
 v is not density-reachable from any core point
3. v is a border point otherwise

Density-Based Cluster Analysis


DBSCAN: Density Estimation Principle

A point u is **density-reachable** from a point v , if either of the following conditions hold:

(a) $u \in N_\varepsilon(v)$, where v is a core point.

(b) There exists a set of points $\{v_1, \dots, v_l\}$, where

$v_{i+1} \in N_\varepsilon(v_i)$ and v_i is core point, $i = 1, \dots, l - 1$, with $v_1 = v$, $v_l = u$.



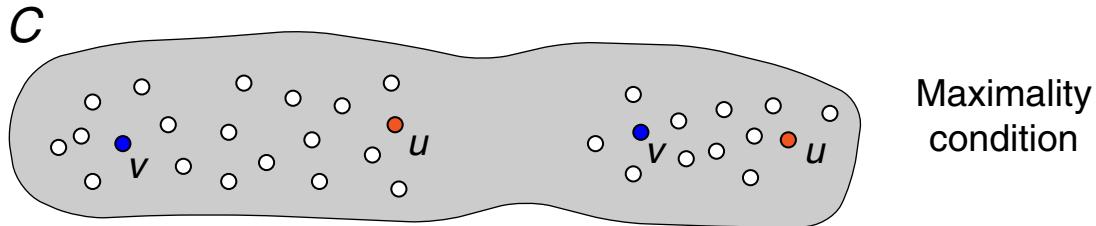
Condition (b) can be considered as the transitive application of Condition (a).

Density-Based Cluster Analysis

DBSCAN: Cluster Interpretation

A cluster $C \subseteq V$ fulfills the following two conditions:

1. $\forall u, v : \text{If } v \in C \text{ and } u \text{ is density-reachable from } v, \text{ then } u \in C.$

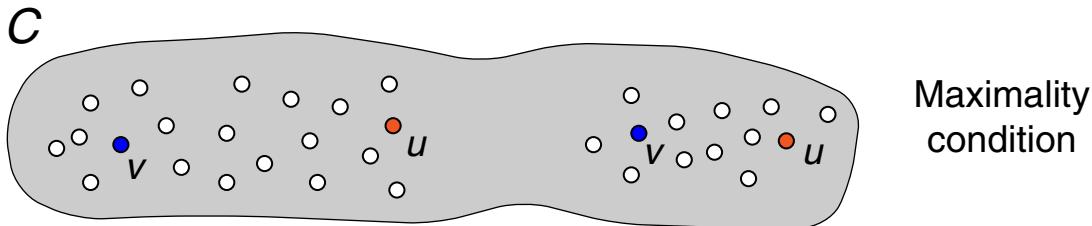


Density-Based Cluster Analysis

DBSCAN: Cluster Interpretation

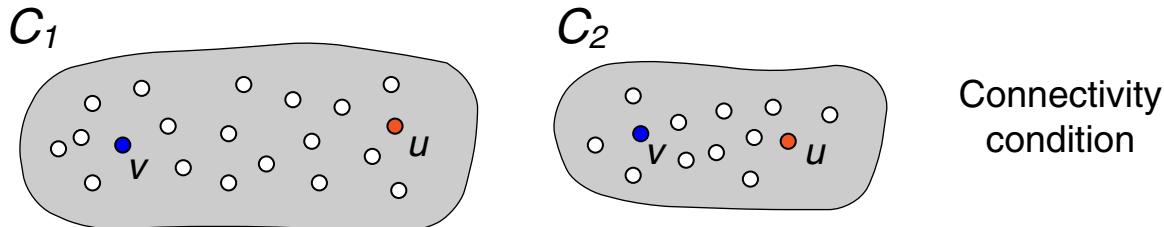
A cluster $C \subseteq V$ fulfills the following two conditions:

1. $\forall u, v : \text{If } v \in C \text{ and } u \text{ is density-reachable from } v, \text{ then } u \in C.$



2. $\forall u, v \in C : u \text{ is density-connected with } v,$ which is defined as follows:

There exists a point t wherefrom u and v are density-reachable.



Remarks:

- ❑ Condition 1 (maximality) states a constraint between any two points.
Condition 2 (connectivity) states an additional constraint with respect to a third point.
- ❑ The maximality condition is problematic if a border point lies in the ε -neighborhoods of two core points that belong to two different clusters. Such a border point would then belong to both clusters; however, the algorithm breaks this tie by assigning this point to the first cluster found.

Density-Based Cluster Analysis

DBSCAN: Algorithm

Input: $G = \langle V, E, w \rangle$. Weighted graph.
 d . Distance measure for two nodes in V .
 ε . Neighborhood radius.
 $MinPts$. Lower bound for point number in ε -neighborhood.

Output: $\gamma : V \rightarrow \mathbb{Z}$. Cluster assignment function.

- 1.
- 2.
3. $v = choose_unclassified_point(V)$
4. $N_\varepsilon(v) = neighborhood(G, d, v, \varepsilon)$
5. **IF** $|N_\varepsilon(v)| \geq MinPts$ **THEN** // v is core point
- 6.
7. $C_i = density_reachable_hull(G, d, N_\varepsilon(v))$ // identify dense region
8. **FOREACH** $v \in C_i$ **DO** $\gamma(v) = i$ // assign points in region to cluster i
9. **ELSE** $\gamma(v) = -1$ // classify v _tentatively_ as noise (-1)
- 10.
- 11.

Density-Based Cluster Analysis

DBSCAN: Algorithm

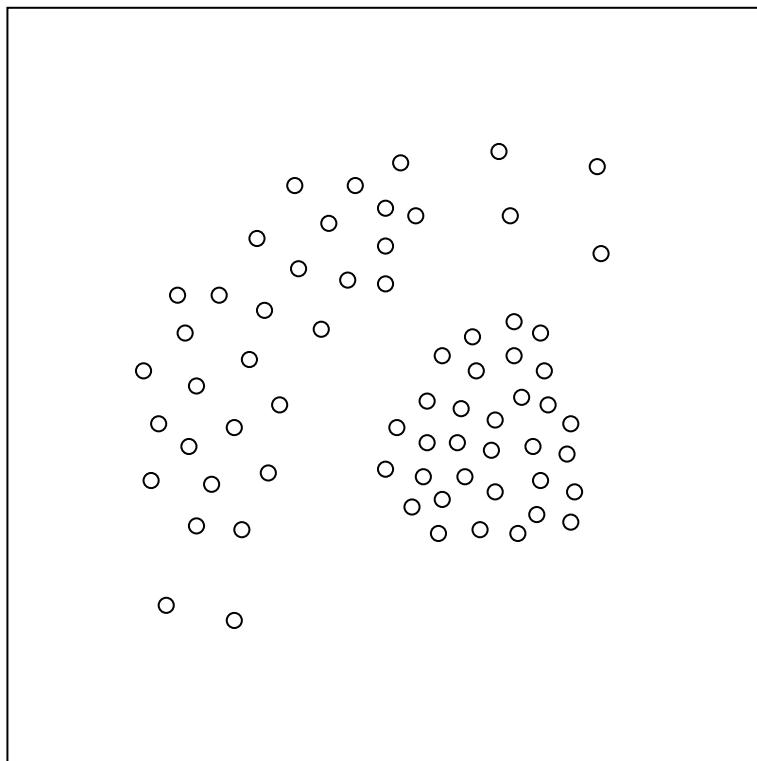
Input: $G = \langle V, E, w \rangle$. Weighted graph.
 d . Distance measure for two nodes in V .
 ε . Neighborhood radius.
 $MinPts$. Lower bound for point number in ε -neighborhood.

Output: $\gamma : V \rightarrow \mathbb{Z}$. Cluster assignment function.

1. $i = 0$
2. **WHILE** $\exists v : (v \in V \text{ AND } \gamma(v) = \perp)$ **DO** // check for unclassified (\perp) points
3. $v = choose_unclassified_point(V)$
4. $N_\varepsilon(v) = neighborhood(G, d, v, \varepsilon)$
5. **IF** $|N_\varepsilon(v)| \geq MinPts$ **THEN** // v is core point
6. $i = i + 1$
7. $C_i = density_reachable_hull(G, d, N_\varepsilon(v))$ // identify dense region
8. **FOREACH** $v \in C_i$ **DO** $\gamma(v) = i$ // assign points in region to cluster i
9. **ELSE** $\gamma(v) = -1$ // classify v _tentatively_ as noise (-1)
10. **ENDDO**
11. **RETURN**(γ)

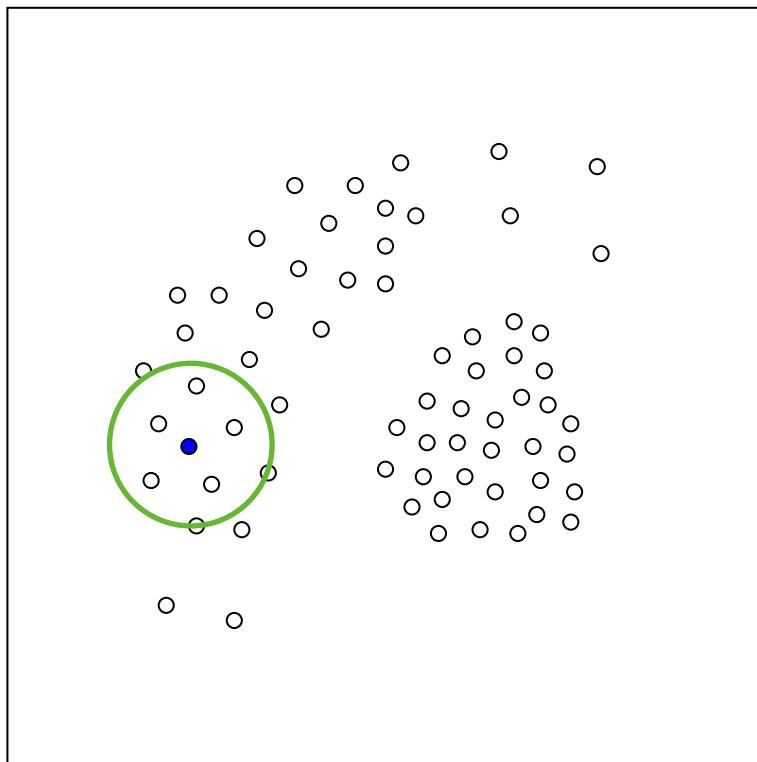
Density-Based Cluster Analysis

DBSCAN



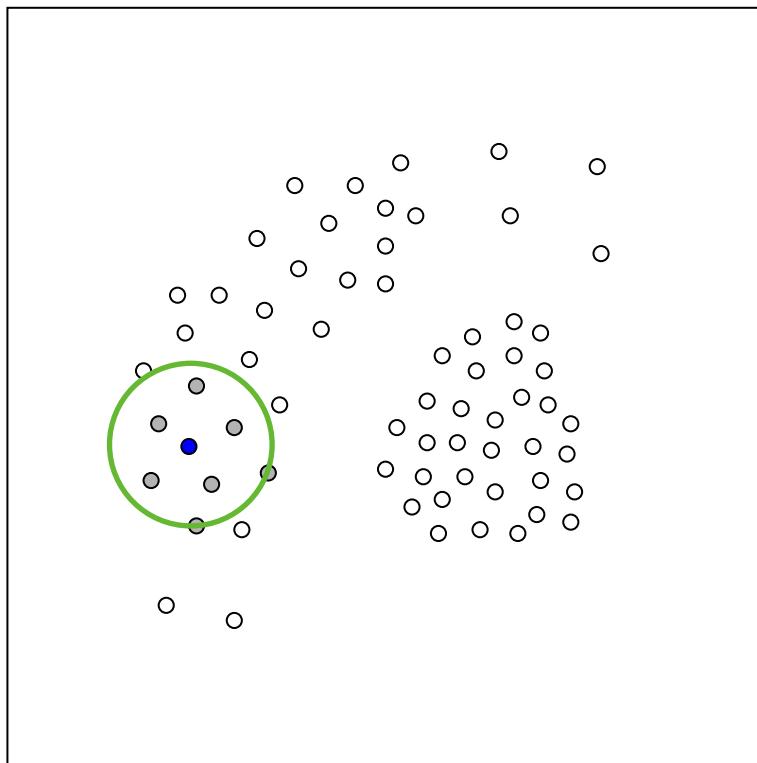
Density-Based Cluster Analysis

DBSCAN



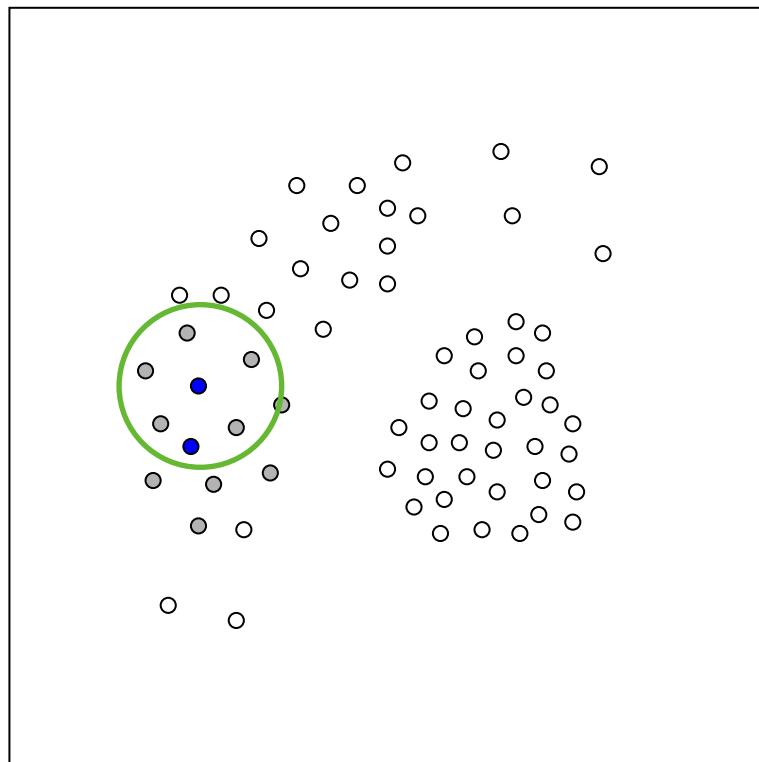
Density-Based Cluster Analysis

DBSCAN



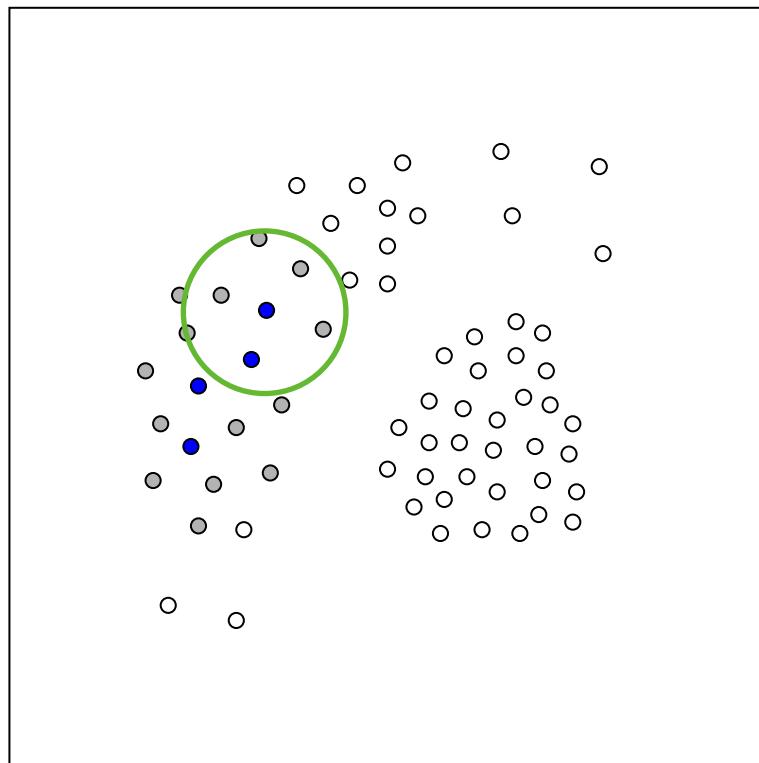
Density-Based Cluster Analysis

DBSCAN



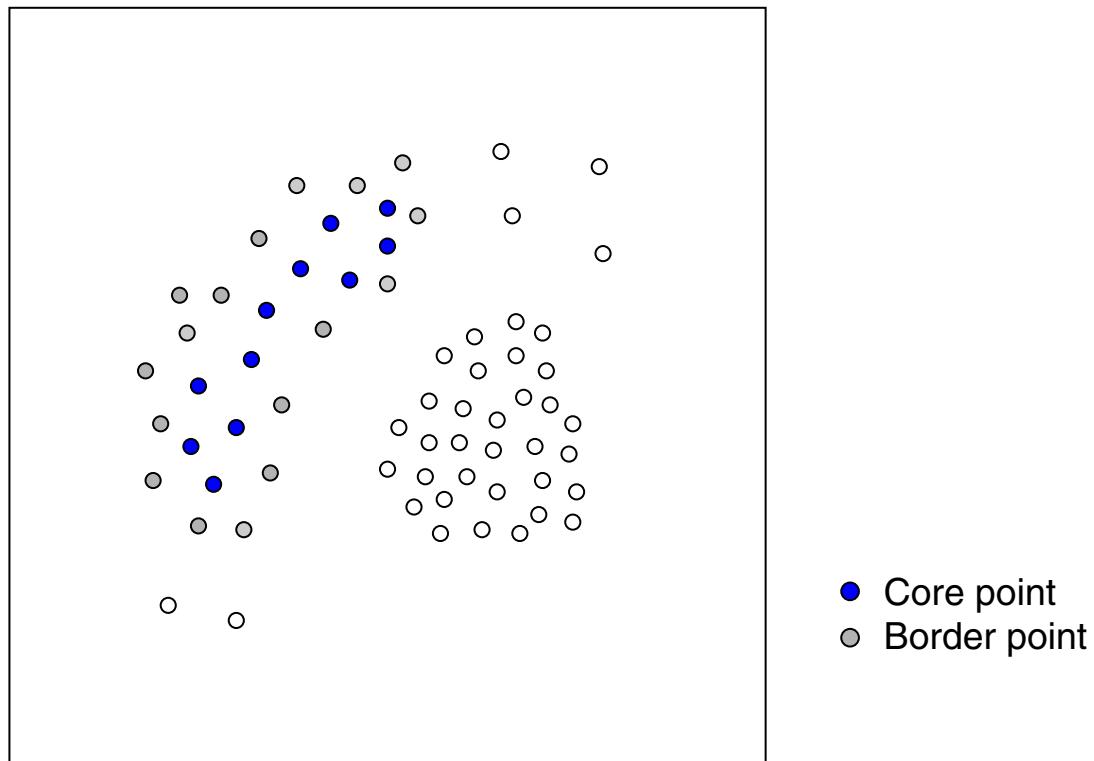
Density-Based Cluster Analysis

DBSCAN



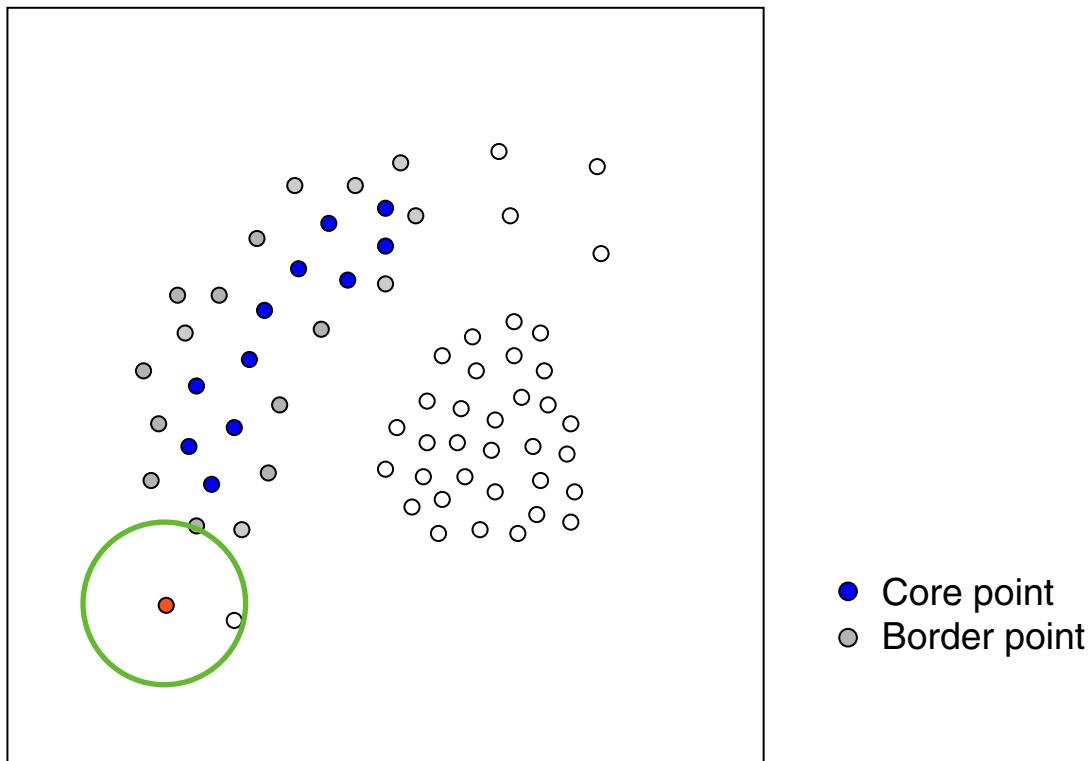
Density-Based Cluster Analysis

DBSCAN



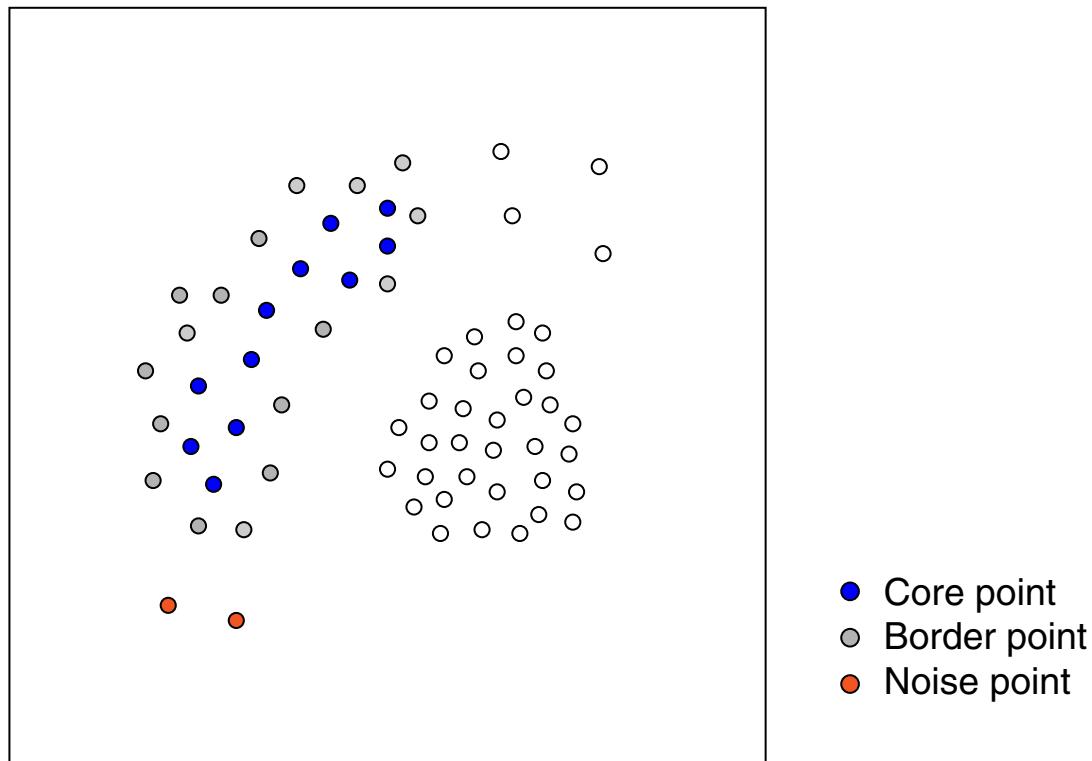
Density-Based Cluster Analysis

DBSCAN



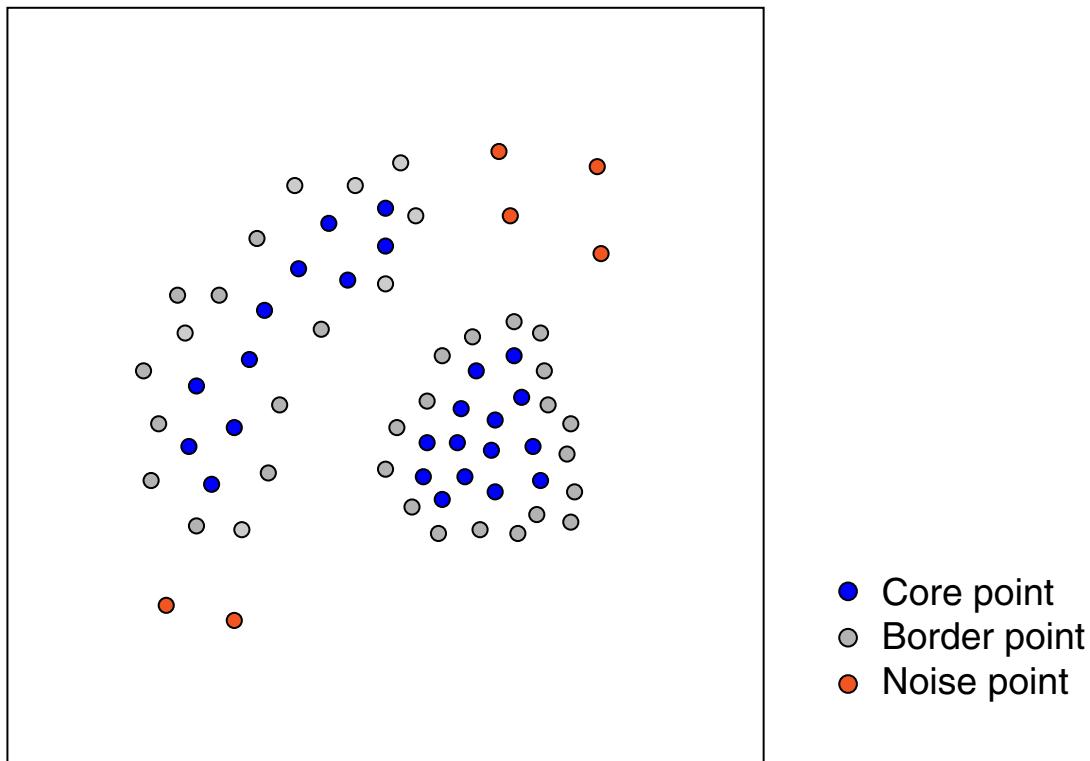
Density-Based Cluster Analysis

DBSCAN



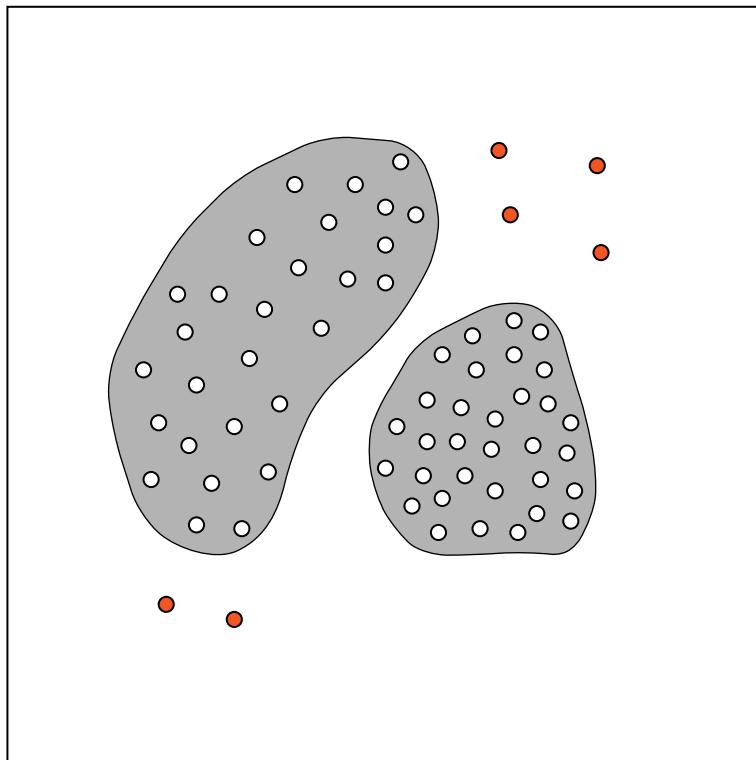
Density-Based Cluster Analysis

DBSCAN



Density-Based Cluster Analysis

DBSCAN



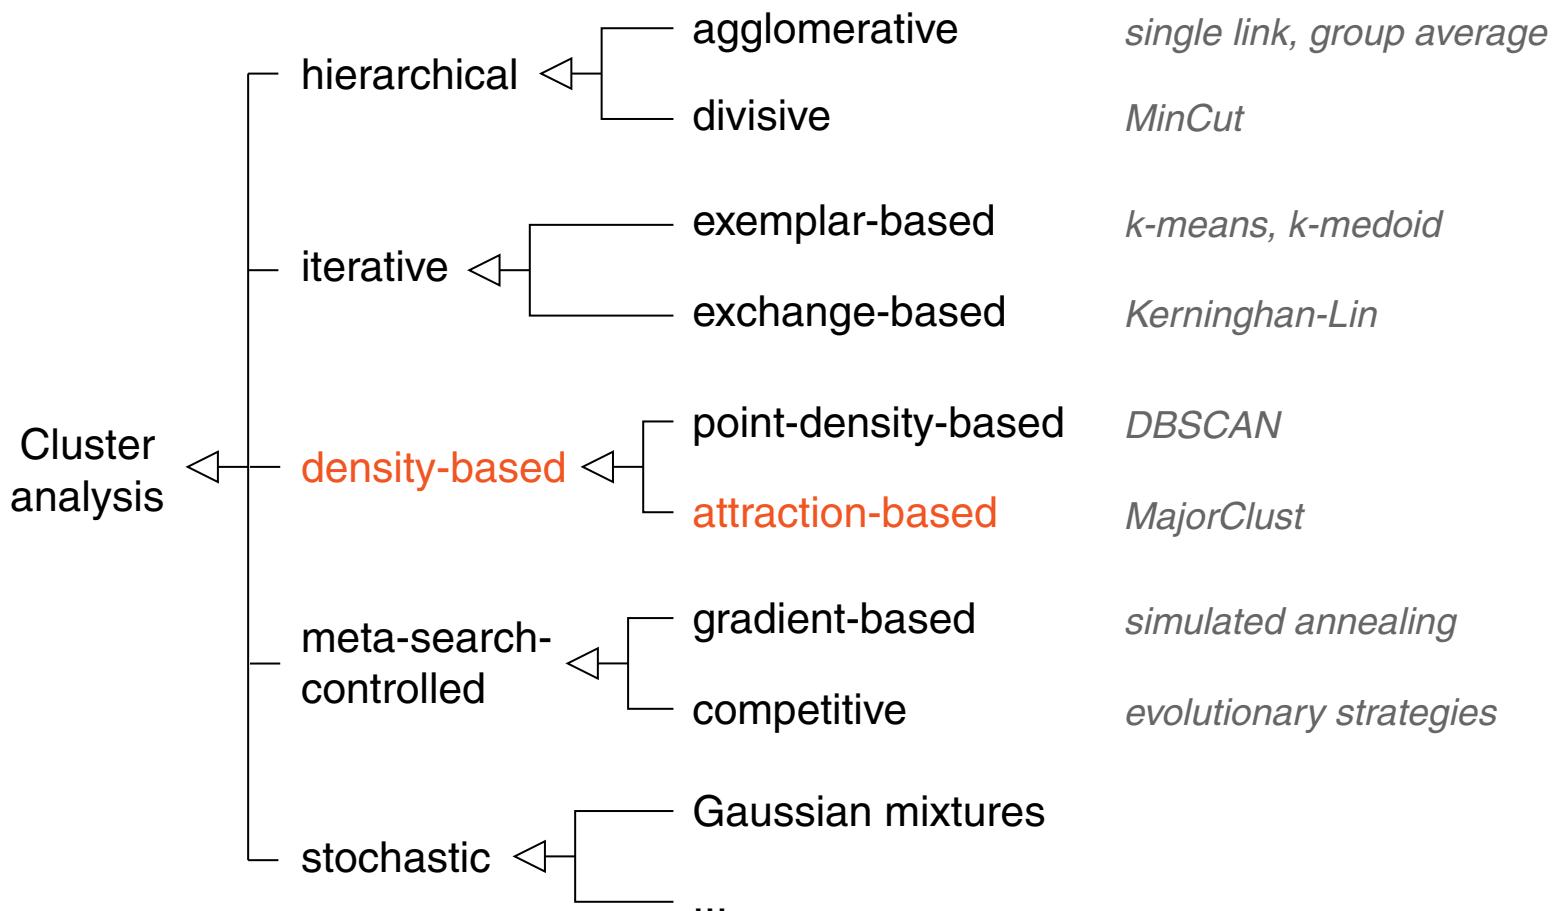
● Noise point

Remarks:

- ❑ Note that points that are labeled as noise can be re-labeled with a cluster number exactly once. I.e., a point will retain its tentative noise label only if it is not density-reachable from any other point.
- ❑ The construction of C_i as the density-reachable hull of $N_\varepsilon(v)$ (Line 7) corresponds to a recursive analysis of the points in $N_\varepsilon(v)$ with regard to their density reachability.
- ❑ A slightly different and compact formulation of the algorithm is given in [Tan/Steinbach/Kumar 2005, p. 528].

Density-Based Cluster Analysis

Merging Principles



Density-Based Cluster Analysis

MajorClust: Density Estimation Principle [Stein/Niggemann 1999]

The weighted edges in a graph $G = \langle V, E, w \rangle$ are interpreted as attracting forces, where members of the same cluster combine their forces.

Unique membership situation, leading to a merge of two clusters:



Density-Based Cluster Analysis

MajorClust: Density Estimation Principle [Stein/Niggemann 1999]

The weighted edges in a graph $G = \langle V, E, w \rangle$ are interpreted as attracting forces, where members of the same cluster combine their forces.

Unique membership situation, leading to a merge of two clusters:



Unique membership situation, leading to a change of cluster membership:



Density-Based Cluster Analysis

MajorClust: Density Estimation Principle [Stein/Niggemann 1999]

The weighted edges in a graph $G = \langle V, E, w \rangle$ are interpreted as attracting forces, where members of the same cluster combine their forces.


Unique membership situation, leading to a merge of two clusters:



Unique membership situation, leading to a change of cluster membership:



Ambiguous membership situation:



Density-Based Cluster Analysis

MajorClust: Algorithm

Input: $G = \langle V, E, w \rangle$. Weighted graph.
 d . Distance measure for two nodes in V .

Output: $\gamma : V \rightarrow \mathbf{N}$. Cluster assignment function.

- 1 .
- 2 .
- 3 .
- 4 .
- 5 . **FOREACH** $v \in V$ **DO**
- 6 . $\gamma^* = \operatorname{argmax}_{i: i \in \{1, \dots, |V|\}} \sum_{\{u, v\}: \{u, v\} \in E \wedge \gamma(u) = i} w(u, v)$ // find strongest cluster for v
- 7 . **IF** $\gamma(v) \neq \gamma^*$ **THEN** $\gamma(v) = \gamma^*$, $t = \text{False}$ **ENDIF** // reassign v
- 8 . **ENDDO**
- 9 .
- 10 .

Density-Based Cluster Analysis

MajorClust: Algorithm

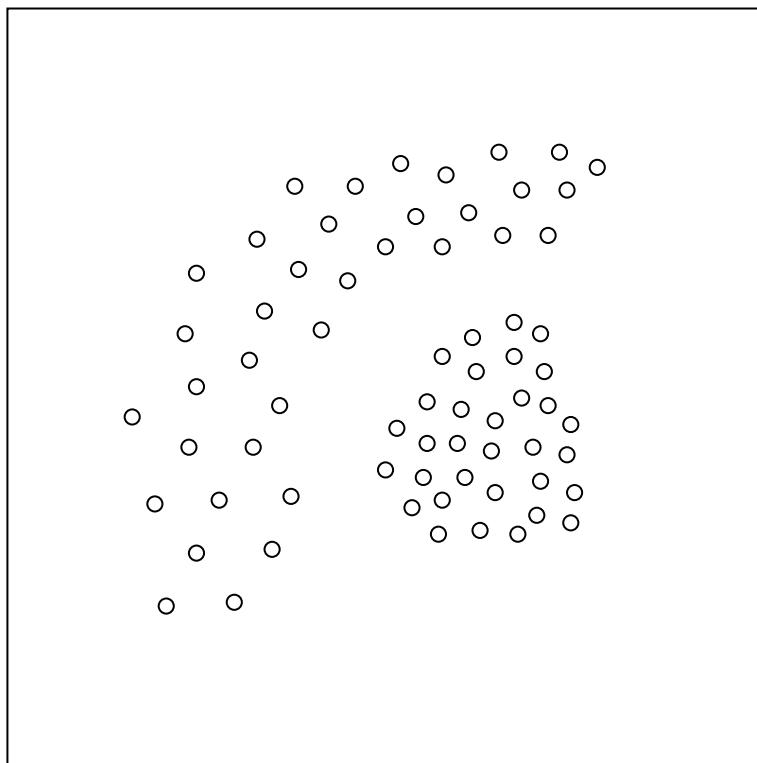
Input: $G = \langle V, E, w \rangle$. Weighted graph.
 d . Distance measure for two nodes in V .

Output: $\gamma : V \rightarrow \mathbb{N}$. Cluster assignment function.

1. $i = 0, t = \text{False}$
2. **FOREACH** $v \in V$ **DO** $i = i + 1, \gamma(v) = i$ **ENDDO** // initial clustering
3. **UNLESS** t **DO**
4. $t = \text{True}$
5. **FOREACH** $v \in V$ **DO**
6. $\gamma^* = \underset{i: i \in \{1, \dots, |V|\}}{\operatorname{argmax}} \sum_{\{u, v\}: \{u, v\} \in E \wedge \gamma(u)=i} w(u, v)$ // find strongest cluster for v
7. **IF** $\gamma(v) \neq \gamma^*$ **THEN** $\gamma(v) = \gamma^*, t = \text{False}$ **ENDIF** // reassign v
8. **ENDDO**
9. **ENDDO**
10. **RETURN**(γ)

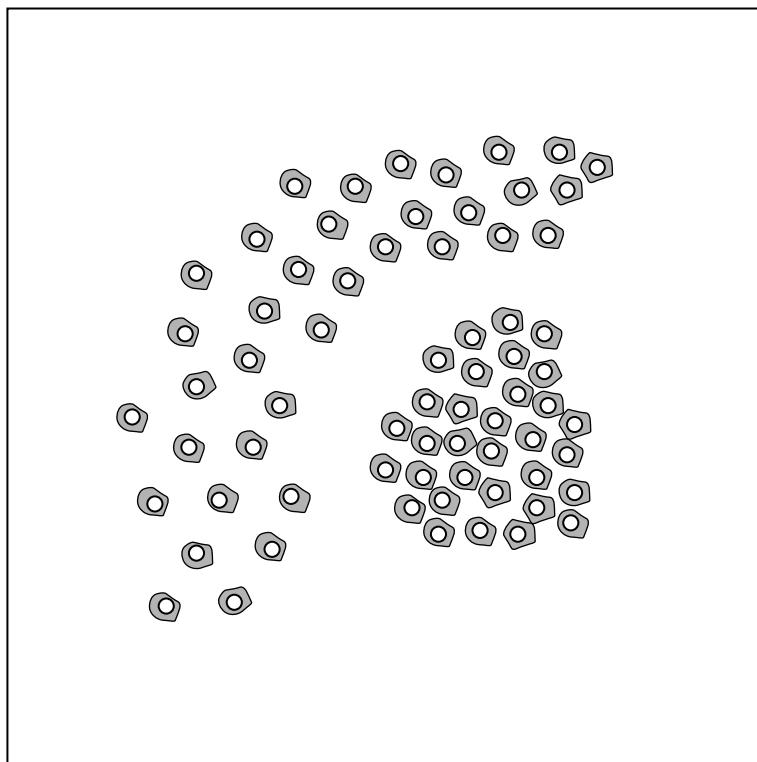
Density-Based Cluster Analysis

MajorClust




Density-Based Cluster Analysis

MajorClust



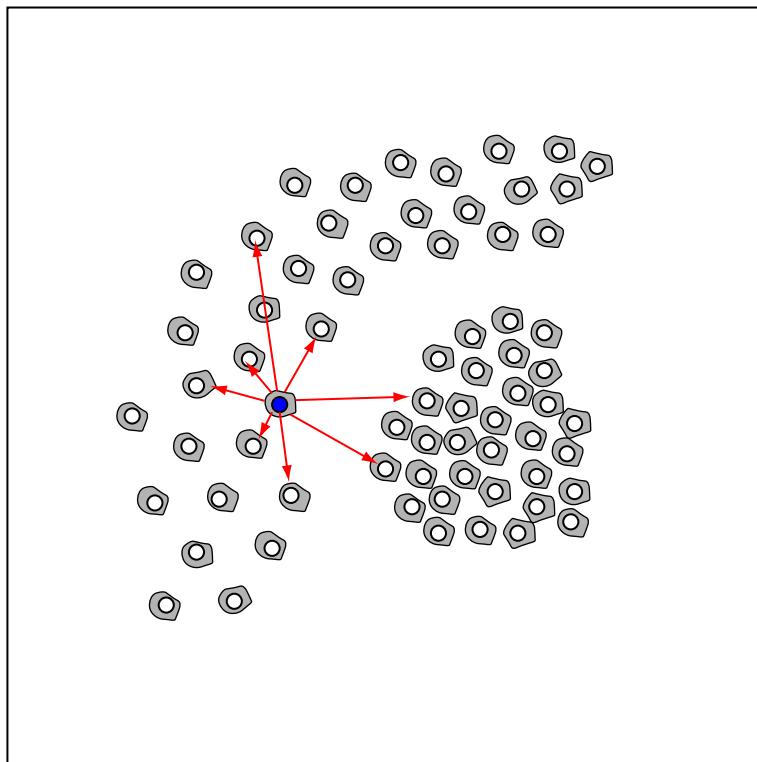
Density-Based Cluster Analysis

MajorClust



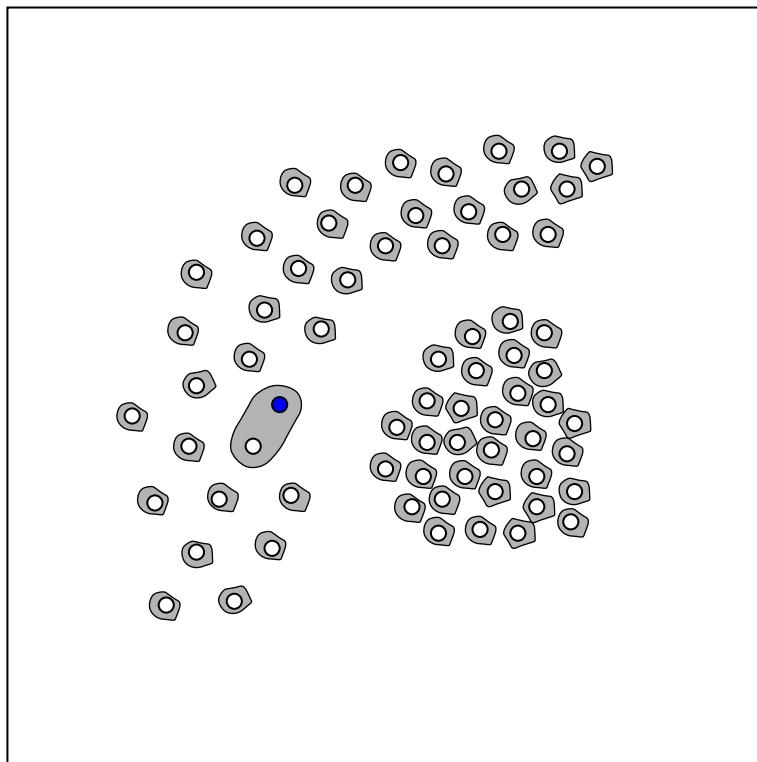
Density-Based Cluster Analysis

MajorClust



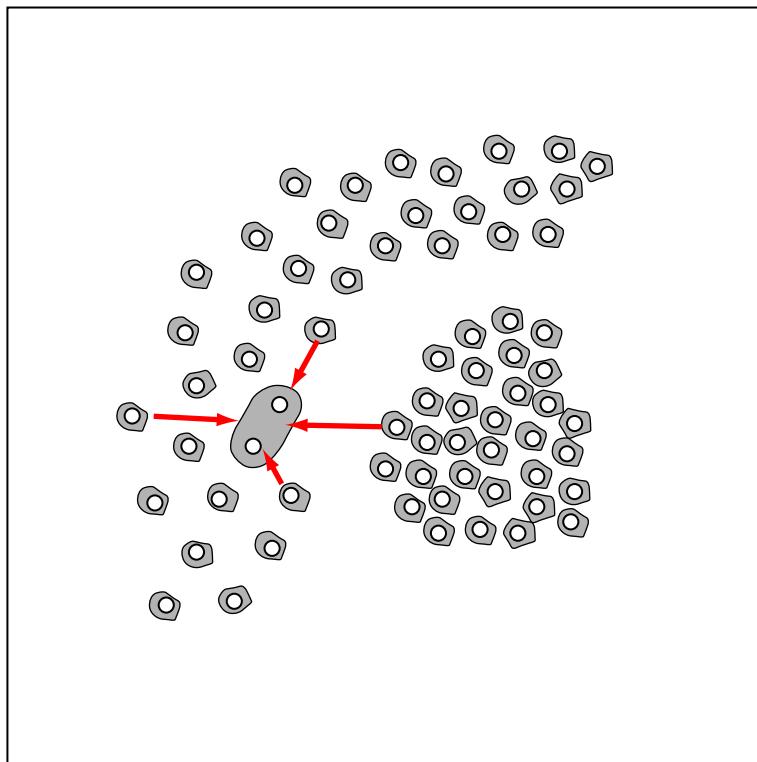
Density-Based Cluster Analysis

MajorClust




Density-Based Cluster Analysis

MajorClust



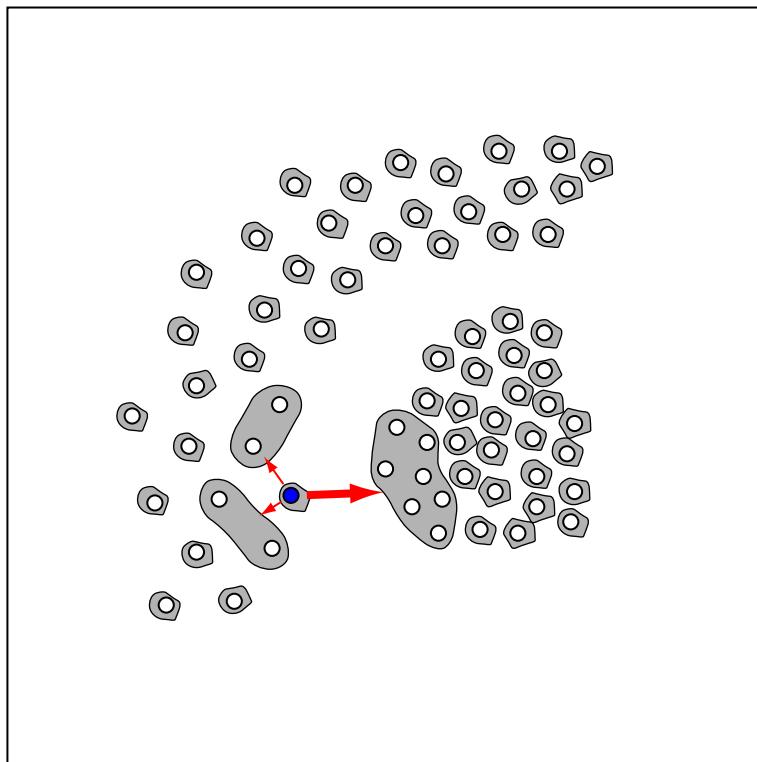
Density-Based Cluster Analysis

MajorClust



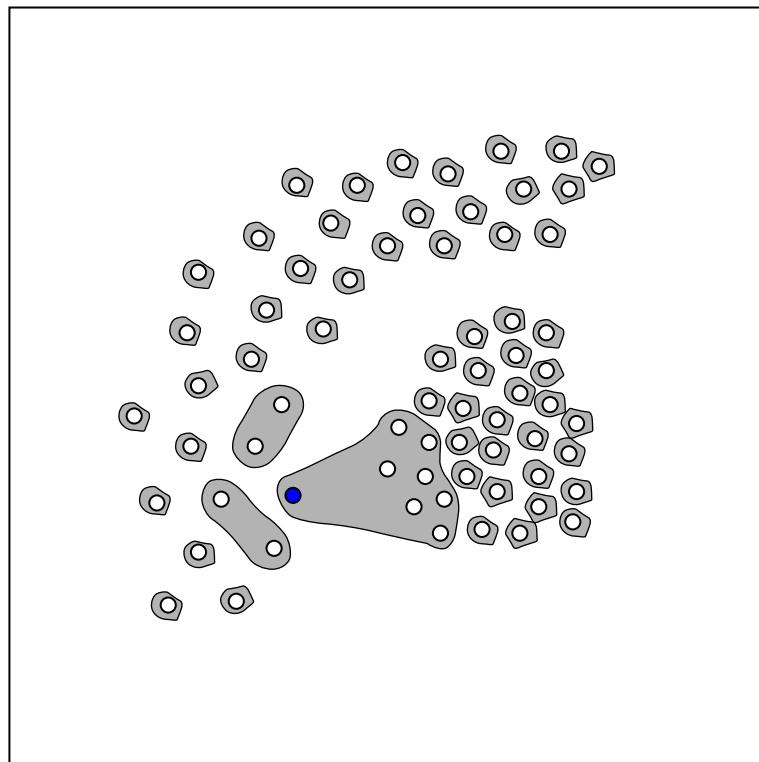
Density-Based Cluster Analysis

MajorClust



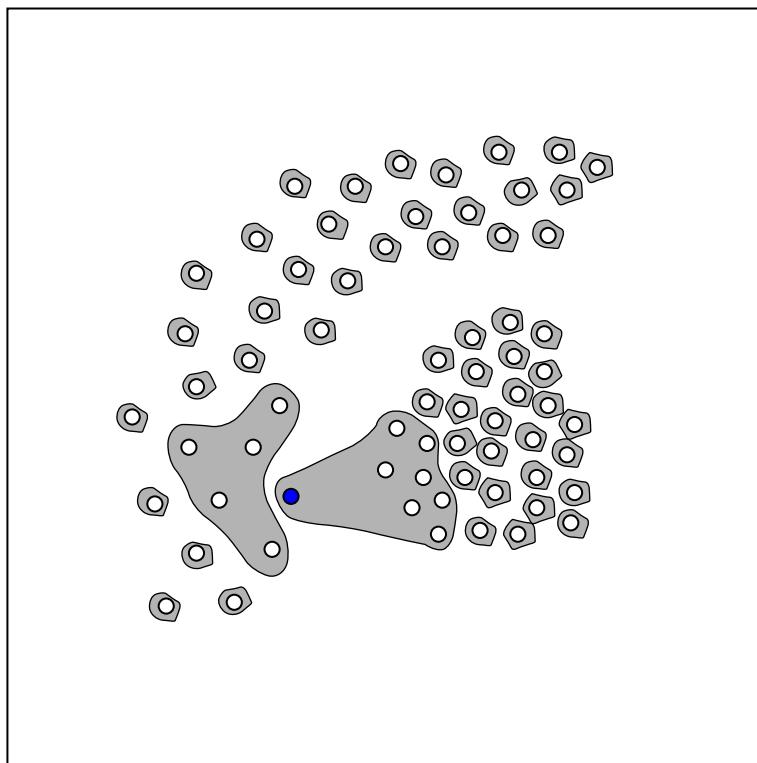
Density-Based Cluster Analysis

MajorClust




Density-Based Cluster Analysis

MajorClust



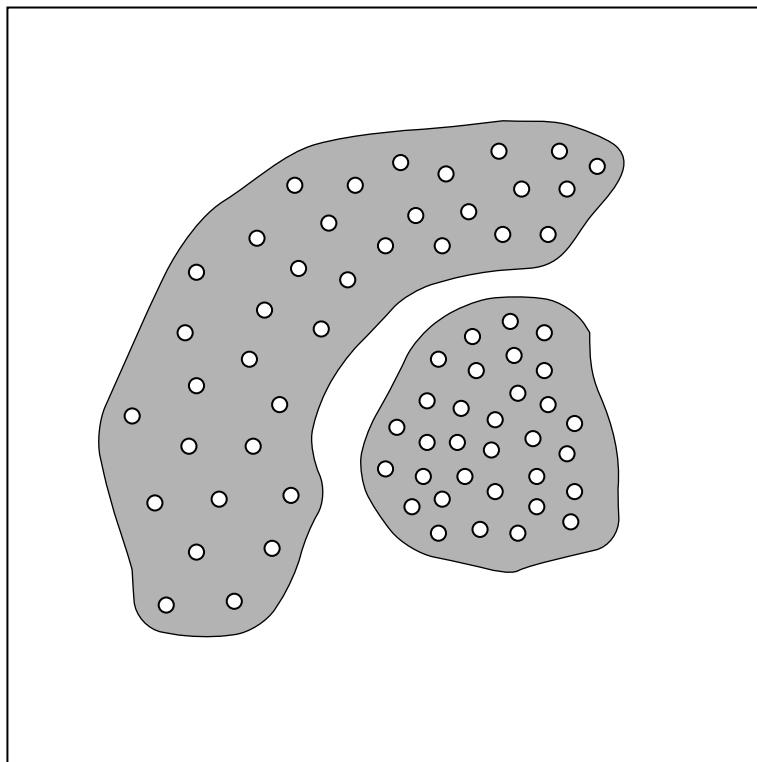
Density-Based Cluster Analysis

MajorClust



Density-Based Cluster Analysis

MajorClust



Remarks:

- MajorClust combines properties from other paradigms:
 - distance-depending analysis (hierarchical paradigm, iterative paradigm)
 - reversible merging decisions (iterative paradigm)
 - distribution-dependent analysis (density paradigm)

Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)

Each clustering $\mathcal{C} = \{C_1, \dots, C_k\}$ induces k subgraphs within $G = \langle V, E, w \rangle$.

MajorClust is a heuristic to maximize the *weighted partial edge connectivity*, $\Lambda(\mathcal{C})$.

$$\Lambda(\mathcal{C}) = \sum_{i=1}^k |C_i| \cdot \lambda_i$$



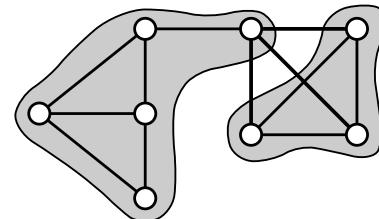
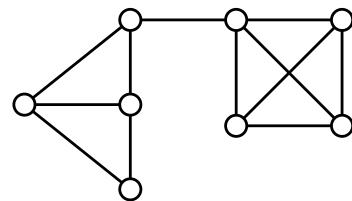
Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)

Each clustering $\mathcal{C} = \{C_1, \dots, C_k\}$ induces k subgraphs within $G = \langle V, E, w \rangle$.

MajorClust is a heuristic to maximize the *weighted partial edge connectivity*, $\Lambda(\mathcal{C})$.

$$\Lambda(\mathcal{C}) = \sum_{i=1}^k |C_i| \cdot \lambda_i$$




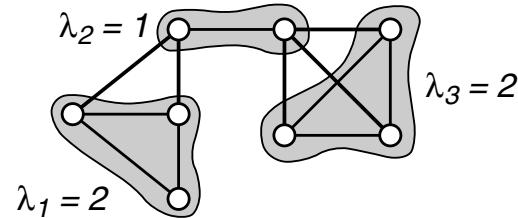

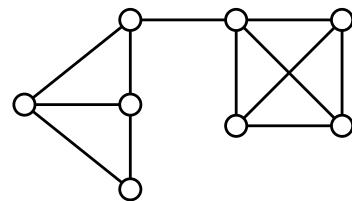
Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)

Each clustering $\mathcal{C} = \{C_1, \dots, C_k\}$ induces k subgraphs within $G = \langle V, E, w \rangle$.

MajorClust is a heuristic to maximize the *weighted partial edge connectivity*, $\Lambda(\mathcal{C})$.

$$\Lambda(\mathcal{C}) = \sum_{i=1}^k |C_i| \cdot \lambda_i$$




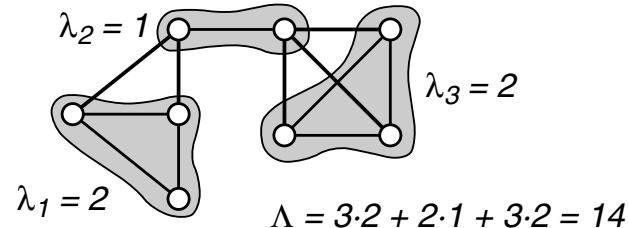

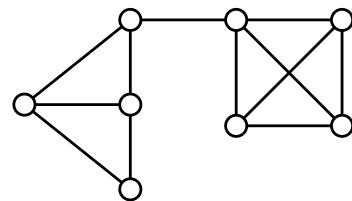
Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)

Each clustering $\mathcal{C} = \{C_1, \dots, C_k\}$ induces k subgraphs within $G = \langle V, E, w \rangle$.

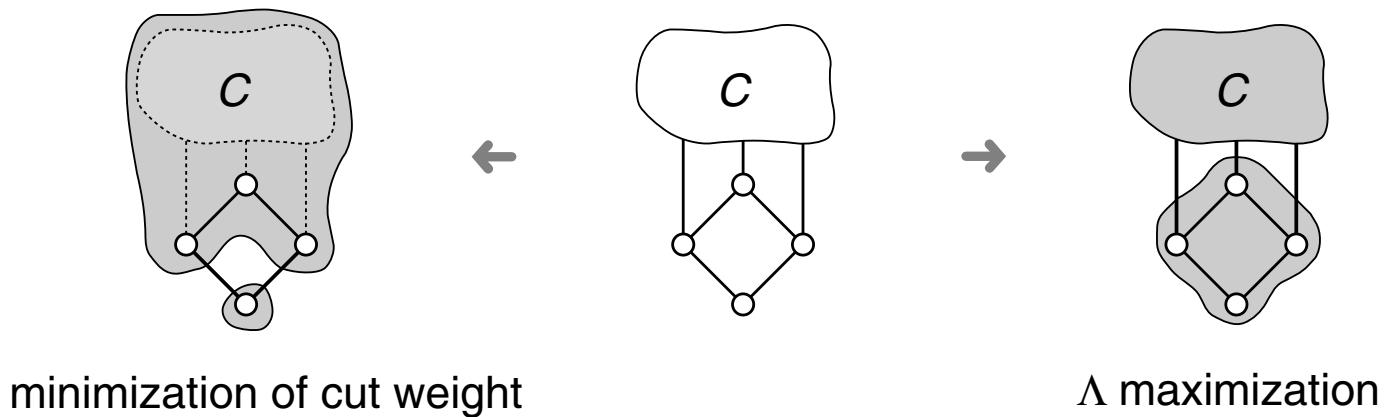
MajorClust is a heuristic to maximize the *weighted partial edge connectivity*, $\Lambda(\mathcal{C})$.

$$\Lambda(\mathcal{C}) = \sum_{i=1}^k |C_i| \cdot \lambda_i$$



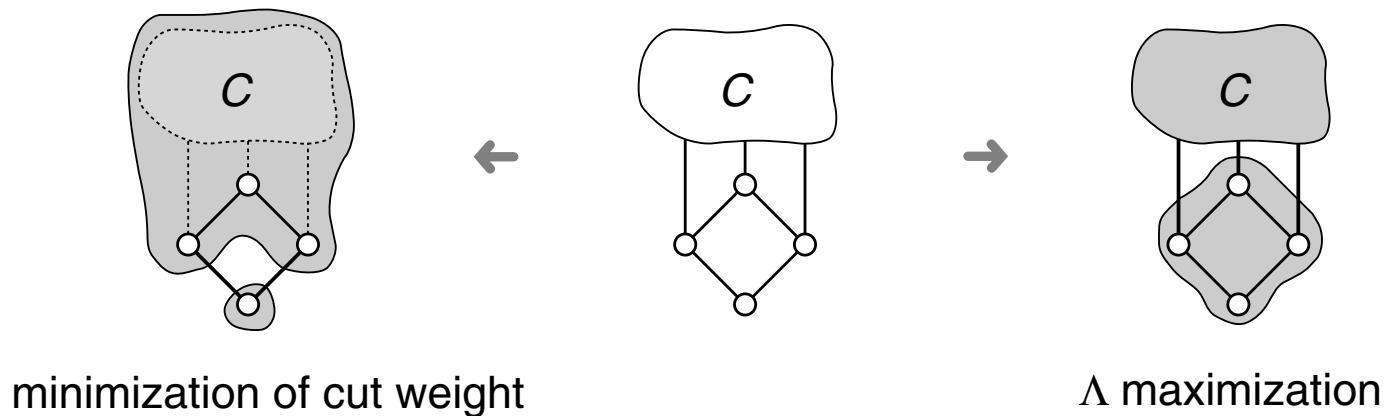
Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)



Density-Based Cluster Analysis

MajorClust: Density Estimation Principle (continued)



Theorem 5 (Strong Splitting Condition) [Stein/Niggemann 1999]

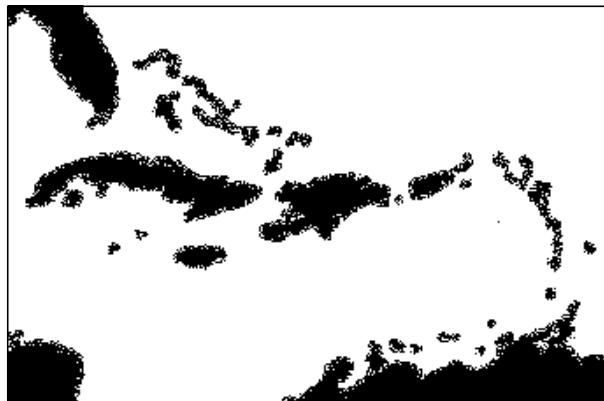
Let $\mathcal{C} = \{C_1, \dots, C_k\}$ be a partitioning of a graph $G = \langle V, E, w \rangle$. Moreover, let $\lambda(G)$ denote the edge connectivity of G , and let $\lambda_1, \dots, \lambda_k$ denote the edge connectivity values of the k subgraphs that are induced by C_1, \dots, C_k .

If the inequality $\lambda(G) < \min\{\lambda_1, \dots, \lambda_k\}$ holds, then the partitioning defined by Λ -maximization corresponds to the minimum cut splitting of G . The inequality is denoted as “Strong Splitting Condition”.

Density-Based Cluster Analysis

DBSCAN versus MajorClust: Low-Dimensional Data

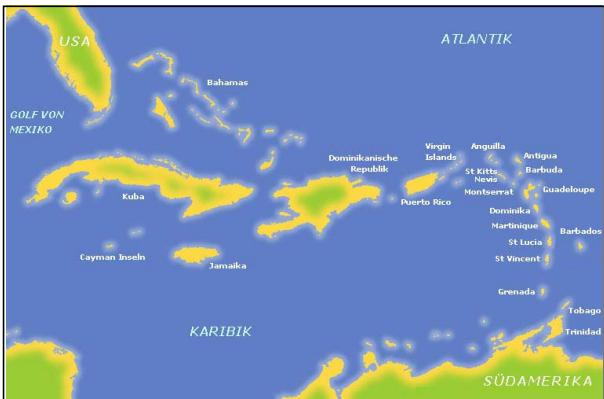
Caribbean Islands, about 20.000 points:



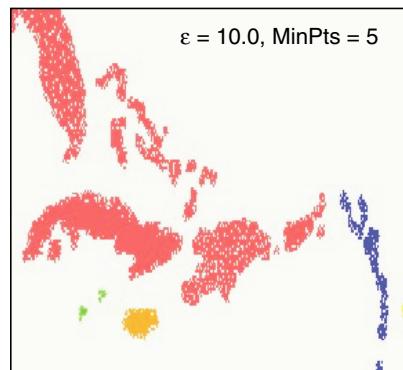
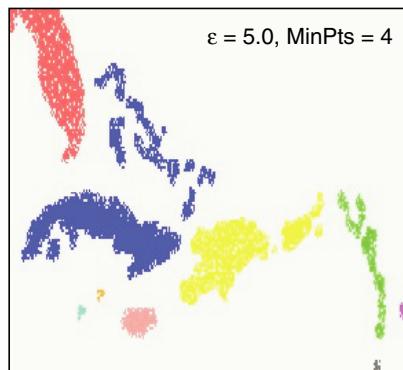
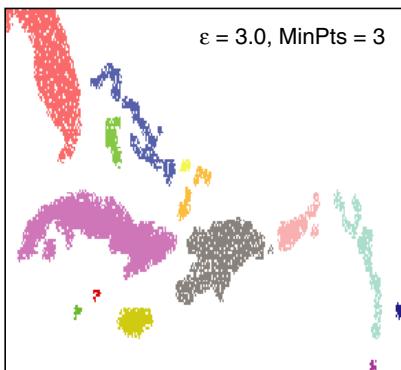
Density-Based Cluster Analysis

DBSCAN versus MajorClust: Low-Dimensional Data (continued)

Caribbean Islands, about 20.000 points:





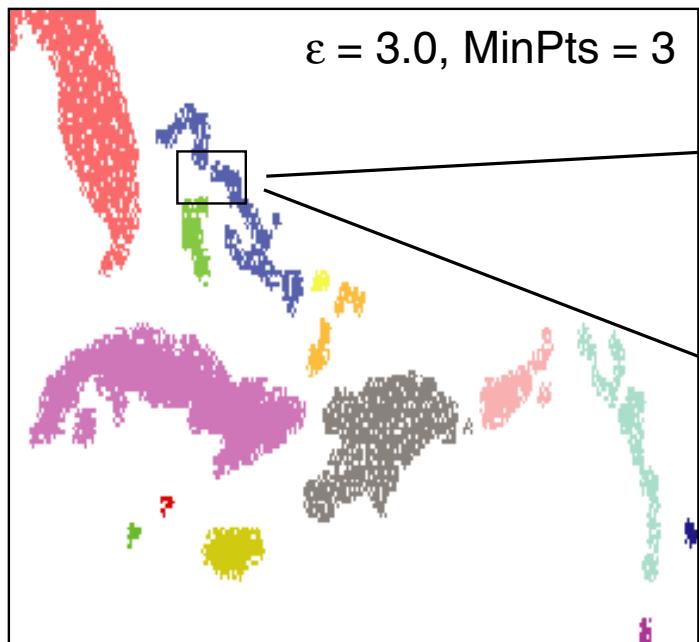
Cluster analysis by DBSCAN:



Density-Based Cluster Analysis

DBSCAN versus MajorClust: Low-Dimensional Data (continued)

The problem of finding useful ε -values for DBSCAN:



Two separate clusters will be detected.

The clusters will be merged.

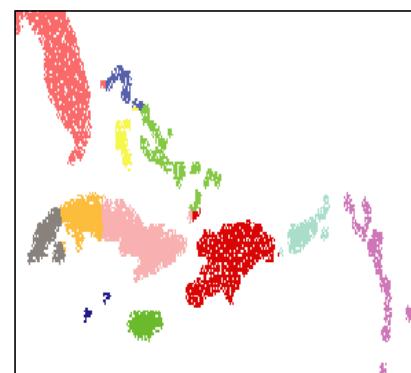
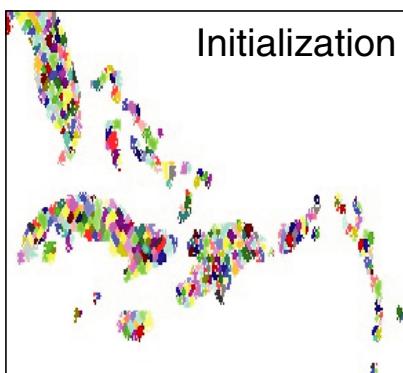
Density-Based Cluster Analysis

DBSCAN versus MajorClust: Low-Dimensional Data (continued)

Caribbean Islands, about 20.000 points:



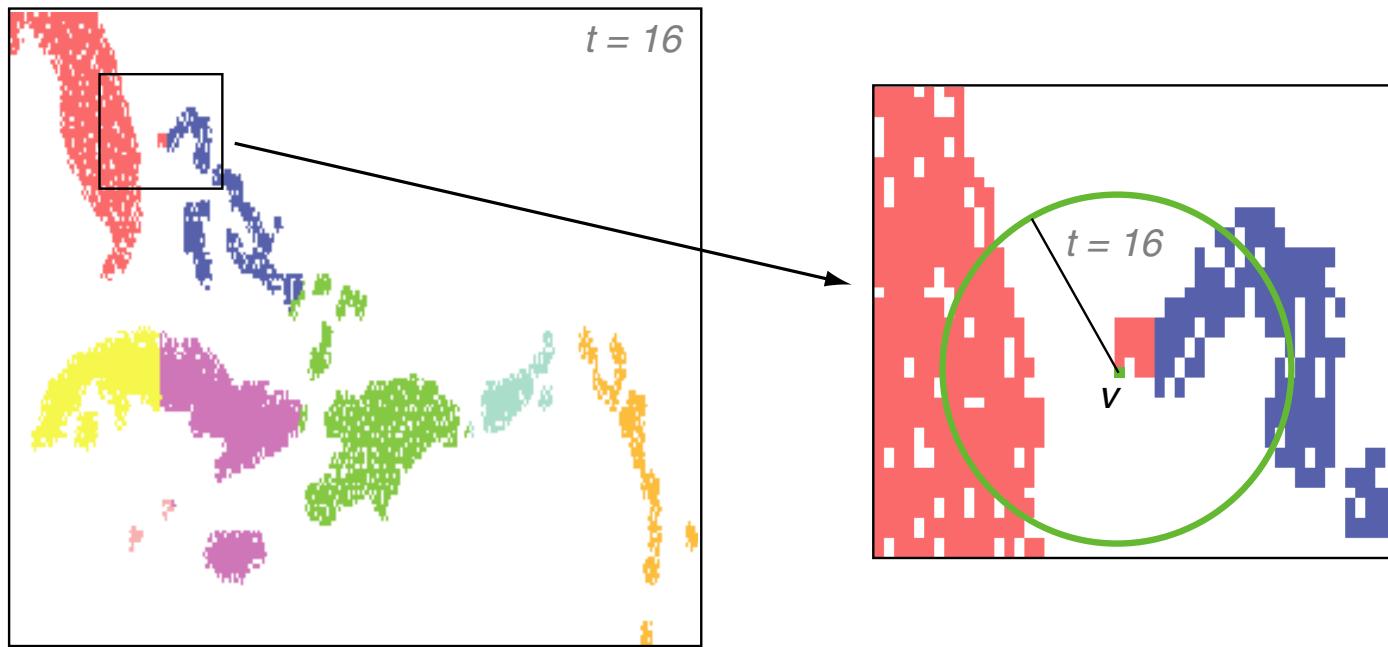
Cluster analysis by MajorClust:



Density-Based Cluster Analysis

DBSCAN versus MajorClust: Low-Dimensional Data (continued)

The problem of the global analysis approach (no restriction by means of an ε -neighborhood) in MajorClust:



Remarks:

- MajorClust is superior to DBSCAN with regard to the identification of differently dense clusters within the same clustering. DBSCAN is more flexible (= can be better adapted) than MajorClust with regard to point densities in different clusterings.
- MajorClust considers always all points of V , while DBSCAN works locally, i.e., on small subsets of V .

Density-Based Cluster Analysis

DBSCAN versus MajorClust: High-Dimensional Data

Typical document categorization setting:

- $10^4 - 10^5$ documents
- 10 - 100 categories: politics, culture, economics, etc.
- documents belong to one category
- dimension of the feature space > 10 000

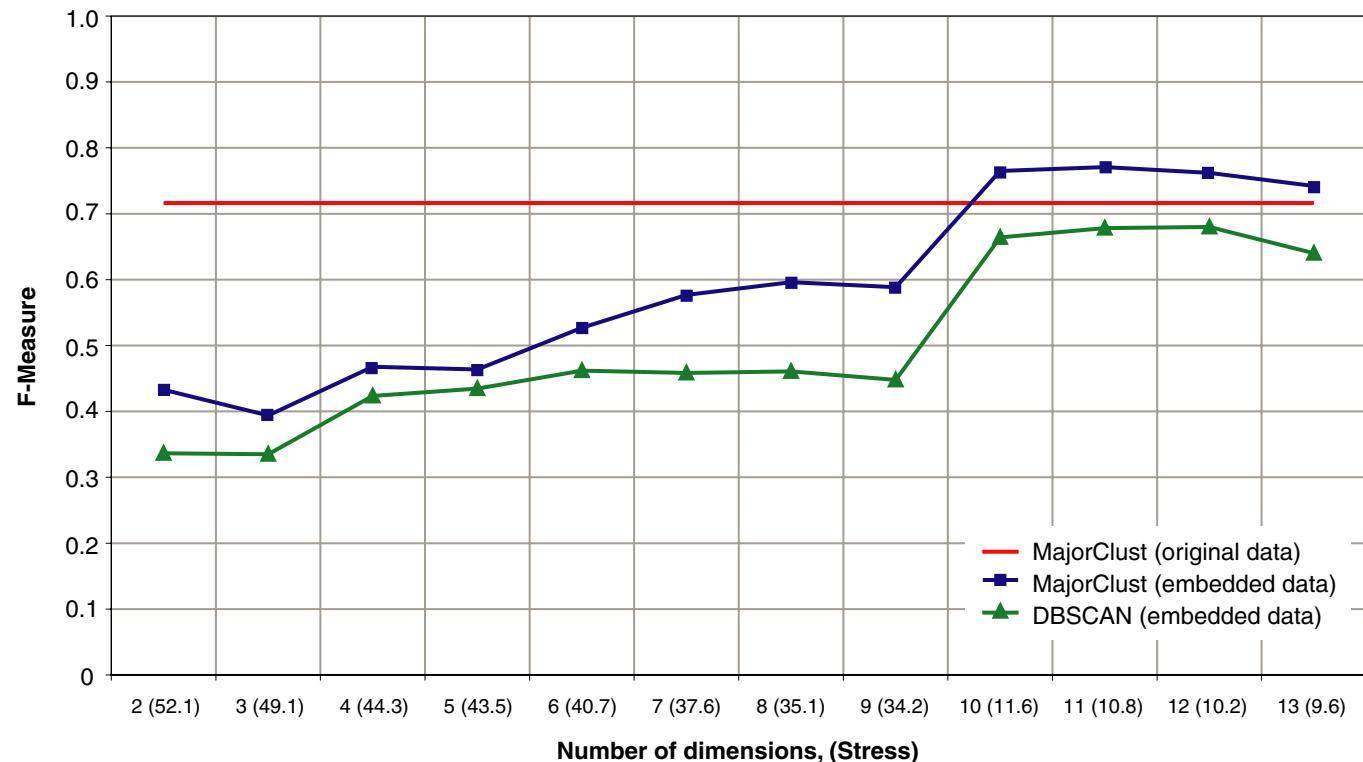
DBSCAN:

- degenerates with increasing number of dimensions
- the degeneration is rooted in the computation of the ε -neighborhood
- dimension reduction provides a way out, e.g. by embedding the data with multi-dimensional scaling, MDS

Density-Based Cluster Analysis

DBSCAN versus MajorClust: High-Dimensional Data (continued)

Classification effectiveness (F measure) over dimension number:



[Stein/Busch 2005]

Remarks:

- ❑ Usually, the neighborhood search in high-dimensional spaces cannot be solved efficiently. Given p dimensions with p about 10 or larger, an exhaustive search, i.e., a linear scan of all feature vectors will be more efficient than the application of a space partitioning data structure (quad-tree, k-d tree, etc,) or a data partitioning data structure (R -tree, Rf -tree, X -tree, etc.).
- ❑ DBSCAN employs the R -tree data structure to compute ε -neighborhoods. This data structure accomplishes the major part of the DBSCAN cluster analysis approach and is ideally suited for treating low-dimensional data efficiently. The application of DBSCAN to high-dimensional data either requires an embedding into a low-dimensional space or to accept the runtime for a naive construction of ε -neighborhoods.
- ❑ Neighborhood search in high-dimensional spaces can be addressed with approximate methods such as locality sensitive hashing (LSH), or Fuzzy fingerprinting. [Weber 1999] [Gionis/Indyk/Motwani 1999-2004] [Stein 2005-2007] [Andoni 2009]