

# Chapter ML:II (continued)

## II. Machine Learning Basics

- ❑ Regression
- ❑ Concept Learning: Search in Hypothesis Space
- ❑ Concept Learning: Search in Version Space
- ❑ Measuring Performance

# Measuring Performance

## True Misclassification Rate

### Definition 8 (True Misclassification Rate)

Let  $X$  be a feature space with a finite number of elements. Moreover, let  $C$  be a set of classes, let  $y : X \rightarrow C$  be a classifier, and let  $c$  be the target concept to be learned. Then the true misclassification rate, denoted as  $Err^*(y)$ , is defined as follows:

$$Err^*(y) = \frac{|\{\mathbf{x} \in X : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|X|}$$

# Measuring Performance

## True Misclassification Rate

### Definition 8 (True Misclassification Rate)

Let  $X$  be a feature space with a finite number of elements. Moreover, let  $C$  be a set of classes, let  $y : X \rightarrow C$  be a classifier, and let  $c$  be the target concept to be learned. Then the true misclassification rate, denoted as  $Err^*(y)$ , is defined as follows:

$$Err^*(y) = \frac{|\{\mathbf{x} \in X : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|X|}$$

Problem:

- Usually the *total function*  $c$  is unknown.

Solution:

- **Estimation** of  $Err^*(y)$  with  $Err(y, D_s)$ , i.e., evaluating  $y$  on a subset  $D_s \subseteq D$  of carefully chosen examples  $D$ . Recall that for the feature vectors in  $D$  the target concept  $c$  is known.

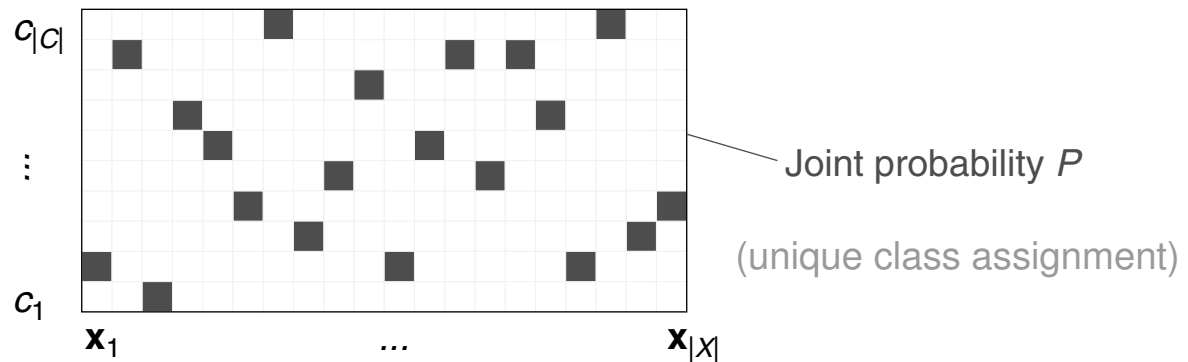
## Remarks:

- ❑ Instead of the term “true misclassification rate” we may also use the term “true misclassification error” or simply “true error”.
- ❑ The English word “rate” can be used to denote both the mathematical concept of a flow quantity (a change of a quantity per time unit) as well as the mathematical concept of a *portion*, a *percentage*, or a *ratio*, which has a stationary (= time-independent) semantics. Note that the latter semantics is meant here when talking about the misclassification rate.
- ❑ Unfortunately, the German word „Rate“ is often (mis)used to denote the mathematical concept of a portion, a percentage, or a ratio. Taking a precise mathematical standpoint, the correct German words are „Anteil“ or „Quote“. I.e., a semantically correct translation of misclassification rate is „Missklassifikationsanteil“, and not „Missklassifikationsrate“.

# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

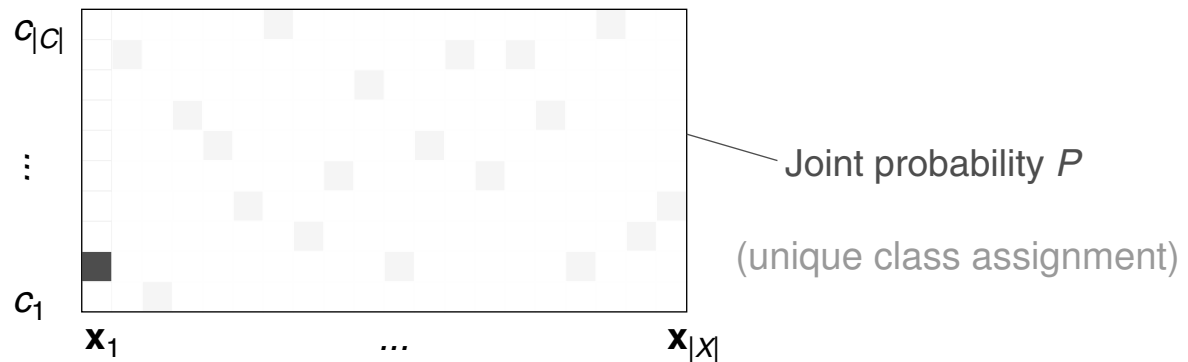
Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:



# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

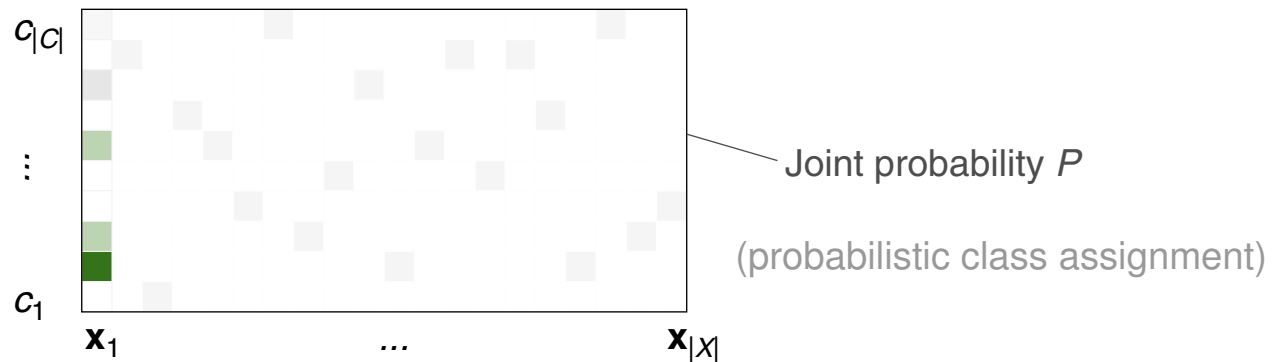
Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:



# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

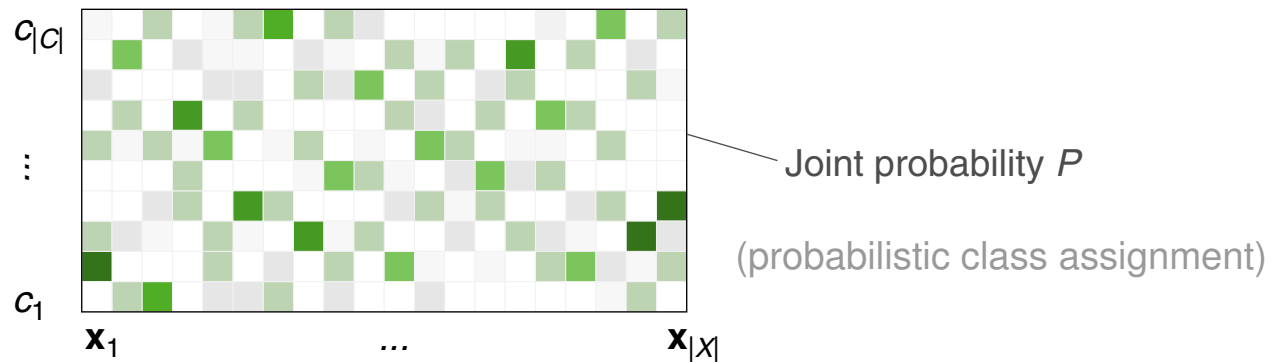
Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:



# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:






# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation


Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:



# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:




$$\underline{Err^*(y)} = \sum_{\mathbf{x} \in X} \sum_{c \in C} P(\mathbf{x}, c) \cdot I(y(\mathbf{x}), c), \quad \text{with } I(y(\mathbf{x}), c) = \begin{cases} 0 & \text{if } y(\mathbf{x}) = c \\ 1 & \text{otherwise} \end{cases}$$

# Measuring Performance

## True Misclassification Rate: Probabilistic Foundation

Let  $X$  and  $C$  be defined as before. Moreover, let  $P$  be a probability measure on  $X \times C$ . Then  $P(\mathbf{x}, c)$  (precisely:  $P(\mathcal{H} = \mathbf{x}, \mathcal{C} = c)$ ) denotes the probability that feature vector  $\mathbf{x} \in X$  belongs to class  $c \in C$ . Illustration:



$$\underline{Err^*(y)} = \sum_{\mathbf{x} \in X} \sum_{c \in C} P(\mathbf{x}, c) \cdot I(y(\mathbf{x}), c), \quad \text{with } I(y(\mathbf{x}), c) = \begin{cases} 0 & \text{if } y(\mathbf{x}) = c \\ 1 & \text{otherwise} \end{cases}$$

$D = \{(\mathbf{x}_1, c_1), \dots, (\mathbf{x}_n, c_n)\} \subseteq X \times C$  is a set of examples whose elements are drawn independently and according to the same  $P$ .

## Remarks:

- ❑  $\mathcal{H}$  and  $\mathcal{C}$  are random variables with domains  $X$  and  $C$  respectively.
- ❑ Let  $A$  and  $B$  denote two events, e.g.,  $A = “\mathcal{H} = \mathbf{x}”$  and  $B = “\mathcal{C} = c”$ . Then the following expressions are syntactic variants, i.e., they are semantically equivalent:  $P(A, B)$ ,  $P(A \text{ and } B)$ ,  $P(A \wedge B)$ .
- ❑ The function  $c(\mathbf{x})$  is modeled as random variable,  $\mathcal{C}$ , since in the real world the classification of a feature vector  $\mathbf{x}$  may not be deterministic but the result of a random (measuring) process. Keyword: label noise.
- ❑  $P(\mathbf{x}, c) = P(c, \mathbf{x}) = P(c \mid \mathbf{x}) \cdot P(\mathbf{x})$ , where  $P(\mathbf{x})$  denotes the a-priori probability for observing  $\mathbf{x}$ , while  $P(c \mid \mathbf{x})$  denotes the *likelihood* for the event of  $c$  given the event of  $\mathbf{x}$ .
- ❑ The elements in  $D$  are considered as random variables that are both independent of each other and identically distributed. This property of a set of random variables is abbreviated with “i.i.d.”
- ❑ If the elements in  $D$  or  $D_s$  were not chosen according to  $P$ , then  $Err(y, D_s)$  could not be used as an estimation of  $Err^*(y)$ . Keyword: sample selection bias
- ❑ See the definition of a probability measure in [\[ML:IV Probability Basics\]](#).

# Measuring Performance

## Training Error [True Misclassification Rate]

- $D = \{(\mathbf{x}_1, c(\mathbf{x}_1)), \dots, (\mathbf{x}_n, c(\mathbf{x}_n))\} \subseteq X \times C$  is a set of examples.
- $D_{tr} = D$  is the training set.
- $y : X \rightarrow C$  is a classifier learned on the basis of  $D_{tr}$ .

Training error = misclassification rate with respect to  $D_{tr}$  :

$$Err(y, D_{tr}) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D_{tr} : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D_{tr}|}$$

# Measuring Performance

## Training Error [True Misclassification Rate]

- $D = \{(\mathbf{x}_1, c(\mathbf{x}_1)), \dots, (\mathbf{x}_n, c(\mathbf{x}_n))\} \subseteq X \times C$  is a set of examples.
- $D_{tr} = D$  is the training set.
- $y : X \rightarrow C$  is a classifier learned on the basis of  $D_{tr}$ .

Training error = misclassification rate with respect to  $D_{tr}$  :

$$Err(y, D_{tr}) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D_{tr} : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D_{tr}|}$$

Problems:

- $Err(y, D_{tr})$  is based on examples that are also exploited to learn  $y$ .
- $Err(y, D_{tr})$  quantifies **memorization** but **not** the **generalization** capability of  $y$ .
- $Err(y, D_{tr})$  is an optimistic estimation, i.e., it is constantly lower compared to the error incurred when applying  $y$  in the wild.

# Measuring Performance

## 2-Fold Cross-Validation (Holdout Estimation) [True Misclassification Rate]

- $D = \{(\mathbf{x}_1, c(\mathbf{x}_1)), \dots, (\mathbf{x}_n, c(\mathbf{x}_n))\} \subseteq X \times C$  is a set of examples.
- $D_{tr} \subset D$  is the training set.
- $y : X \rightarrow C$  is a classifier learned on the basis of  $D_{tr}$ .
- $D_{ts} \subset D$  with  $D_{ts} \cap D_{tr} = \emptyset$  is a test set.

Holdout estimation = misclassification rate with respect to  $D_{ts}$  :

$$Err(y, D_{ts}) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D_{ts} : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D_{ts}|}$$

# Measuring Performance

## 2-Fold Cross-Validation (Holdout Estimation) [True Misclassification Rate]

- $D = \{(\mathbf{x}_1, c(\mathbf{x}_1)), \dots, (\mathbf{x}_n, c(\mathbf{x}_n))\} \subseteq X \times C$  is a set of examples.
- $D_{tr} \subset D$  is the training set.
- $y : X \rightarrow C$  is a classifier learned on the basis of  $D_{tr}$ .
- $D_{ts} \subset D$  with  $D_{ts} \cap D_{tr} = \emptyset$  is a test set.

Holdout estimation = misclassification rate with respect to  $D_{ts}$  :

$$Err(y, D_{ts}) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D_{ts} : y(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D_{ts}|}$$

Requirements:

- $D_{tr}$  and  $D_{ts}$  must be **goverend by the same distribution**.
- $D_{tr}$  and  $D_{ts}$  should have similar sizes.



## Remarks:

- ❑ A typical value for splitting  $D$  into training set  $D_{tr}$  and test set  $D_{ts}$  is 2:1.
- ❑ When splitting  $D$  into  $D_{tr}$  and  $D_{ts}$  one has to ensure that the underlying distribution is maintained. Keywords: stratification, sample selection bias

# Measuring Performance

## $k$ -Fold Cross-Validation [Holdout Estimation]

- Form  $k$  test sets by splitting  $D$  into disjoint sets  $D_1, \dots, D_k$  of similar size.
- For  $i = 1, \dots, k$  do:
  1.  $y_i : X \rightarrow C$  is a classifier learned on the basis of  $D \setminus D_i$
  2. 
$$Err(y_i, D_i) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D_i : y_i(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D_i|}$$

Cross-validated misclassification rate:

$$Err_{cv}(y, D) = \frac{1}{k} \sum_{i=1}^k Err(y_i, D_i)$$

# Measuring Performance

## $n$ -Fold Cross-Validation (Leave One Out)

Special case with  $k = n$  :

- Determine the cross-validated misclassification rate for  $D \setminus D_i$  where  $D_i = \{(\mathbf{x}_i, c(\mathbf{x}_i))\}$ ,  $i \in \{1, \dots, n\}$  .

# Measuring Performance

## $n$ -Fold Cross-Validation (Leave One Out)

Special case with  $k = n$  :

- Determine the cross-validated misclassification rate for  $D \setminus D_i$  where  $D_i = \{(\mathbf{x}_i, c(\mathbf{x}_i))\}$ ,  $i \in \{1, \dots, n\}$  .

Problems:

- High computational effort if  $D$  is large.
- Singleton test sets ( $|D_i| = 1$ ) are never stratified since they contain a single class only.

## Remarks:

- ❑ For large  $k$  the set  $D \setminus D_i$  is of similar size as  $D$ . Hence  $Err(y_i, D_i)$  is close to  $Err(y, D)$ , where  $y$  is the classifier learned on the basis of the entire set  $D$ .
- ❑  $n$ -fold cross-validation is a special case of exhaustive cross-validation methods, which learn and test on all possible ways to divide the original sample into a training and a validation set.  
[\[Wikipedia\]](#)

# Measuring Performance

## Bootstrapping [Holdout Estimation]

Resampling the example set  $D$ :

- For  $j = 1, \dots, l$  do:
  1. Form training set  $D_j$  by drawing  $m$  examples from  $D$  with replacement.
  2.  $y_j : X \rightarrow C$  is a classifier learned on the basis of  $D_j$
  3.  $Err(y_j, D \setminus D_j) = \frac{|\{(\mathbf{x}, c(\mathbf{x})) \in D \setminus D_j : y_j(\mathbf{x}) \neq c(\mathbf{x})\}|}{|D \setminus D_j|}$

Bootstrapped misclassification rate:

$$Err_{bt}(y, D) = \frac{1}{l} \sum_{j=1}^l Err(y_j, D \setminus D_j)$$

## Remarks:

- ❑ Let  $|D| = n$ . The probability that an example is not considered is  $(1 - 1/n)^m$ . Similarly, the probability that an example is considered at least once is  $1 - (1 - 1/n)^m$ .
- ❑ If  $m$  gets closer to  $n$ , then  $1 - (1 - 1/n)^m \approx 1 - 1/e \approx 0.632$ . I.e., each training set contains about 63.2% of the examples in  $D$ .
- ❑ The classifiers  $y_1, \dots, y_l$  can be used in a combined fashion, called *ensemble*, where the class is determined by means of a majority decision:

$$y(\mathbf{x}) = \operatorname{argmax}_{c \in C} |\{j \in \{1, \dots, l\} : y_j(\mathbf{x}) = c\}|$$

# Measuring Performance

## Misclassification Costs [Holdout Estimation]

Use of a cost measure for the misclassification of a feature vector  $\mathbf{x}$  in class  $c'$  instead of in class  $c$ :

$$\text{cost}(c' \mid c) \begin{cases} \geq 0 & \text{if } c' \neq c \\ = 0 & \text{otherwise} \end{cases}$$

Estimation of  $\text{Err}_{\text{cost}}^*(y)$  based on a sample  $D_s \subseteq D$ :

$$\text{Err}_{\text{cost}}(y, D_s) = \frac{1}{|D_s|} \cdot \sum_{(\mathbf{x}, c(\mathbf{x})) \in D_s} \text{cost}(y(\mathbf{x}) \mid c(\mathbf{x}))$$



## Remarks:

- ❑ The misclassification rate  $Err$  is a special case of  $Err_{cost}$  with  $cost(c' | c) = 1$  for  $c' \neq c$ .