

# Chapter ML:VI

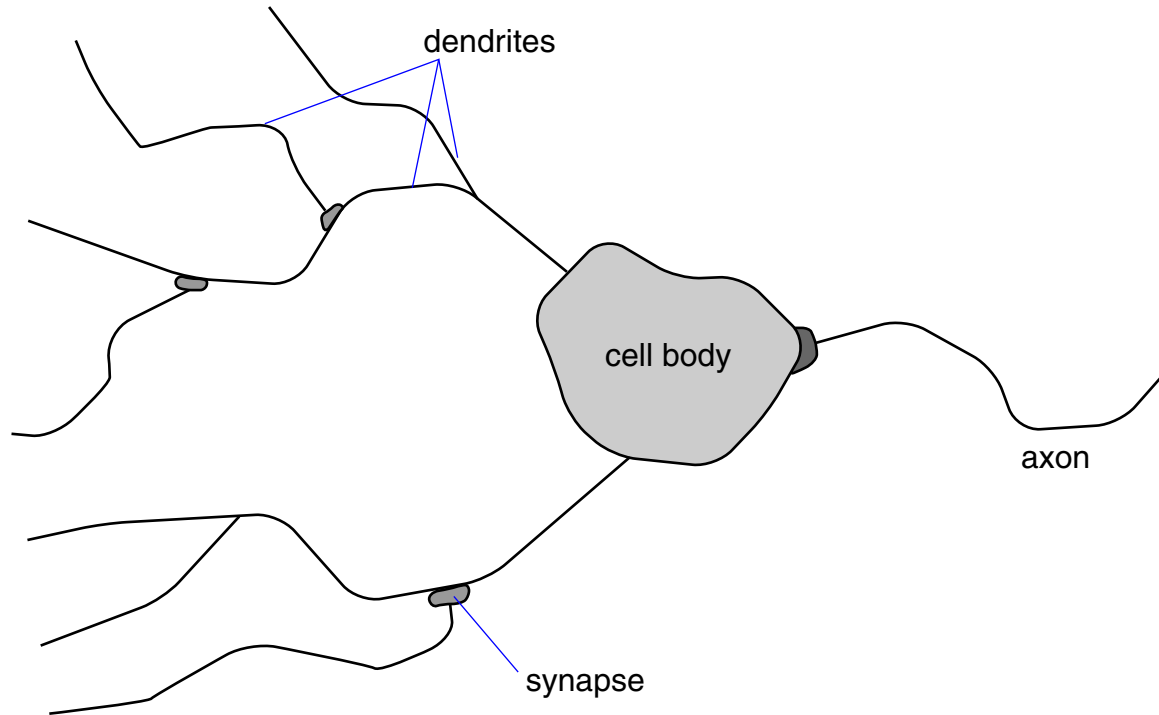
## VI. Neural Networks

- ❑ Perceptron Learning
- ❑ Gradient Descent
- ❑ Multilayer Perceptron
- ❑ Radial Basis Functions

# Perceptron Learning

## The Biological Model

Simplified model of a neuron:



# Perceptron Learning

## The Biological Model (continued)

### Neuron characteristics:

- ❑ The numerous dendrites of a neuron serve as its input channels for electrical signals.
- ❑ At particular contact points between the dendrites, the so-called synapses, electrical signals can be initiated.
- ❑ A synapse can initiate signals of different strengths, where the strength is encoded by the frequency of a pulse train.
- ❑ The cell body of a neuron accumulates the incoming signals.
- ❑ If a particular stimulus threshold is exceeded, the cell body generates a signal, which is output via the axon.
- ❑ The processing of the signals is unidirectional. (from left to right in the figure)

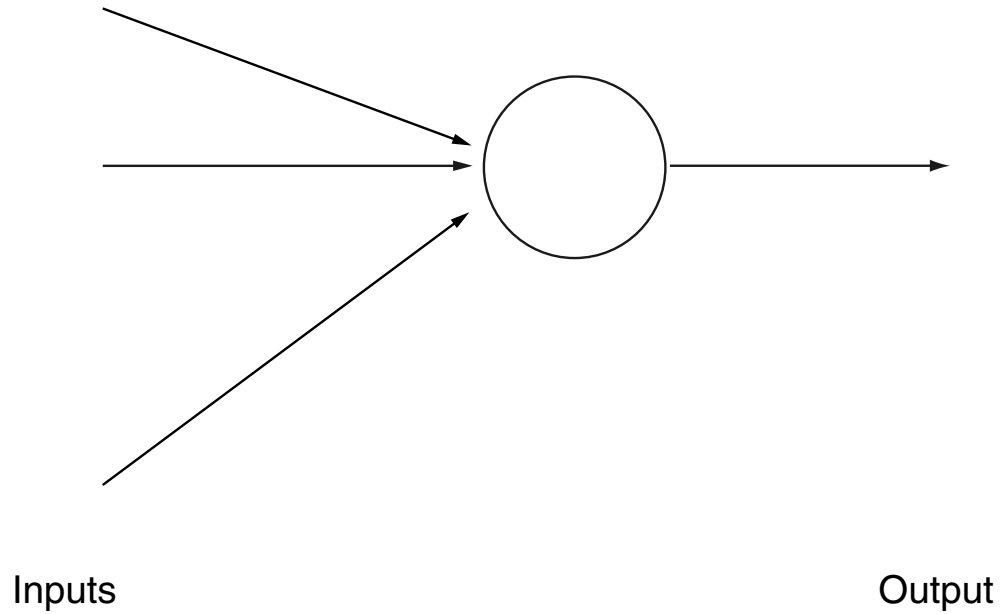
# Perceptron Learning

## History

- 1943 Warren McCulloch and Walter Pitts present a model of the neuron.
- 1949 Donald Hebb postulates a new learning paradigm: reinforcement only for active neurons. (those neurons that are involved in a decision process)
- 1958 Frank Rosenblatt develops the perceptron model.
- 1962 Rosenblatt proves the perceptron convergence theorem.
- 1969 Marvin Minsky and Seymour Papert publish a book on the limitations of the perceptron model.
- 1970
- :
- 1985
- 1986 David Rumelhart and James McClelland present the multilayer perceptron.

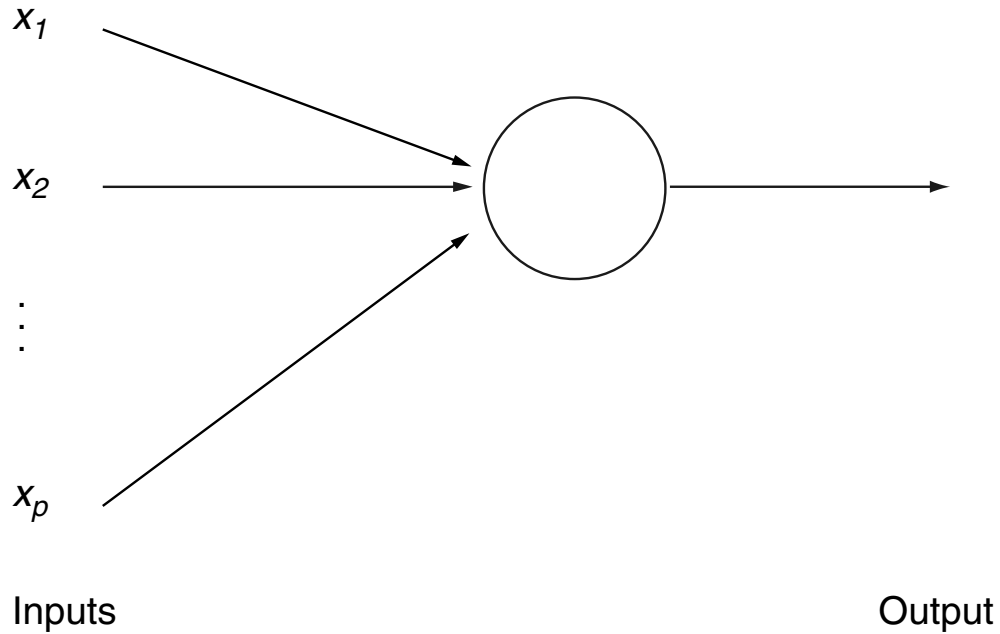
# Perceptron Learning

The Perceptron of Rosenblatt [1958]



# Perceptron Learning


The Perceptron of Rosenblatt [1958]



$$x_j, w_j \in \mathbf{R}, \quad j = 1 \dots p$$

# Perceptron Learning

The Perceptron of Rosenblatt [1958]



$$x_j, w_j \in \mathbf{R}, \quad j = 1 \dots p$$

## Remarks:

- ❑ The perceptron of Rosenblatt is based on the neuron model of McCulloch and Pitts.
- ❑ The perceptron is a “feed forward system”.



# Perceptron Learning

## Specification of Classification Problems [ML Introduction]

Characterization of the model (model world):

- $X$  is a set of feature vectors, also called feature space.  $X \subseteq \mathbf{R}^p$
- $C = \{0, 1\}$  is a set of classes.  $C = \{-1, 1\}$  in the regression setting.
- $c : X \rightarrow C$  is the ideal classifier for  $X$ .  $c$  is approximated by  $y$  (perceptron).
- $D = \{(\mathbf{x}_1, c(\mathbf{x}_1)), \dots, (\mathbf{x}_n, c(\mathbf{x}_n))\} \subseteq X \times C$  is a set of examples.

How could the hypothesis space  $H$  look like?

# Perceptron Learning

## Computation in the Perceptron [Regression]

If  $\sum_{j=1}^p w_j x_j \geq \theta$  then  $y(\mathbf{x}) = 1$ , and


if  $\sum_{j=1}^p w_j x_j < \theta$  then  $y(\mathbf{x}) = 0$ .

# Perceptron Learning

## Computation in the Perceptron [Regression]

If  $\sum_{j=1}^p w_j x_j \geq \theta$  then  $y(\mathbf{x}) = 1$ , and

if  $\sum_{j=1}^p w_j x_j < \theta$  then  $y(\mathbf{x}) = 0$ .



where  $\sum_{j=1}^p w_j x_j = \mathbf{w}^T \mathbf{x}$ . (or other notations for the scalar product)


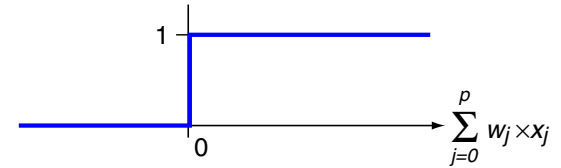
→ A hypothesis is determined by  $\theta, w_1, \dots, w_p$ .

# Perceptron Learning

## Computation in the Perceptron (continued)

$$y(\mathbf{x}) = \text{heaviside}\left(\sum_{j=1}^p w_j x_j - \theta\right)$$

$$= \text{heaviside}\left(\sum_{j=0}^p w_j x_j\right) \quad \text{with } w_0 = -\theta, x_0 = 1$$



→ A hypothesis is determined by  $w_0, w_1, \dots, w_p$ .

## Remarks:

- ❑ If the weight vector is extended by  $w_0 = -\theta$ , and, if the feature vectors are extended by the constant feature  $x_0 = 1$ , the learning algorithm gets a canonical form. Implementations of neural networks introduce this extension often implicitly.
- ❑ Be careful with regard to the dimensionality of the weight vector: it is always denoted as  $w$  here, regardless of the fact whether the  $w_0$ -dimension, with  $w_0 = -\theta$ , is included.
- ❑ The function *heaviside* is named after the mathematician Oliver Heaviside.  
[Heaviside: [step function](#) [O. Heaviside](#)]

# Perceptron Learning

## Weight Adaptation [Algorithms: *BGD* *IGD*]

Algorithm:	<i>PT</i>	Perceptron Training
Input:	$D$	Training examples $(\mathbf{x}, c(\mathbf{x}))$ with $ \mathbf{x}  = p + 1$ , $c(\mathbf{x}) \in \{0, 1\}$ .
	$\eta$	Learning rate, a small positive constant.
Internal:	$y(D)$	Set of $y(\mathbf{x})$ -values computed from the elements $\mathbf{x}$ in $D$ given some $\mathbf{w}$ .
Output:	$\mathbf{w}$	Weight vector.

$PT(D, \eta)$

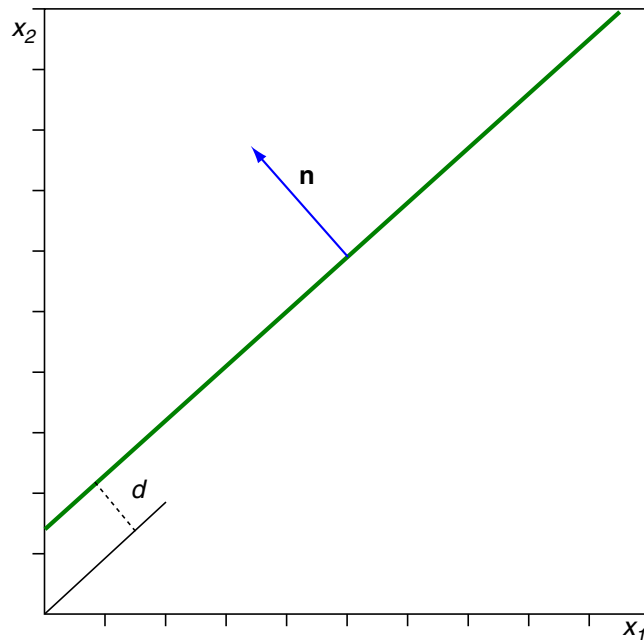
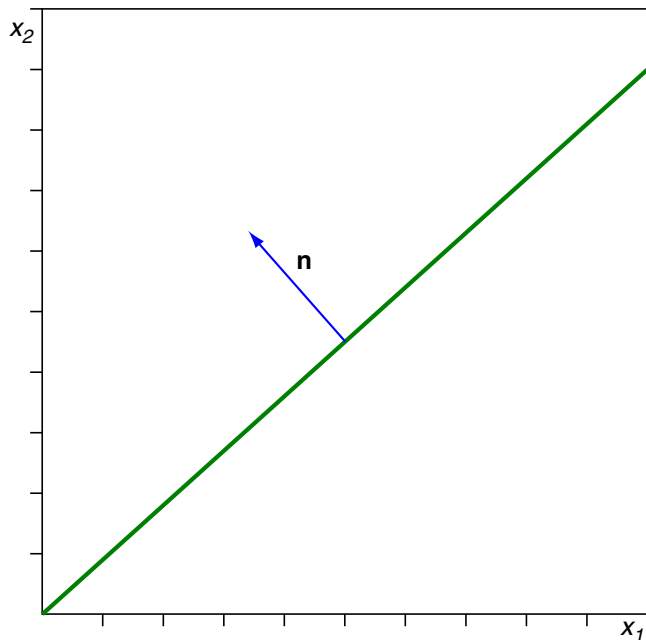
1. *initialize\_random\_weights*( $\mathbf{w}$ ),  $t = 0$
2. **REPEAT**
3.    $t = t + 1$
4.    $(\mathbf{x}, c(\mathbf{x})) = \text{random\_select}(D)$
5.    $\text{error} = c(\mathbf{x}) - \text{heaviside}(\mathbf{w}^T \mathbf{x})$  //  $c(\mathbf{x}) \in \{0, 1\}$ , *heaviside*  $\in \{0, 1\}$ , *error*  $\in \{0, 1, -1\}$
6.    $\Delta \mathbf{w} = \eta \cdot \text{error} \cdot \mathbf{x}$
7.    $\mathbf{w} = \mathbf{w} + \Delta \mathbf{w}$
8. **UNTIL**(*convergence*( $D, y(D)$ )) **OR**  $t > t_{\max}$
9. *return*( $\mathbf{w}$ )

## Remarks:

- ❑ The variable  $t$  denotes the time. The learning algorithm gets an example presented at each point in time and, as a consequence, may adapt the weight vector.
- ❑ The weight adaptation rule compares the true class  $c(\mathbf{x})$  (the ground truth) to the class computed by the perceptron. In case of a wrong classification of a feature vector  $\mathbf{x}$ , *error* is either  $-1$  or  $+1$ , regardless of the exact numeric difference between  $c(\mathbf{x})$  and  $\mathbf{w}^T \mathbf{x}$ .
- ❑  $y(D)$  is the set of  $y(\mathbf{x})$ -values given  $\mathbf{w}$  for the elements  $\mathbf{x}$  in  $D$ .

# Perceptron Learning

## Weight Adaptation: Illustration in Input Space





Definition of an (affine) hyperplane:  $L = \{\mathbf{x} \mid \mathbf{n}^T \mathbf{x} = d\}$  [\[Wikipedia\]](#)

- ❑  $\mathbf{n}$  denotes a normal vector that is perpendicular to the hyperplane  $L$ .
- ❑ If  $\|\mathbf{n}\| = 1$  then  $|\mathbf{n}^T \mathbf{x} - d|$  gives the distance of any point  $\mathbf{x}$  to  $L$ .
- ❑ If  $\text{sgn}(\mathbf{n}^T \mathbf{x}_1 - d) = \text{sgn}(\mathbf{n}^T \mathbf{x}_2 - d)$ , then  $\mathbf{x}_1$  and  $\mathbf{x}_2$  lie on the same side of the hyperplane.



# Perceptron Learning

## Weight Adaptation: Illustration in Input Space (continued)



Definition of an (affine) hyperplane:  $\mathbf{w}^T \mathbf{x} = 0 \Leftrightarrow \sum_{j=1}^p w_j x_j = \theta = -w_0$

(hyperplane definition as before, with notation taken from the classification problem setting)

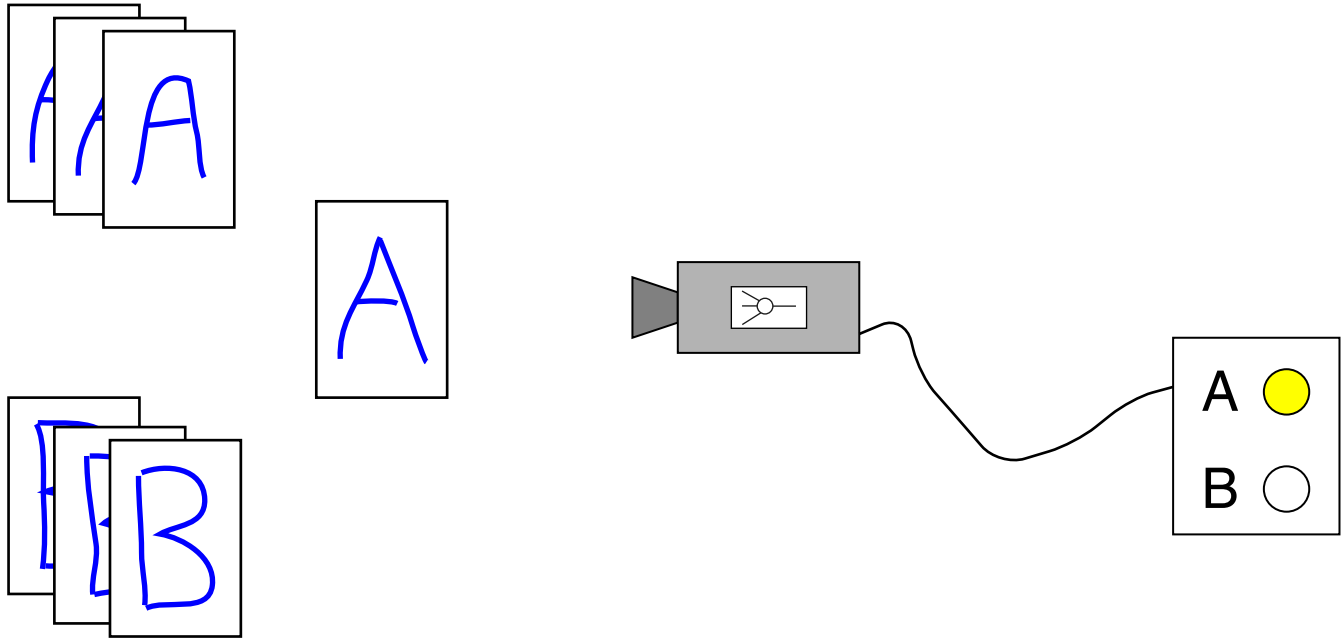
## Remarks:

- ❑ A perceptron defines a hyperplane that is perpendicular (= normal) to  $(w_1, \dots, w_p)^T$ .
- ❑  $\theta$  or  $-w_0$  specify the offset of the hyperplane from the origin, along  $(w_1, \dots, w_p)^T$  and as multiple of  $1/\|(w_1, \dots, w_p)^T\|$ .
- ❑ The set of possible weight vectors  $\mathbf{w} = (w_0, w_1, \dots, w_p)^T$  form the hypothesis space  $H$ .
- ❑ Weight adaptation means learning, and the shown learning paradigm is supervised.
- ❑ For the weight adaptation in Line 6–7 of the [PT Algorithm](#), note that if some  $x_j$  is zero,  $\Delta w_j$  will be zero as well. Keyword: Hebbian learning [\[Hebb 1949\]](#)
- ❑ Note that here (and in the following illustrations) the hyperplane movement is not the result of solving a regression problem in the  $(p + 1)$ -dimensional input-output-space, where the sum of the residuals is to be minimized.

Rather, the [PT Algorithm](#) takes each missclassified example  $\mathbf{x}$  as an event to update the hyperplane's normal vector by a fixed amount that is proportional to  $\mathbf{x}$ . In particular, the update,  $\Delta \mathbf{w}$ , does not exploit the residual associated with  $\mathbf{x}$  at time  $t$ , i.e., the absolute value of the distance of  $\mathbf{x}$  from the hyperplane is disregarded.

# Perceptron Learning

## Example



- ❑ The examples are presented to the perceptron.
- ❑ The perceptron computes a value that is interpreted as class label.

# Perceptron Learning

## Example (continued)

Encoding:

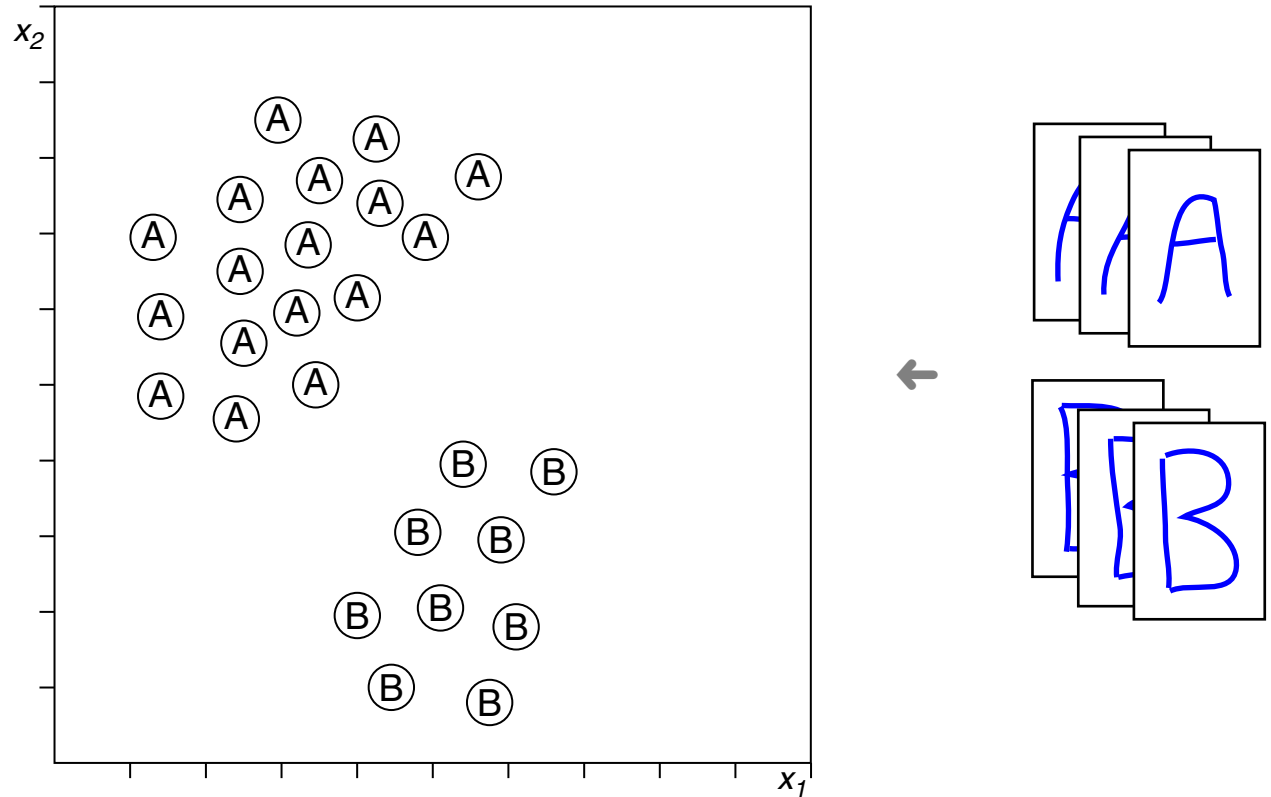
- ❑ The encoding of the examples is based on expressive features such as the number of line crossings, most acute angle, longest line, etc.
- ❑ The class label,  $c(\mathbf{x})$ , is encoded as a number. Examples from  $A$  are labeled with 1, examples from  $B$  are labeled with 0.

$$\underbrace{\begin{pmatrix} x_{1_1} \\ x_{1_2} \\ \vdots \\ x_{1_p} \end{pmatrix} \quad \cdots \quad \begin{pmatrix} x_{k_1} \\ x_{k_2} \\ \vdots \\ x_{k_p} \end{pmatrix}}_{\text{Class } A \simeq c(\mathbf{x}) = 1}$$

$$\underbrace{\begin{pmatrix} x_{l_1} \\ x_{l_2} \\ \vdots \\ x_{l_p} \end{pmatrix} \quad \cdots \quad \begin{pmatrix} x_{m_1} \\ x_{m_2} \\ \vdots \\ x_{m_p} \end{pmatrix}}_{\text{Class } B \simeq c(\mathbf{x}) = 0}$$

# Perceptron Learning

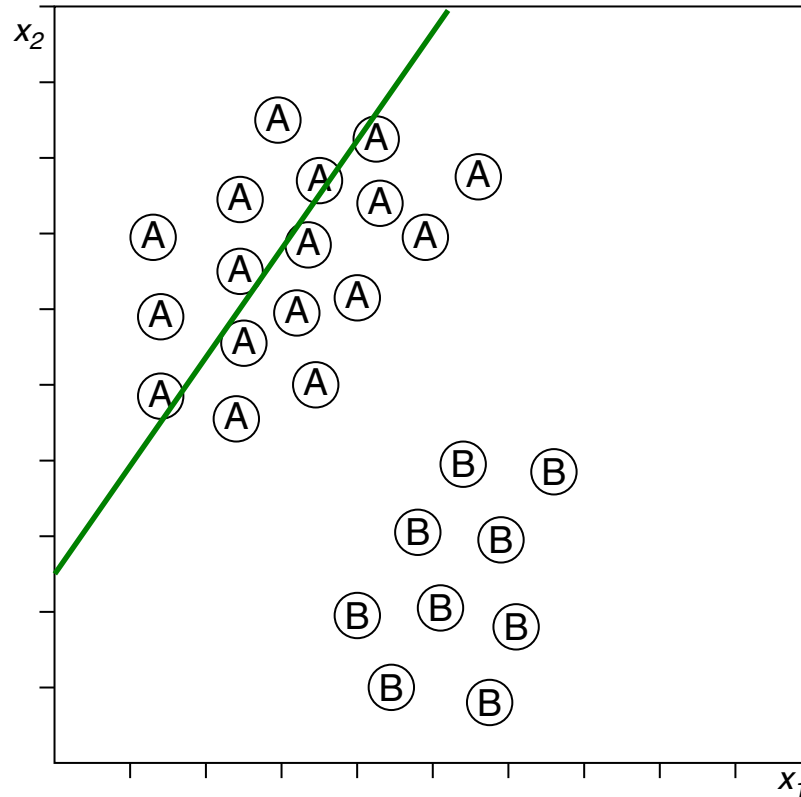
Example: Illustration in Input Space [PT Algorithm]



A possible configuration of encoded objects in the feature space  $X$ .

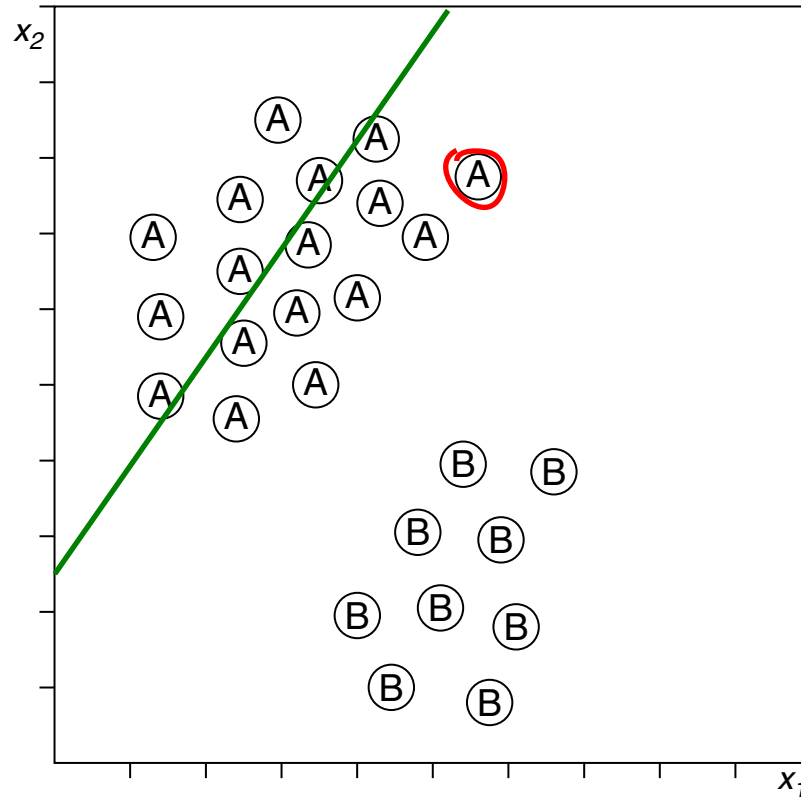
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



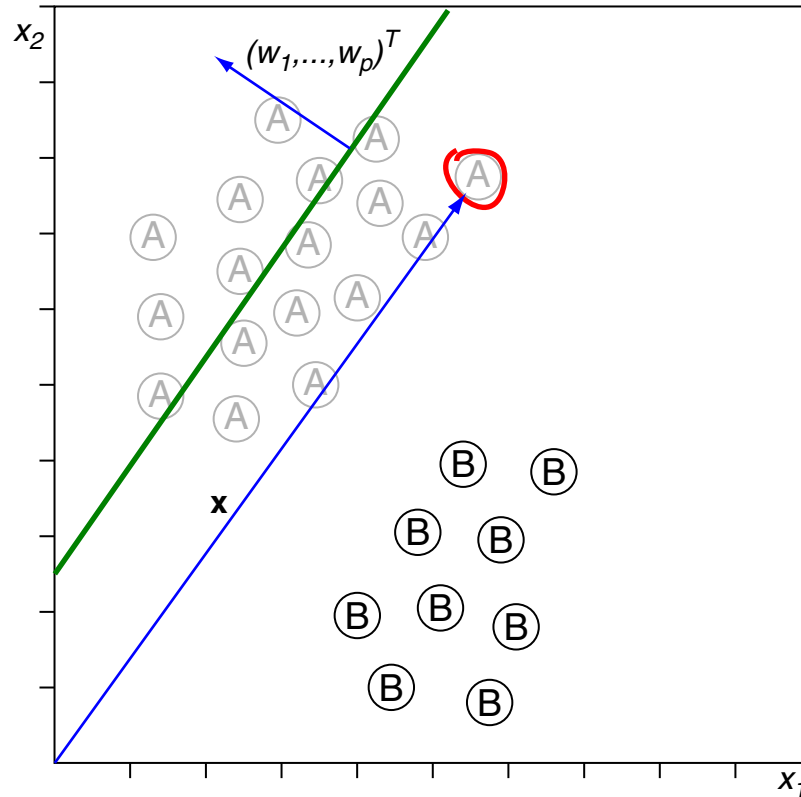
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



# Perceptron Learning

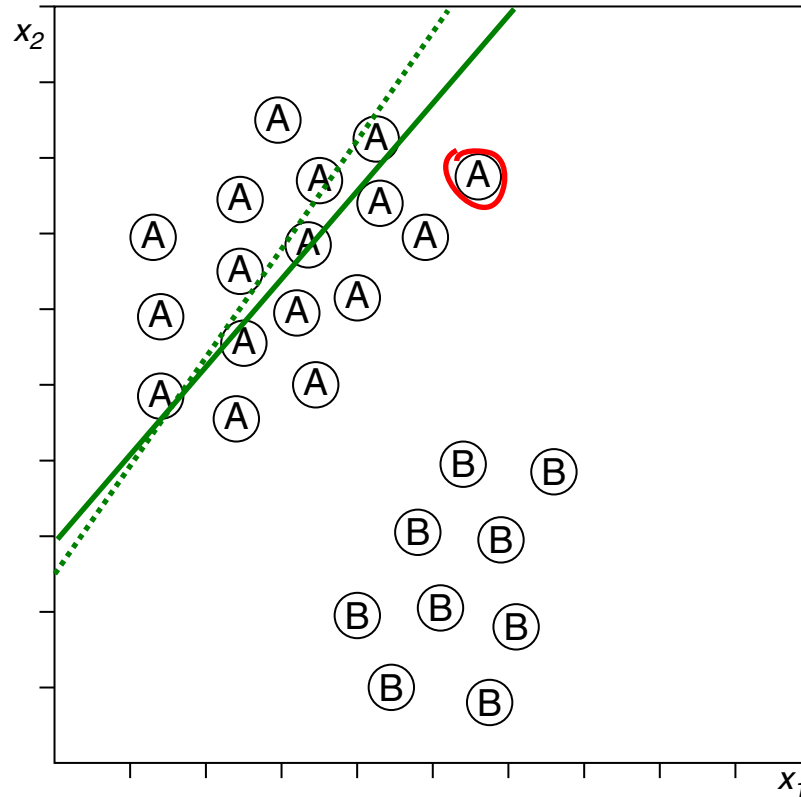
Example: Illustration in Input Space [PT Algorithm]





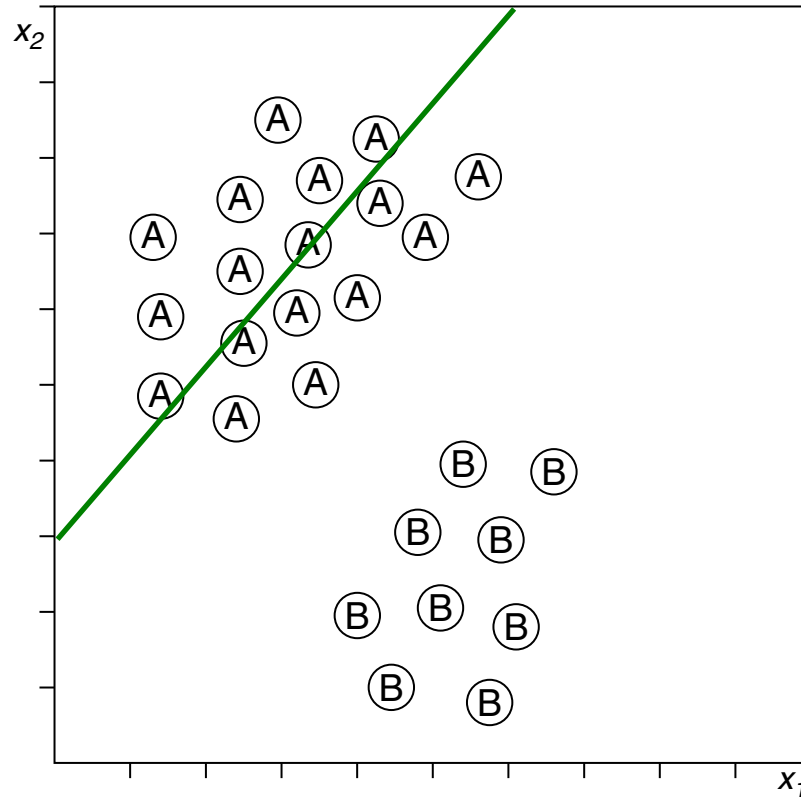
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



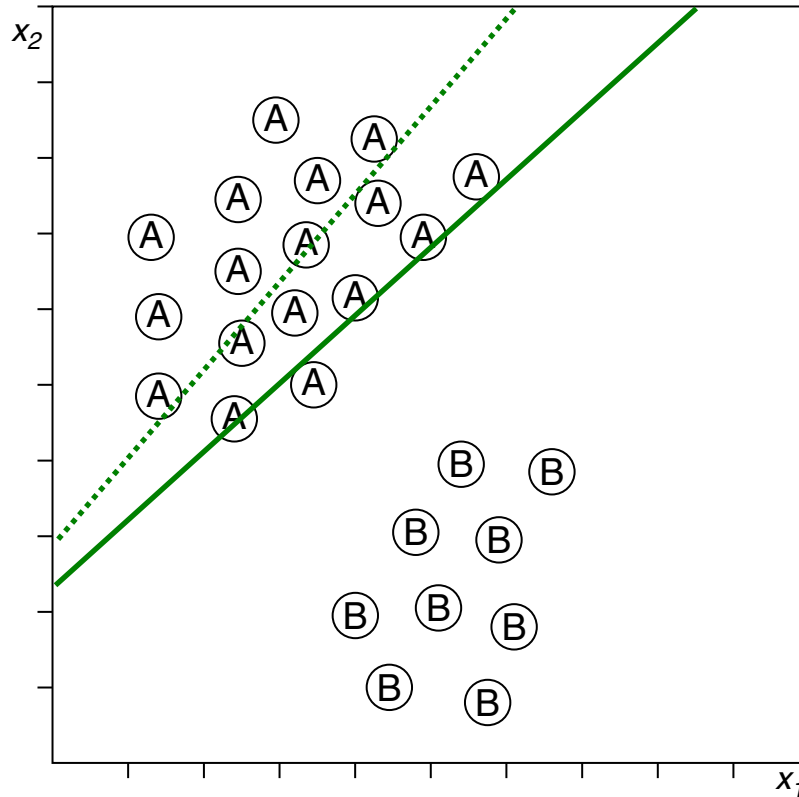
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



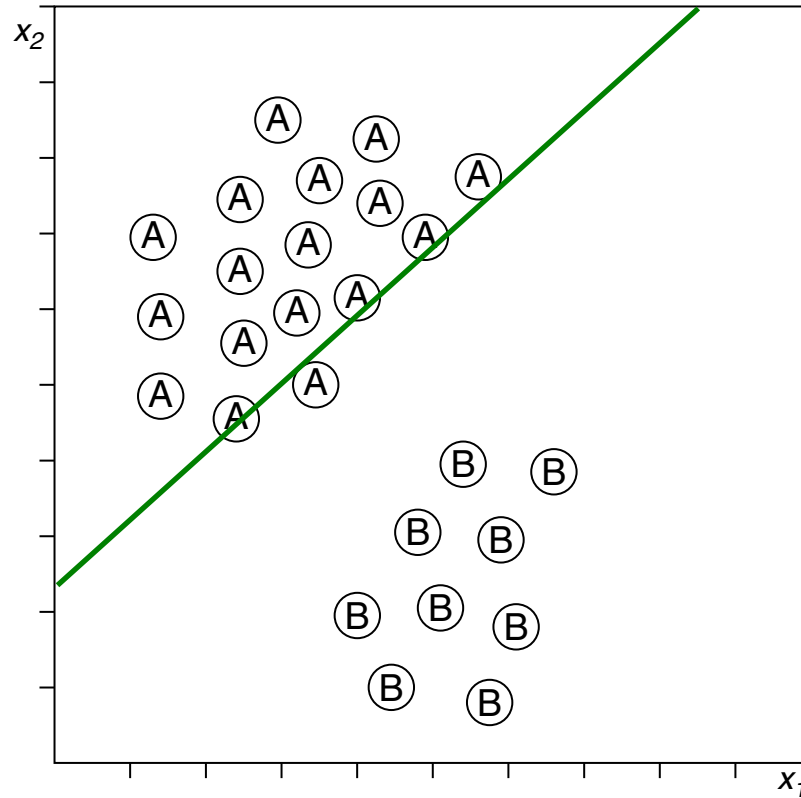
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)




# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



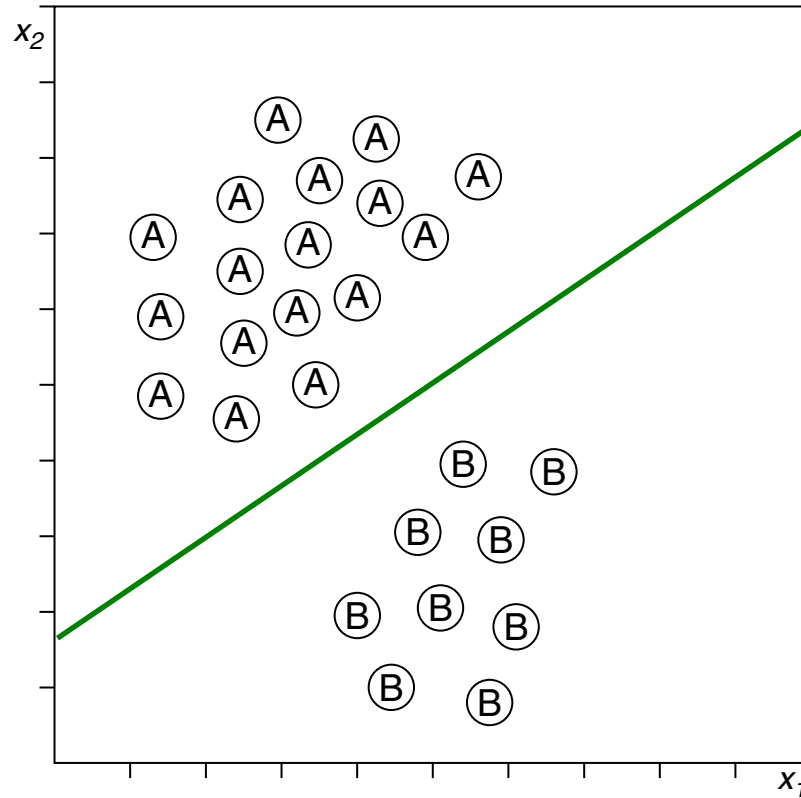
# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



# Perceptron Learning

Example: Illustration in Input Space [\[PT Algorithm\]](#)



# Perceptron Learning

## Perceptron Convergence Theorem [\[Discussion\]](#)

Questions:

1. Which kind of learning tasks can be addressed with the functions in the hypothesis space  $H$ ?
2. Can the [\*PT\* Algorithm](#) construct such a function for a given task?

# Perceptron Learning

## Perceptron Convergence Theorem [\[Discussion\]](#)

Questions:

1. Which kind of learning tasks can be addressed with the functions in the hypothesis space  $H$ ?
2. Can the [PT Algorithm](#) construct such a function for a given task?

### Theorem 1 (Perceptron Convergence [\[Rosenblatt 1962\]](#))

Let  $X_0$  and  $X_1$  be two finite sets with vectors of the form  $\mathbf{x} = (1, x_1, \dots, x_p)^T$ , let  $X_1 \cap X_0 = \emptyset$ , and let  $\hat{\mathbf{w}}$  define a separating hyperplane with respect to  $X_0$  and  $X_1$ . Moreover, let  $D$  be a set of examples of the form  $(\mathbf{x}, 0)$ ,  $\mathbf{x} \in X_0$  and  $(\mathbf{x}, 1)$ ,  $\mathbf{x} \in X_1$ . Then holds:

If the examples in  $D$  are processed with the [PT Algorithm](#), the constructed weight vector  $\mathbf{w}$  will converge within a finite number of iterations.



# Perceptron Learning

## Perceptron Convergence Theorem: Proof

### Preliminaries:

- The sets  $X_1$  and  $X_0$  are separated by a hyperplane  $\hat{\mathbf{w}}$ . The proof requires that for all  $\mathbf{x} \in X_1$  the inequality  $\hat{\mathbf{w}}^T \mathbf{x} > 0$  holds. This condition is always fulfilled, as the following consideration shows.

Let  $\mathbf{x}' \in X_1$  with  $\hat{\mathbf{w}}^T \mathbf{x}' = 0$ . Since  $X_0$  is finite, the members  $\mathbf{x} \in X_0$  have a minimum positive distance  $\delta$  with regard to the hyperplane  $\hat{\mathbf{w}}$ . Hence,  $\hat{\mathbf{w}}$  can be moved by  $\frac{\delta}{2}$  towards  $X_0$ , resulting in a new hyperplane  $\hat{\mathbf{w}}'$  that still fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} < 0$  for all  $\mathbf{x} \in X_0$ , but that now also fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X_1$ .

# Perceptron Learning

## Perceptron Convergence Theorem: Proof

### Preliminaries:

- The sets  $X_1$  and  $X_0$  are separated by a hyperplane  $\hat{\mathbf{w}}$ . The proof requires that for all  $\mathbf{x} \in X_1$  the inequality  $\hat{\mathbf{w}}^T \mathbf{x} > 0$  holds. This condition is always fulfilled, as the following consideration shows.

Let  $\mathbf{x}' \in X_1$  with  $\hat{\mathbf{w}}^T \mathbf{x}' = 0$ . Since  $X_0$  is finite, the members  $\mathbf{x} \in X_0$  have a minimum positive distance  $\delta$  with regard to the hyperplane  $\hat{\mathbf{w}}$ . Hence,  $\hat{\mathbf{w}}$  can be moved by  $\frac{\delta}{2}$  towards  $X_0$ , resulting in a new hyperplane  $\hat{\mathbf{w}}'$  that still fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} < 0$  for all  $\mathbf{x} \in X_0$ , but that now also fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X_1$ .

- By defining  $X' = X_1 \cup \{-\mathbf{x} \mid \mathbf{x} \in X_0\}$ , the searched  $\mathbf{w}$  fulfills  $\mathbf{w}^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X'$ . Then, with  $c(\mathbf{x}) = 1$  for all  $\mathbf{x} \in X'$ , *error*  $\in \{0, 1\}$  (instead of  $\{0, 1, -1\}$ ). [[PT Algorithm](#), Line 5]

# Perceptron Learning

## Perceptron Convergence Theorem: Proof

### Preliminaries:

- The sets  $X_1$  and  $X_0$  are separated by a hyperplane  $\hat{\mathbf{w}}$ . The proof requires that for all  $\mathbf{x} \in X_1$  the inequality  $\hat{\mathbf{w}}^T \mathbf{x} > 0$  holds. This condition is always fulfilled, as the following consideration shows.

Let  $\mathbf{x}' \in X_1$  with  $\hat{\mathbf{w}}^T \mathbf{x}' = 0$ . Since  $X_0$  is finite, the members  $\mathbf{x} \in X_0$  have a minimum positive distance  $\delta$  with regard to the hyperplane  $\hat{\mathbf{w}}$ . Hence,  $\hat{\mathbf{w}}$  can be moved by  $\frac{\delta}{2}$  towards  $X_0$ , resulting in a new hyperplane  $\hat{\mathbf{w}}'$  that still fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} < 0$  for all  $\mathbf{x} \in X_0$ , but that now also fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X_1$ .

- By defining  $X' = X_1 \cup \{-\mathbf{x} \mid \mathbf{x} \in X_0\}$ , the searched  $\mathbf{w}$  fulfills  $\mathbf{w}^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X'$ . Then, with  $c(\mathbf{x}) = 1$  for all  $\mathbf{x} \in X'$ , *error*  $\in \{0, 1\}$  (instead of  $\{0, 1, -1\}$ ). [[PT Algorithm](#), Line 5]
- The *PT* Algorithm performs a number of iterations, where  $\mathbf{w}(t)$  denotes the weight vector for iteration  $t$ , which form the basis for the weight vector  $\mathbf{w}(t+1)$ .  $\mathbf{x}(t) \in X'$  denotes the feature vector chosen in round  $t$ . The first (and randomly chosen) weight vector is denoted as  $\mathbf{w}(0)$ .

# Perceptron Learning

## Perceptron Convergence Theorem: Proof

### Preliminaries:

- The sets  $X_1$  and  $X_0$  are separated by a hyperplane  $\hat{\mathbf{w}}$ . The proof requires that for all  $\mathbf{x} \in X_1$  the inequality  $\hat{\mathbf{w}}^T \mathbf{x} > 0$  holds. This condition is always fulfilled, as the following consideration shows.

Let  $\mathbf{x}' \in X_1$  with  $\hat{\mathbf{w}}^T \mathbf{x}' = 0$ . Since  $X_0$  is finite, the members  $\mathbf{x} \in X_0$  have a minimum positive distance  $\delta$  with regard to the hyperplane  $\hat{\mathbf{w}}$ . Hence,  $\hat{\mathbf{w}}$  can be moved by  $\frac{\delta}{2}$  towards  $X_0$ , resulting in a new hyperplane  $\hat{\mathbf{w}}'$  that still fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} < 0$  for all  $\mathbf{x} \in X_0$ , but that now also fulfills  $(\hat{\mathbf{w}}')^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X_1$ .

- By defining  $X' = X_1 \cup \{-\mathbf{x} \mid \mathbf{x} \in X_0\}$ , the searched  $\mathbf{w}$  fulfills  $\mathbf{w}^T \mathbf{x} > 0$  for all  $\mathbf{x} \in X'$ . Then, with  $c(\mathbf{x}) = 1$  for all  $\mathbf{x} \in X'$ , *error*  $\in \{0, 1\}$  (instead of  $\{0, 1, -1\}$ ). [[PT Algorithm](#), Line 5]
- The *PT* Algorithm performs a number of iterations, where  $\mathbf{w}(t)$  denotes the weight vector for iteration  $t$ , which form the basis for the weight vector  $\mathbf{w}(t+1)$ .  $\mathbf{x}(t) \in X'$  denotes the feature vector chosen in round  $t$ . The first (and randomly chosen) weight vector is denoted as  $\mathbf{w}(0)$ .
- Recall the Cauchy-Schwarz inequality:  $\|\mathbf{a}\|^2 \cdot \|\mathbf{b}\|^2 \geq (\mathbf{a}^T \mathbf{b})^2$ , where  $\|\mathbf{x}\| := \sqrt{\mathbf{x}^T \mathbf{x}}$  denotes the Euclidean norm.

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

Line of argument:

- (a) We state a lower bound for how much  $\|\mathbf{w}\|$  must change from its initial value after  $n$  iterations (to become a separating hyperplane). The derivation of this lower bound exploits the presupposed linear separability of  $X_0$  and  $X_1$ .
- (b) We state an upper bound for how much  $\|\mathbf{w}\|$  can change from its initial value after  $n$  iterations. The derivation of this upper bound exploits the finiteness of  $X_0$  and  $X_1$ , which in turn guarantees the existence of an upper bound for the norm of the maximum feature vector.
- (c) We observe that the lower bound grows quadratically in  $n$ , whereas the upper bound grows linearly. From the relation “lower bound  $<$  upper bound” we derive a finite upper bound for  $n$ .

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

1. The PT Algorithm computes in iteration  $t$  the scalar product  $\mathbf{w}(t)^T \mathbf{x}(t)$ . If classified correctly,  $\mathbf{w}(t)^T \mathbf{x}(t) > 0$  and  $\mathbf{w}$  is unchanged. Otherwise,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$  [Line 5-7].

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

1. The [PT Algorithm](#) computes in iteration  $t$  the scalar product  $\mathbf{w}(t)^T \mathbf{x}(t)$ . If classified correctly,  $\mathbf{w}(t)^T \mathbf{x}(t) > 0$  and  $\mathbf{w}$  is unchanged. Otherwise,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$  [Line 5-7].
2. A sequence of  $n$  incorrectly classified feature vectors,  $(\mathbf{x}(t))$ , along with the weight adaptation,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ , results in the series  $\mathbf{w}(n)$  :  
$$\begin{aligned}\mathbf{w}(1) &= \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) \\ \mathbf{w}(2) &= \mathbf{w}(1) + \eta \cdot \mathbf{x}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \eta \cdot \mathbf{x}(1) \\ &\vdots \\ \mathbf{w}(n) &= \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \dots + \eta \cdot \mathbf{x}(n-1)\end{aligned}$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

1. The [PT Algorithm](#) computes in iteration  $t$  the scalar product  $\mathbf{w}(t)^T \mathbf{x}(t)$ . If classified correctly,  $\mathbf{w}(t)^T \mathbf{x}(t) > 0$  and  $\mathbf{w}$  is unchanged. Otherwise,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$  [Line 5-7].
2. A sequence of  $n$  incorrectly classified feature vectors,  $(\mathbf{x}(t))$ , along with the weight adaptation,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ , results in the series  $\mathbf{w}(n)$  :  
$$\begin{aligned}\mathbf{w}(1) &= \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) \\ \mathbf{w}(2) &= \mathbf{w}(1) + \eta \cdot \mathbf{x}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \eta \cdot \mathbf{x}(1) \\ &\vdots \\ \mathbf{w}(n) &= \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \dots + \eta \cdot \mathbf{x}(n-1)\end{aligned}$$
3. The hyperplane defined by  $\hat{\mathbf{w}}$  separates  $X_1$  and  $X_0$ :  $\forall \mathbf{x} \in X' : \hat{\mathbf{w}}^T \mathbf{x} > 0$   
Let  $\delta := \min_{\mathbf{x} \in X'} \hat{\mathbf{w}}^T \mathbf{x}$ . Observe that  $\delta > 0$  holds.



# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

1. The PT Algorithm computes in iteration  $t$  the scalar product  $\mathbf{w}(t)^T \mathbf{x}(t)$ . If classified correctly,  $\mathbf{w}(t)^T \mathbf{x}(t) > 0$  and  $\mathbf{w}$  is unchanged. Otherwise,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$  [Line 5-7].

2. A sequence of  $n$  incorrectly classified feature vectors,  $(\mathbf{x}(t))$ , along with the weight adaptation,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ , results in the series  $\mathbf{w}(n)$  :

$$\mathbf{w}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0)$$

$$\mathbf{w}(2) = \mathbf{w}(1) + \eta \cdot \mathbf{x}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \eta \cdot \mathbf{x}(1)$$

$$\vdots$$

$$\mathbf{w}(n) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \dots + \eta \cdot \mathbf{x}(n-1)$$

3. The hyperplane defined by  $\hat{\mathbf{w}}$  separates  $X_1$  and  $X_0$ :  $\forall \mathbf{x} \in X' : \hat{\mathbf{w}}^T \mathbf{x} > 0$

Let  $\delta := \min_{\mathbf{x} \in X'} \hat{\mathbf{w}}^T \mathbf{x}$ . Observe that  $\delta > 0$  holds.

4. Analyze the scalar product of  $\mathbf{w}(n)$  and  $\hat{\mathbf{w}}$  :

$$\hat{\mathbf{w}}^T \mathbf{w}(n) = \hat{\mathbf{w}}^T \mathbf{w}(0) + \eta \cdot \hat{\mathbf{w}}^T \mathbf{x}(0) + \dots + \eta \cdot \hat{\mathbf{w}}^T \mathbf{x}(n-1)$$

$$\Rightarrow \hat{\mathbf{w}}^T \mathbf{w}(n) \geq \hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta \geq 0 \quad (\text{for } n \geq n_0 \text{ with sufficiently large } n_0 \in \mathbb{N})$$

$$\Rightarrow (\hat{\mathbf{w}}^T \mathbf{w}(n))^2 \geq (\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

1. The PT Algorithm computes in iteration  $t$  the scalar product  $\mathbf{w}(t)^T \mathbf{x}(t)$ . If classified correctly,  $\mathbf{w}(t)^T \mathbf{x}(t) > 0$  and  $\mathbf{w}$  is unchanged. Otherwise,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$  [Line 5-7].

2. A sequence of  $n$  incorrectly classified feature vectors,  $(\mathbf{x}(t))$ , along with the weight adaptation,  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ , results in the series  $\mathbf{w}(n)$ :

$$\mathbf{w}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0)$$

$$\mathbf{w}(2) = \mathbf{w}(1) + \eta \cdot \mathbf{x}(1) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \eta \cdot \mathbf{x}(1)$$

$$\vdots$$

$$\mathbf{w}(n) = \mathbf{w}(0) + \eta \cdot \mathbf{x}(0) + \dots + \eta \cdot \mathbf{x}(n-1)$$

3. The hyperplane defined by  $\hat{\mathbf{w}}$  separates  $X_1$  and  $X_0$ :  $\forall \mathbf{x} \in X' : \hat{\mathbf{w}}^T \mathbf{x} > 0$

Let  $\delta := \min_{\mathbf{x} \in X'} \hat{\mathbf{w}}^T \mathbf{x}$ . Observe that  $\delta > 0$  holds.

4. Analyze the scalar product of  $\mathbf{w}(n)$  and  $\hat{\mathbf{w}}$ :

$$\hat{\mathbf{w}}^T \mathbf{w}(n) = \hat{\mathbf{w}}^T \mathbf{w}(0) + \eta \cdot \hat{\mathbf{w}}^T \mathbf{x}(0) + \dots + \eta \cdot \hat{\mathbf{w}}^T \mathbf{x}(n-1)$$

$$\Rightarrow \hat{\mathbf{w}}^T \mathbf{w}(n) \geq \hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta \geq 0 \quad (\text{for } n \geq n_0 \text{ with sufficiently large } n_0 \in \mathbb{N})$$

$$\Rightarrow (\hat{\mathbf{w}}^T \mathbf{w}(n))^2 \geq (\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2$$

5. Apply the Cauchy-Schwarz inequality:

$$\|\hat{\mathbf{w}}\|^2 \cdot \|\mathbf{w}(n)\|^2 \geq (\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2 \Rightarrow \|\mathbf{w}(n)\|^2 \geq \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2}$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

6. Consider again the weight adaptation  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ :

$$\begin{aligned} \|\mathbf{w}(t+1)\|^2 &= \|\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)\|^2 \\ &= (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t))^T (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)) \\ &= \mathbf{w}(t)^T \mathbf{w}(t) + \eta^2 \cdot \mathbf{x}(t)^T \mathbf{x}(t) + 2\eta \cdot \mathbf{w}(t)^T \mathbf{x}(t) \\ &\leq \|\mathbf{w}(t)\|^2 + \|\eta \cdot \mathbf{x}(t)\|^2 \quad (\text{since } \mathbf{w}(t)^T \mathbf{x}(t) < 0) \end{aligned}$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

6. Consider again the weight adaptation  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ :

$$\begin{aligned} \|\mathbf{w}(t+1)\|^2 &= \|\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)\|^2 \\ &= (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t))^T (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)) \\ &= \mathbf{w}(t)^T \mathbf{w}(t) + \eta^2 \cdot \mathbf{x}(t)^T \mathbf{x}(t) + 2\eta \cdot \mathbf{w}(t)^T \mathbf{x}(t) \\ &\leq \|\mathbf{w}(t)\|^2 + \|\eta \cdot \mathbf{x}(t)\|^2 \quad (\text{since } \mathbf{w}(t)^T \mathbf{x}(t) < 0) \end{aligned}$$

7. Consider the series  $\mathbf{w}(n)$  from Step 2:

$$\begin{aligned} \|\mathbf{w}(n)\|^2 &\leq \|\mathbf{w}(n-1)\|^2 + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &\leq \|\mathbf{w}(n-2)\|^2 + \|\eta \cdot \mathbf{x}(n-2)\|^2 + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &\leq \|\mathbf{w}(0)\|^2 + \|\eta \cdot \mathbf{x}(0)\|^2 + \dots + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &= \|\mathbf{w}(0)\|^2 + \sum_{j=0}^{n-1} \|\eta \cdot \mathbf{x}(j)\|^2 \end{aligned}$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

6. Consider again the weight adaptation  $\mathbf{w}(t+1) = \mathbf{w}(t) + \eta \cdot \mathbf{x}(t)$ :

$$\begin{aligned} \|\mathbf{w}(t+1)\|^2 &= \|\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)\|^2 \\ &= (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t))^T (\mathbf{w}(t) + \eta \cdot \mathbf{x}(t)) \\ &= \mathbf{w}(t)^T \mathbf{w}(t) + \eta^2 \cdot \mathbf{x}(t)^T \mathbf{x}(t) + 2\eta \cdot \mathbf{w}(t)^T \mathbf{x}(t) \\ &\leq \|\mathbf{w}(t)\|^2 + \|\eta \cdot \mathbf{x}(t)\|^2 \quad (\text{since } \mathbf{w}(t)^T \mathbf{x}(t) < 0) \end{aligned}$$

7. Consider the series  $\mathbf{w}(n)$  from Step 2:

$$\begin{aligned} \|\mathbf{w}(n)\|^2 &\leq \|\mathbf{w}(n-1)\|^2 + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &\leq \|\mathbf{w}(n-2)\|^2 + \|\eta \cdot \mathbf{x}(n-2)\|^2 + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &\leq \|\mathbf{w}(0)\|^2 + \|\eta \cdot \mathbf{x}(0)\|^2 + \dots + \|\eta \cdot \mathbf{x}(n-1)\|^2 \\ &= \|\mathbf{w}(0)\|^2 + \sum_{j=0}^{n-1} \|\eta \cdot \mathbf{x}(j)\|^2 \end{aligned}$$

8. With  $\varepsilon := \max_{\mathbf{x} \in X'} \|\mathbf{x}\|^2$  follows  $\|\mathbf{w}(n)\|^2 \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

9. Both inequalities (see Step 5 and Step 8) must be fulfilled:

$$\|\mathbf{w}(n)\|^2 \geq \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \quad \text{and} \quad \|\mathbf{w}(n)\|^2 \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\Rightarrow \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \leq \|\mathbf{w}(n)\|^2 \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\Rightarrow \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\begin{aligned} \text{Set } \mathbf{w}(0) = \mathbf{0} : \quad & \Rightarrow \frac{n^2\eta^2\delta^2}{\|\hat{\mathbf{w}}\|^2} \leq n\eta^2\varepsilon \\ & \Leftrightarrow n \leq \frac{\varepsilon}{\delta^2} \cdot \|\hat{\mathbf{w}}\|^2 \end{aligned}$$

# Perceptron Learning

## Perceptron Convergence Theorem: Proof (continued)

9. Both inequalities (see Step 5 and Step 8) must be fulfilled:

$$\|\mathbf{w}(n)\|^2 \geq \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \quad \text{and} \quad \|\mathbf{w}(n)\|^2 \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\Rightarrow \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \leq \|\mathbf{w}(n)\|^2 \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\Rightarrow \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \leq \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon$$

$$\begin{aligned} \text{Set } \mathbf{w}(0) = \mathbf{0} : \quad & \Rightarrow \frac{n^2\eta^2\delta^2}{\|\hat{\mathbf{w}}\|^2} \leq n\eta^2\varepsilon \\ & \Leftrightarrow n \leq \frac{\varepsilon}{\delta^2} \cdot \|\hat{\mathbf{w}}\|^2 \end{aligned}$$

→ The PT Algorithm terminates within a finite number of iterations.

$$\text{Observe: } \frac{(\hat{\mathbf{w}}^T \mathbf{w}(0) + n\eta\delta)^2}{\|\hat{\mathbf{w}}\|^2} \in \Theta(n^2) \quad \text{and} \quad \|\mathbf{w}(0)\|^2 + n\eta^2\varepsilon \in \Theta(n)$$

# Perceptron Learning

## Perceptron Convergence Theorem: Discussion [\[Theorem\]](#)

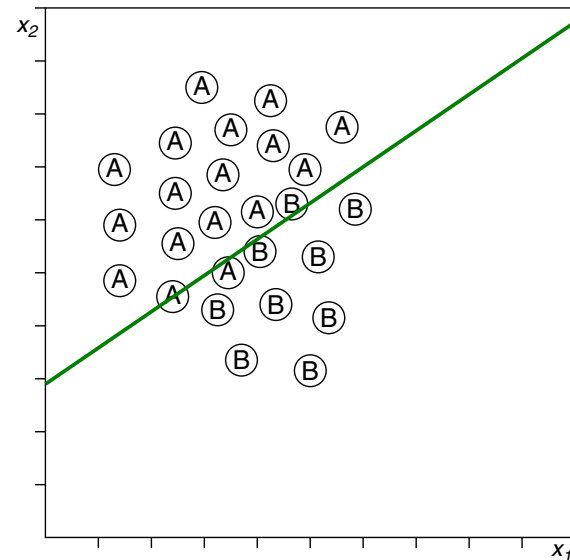
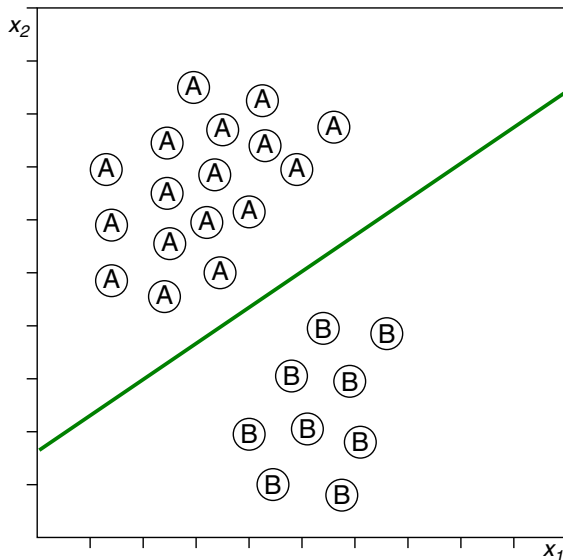
- ❑ If a separating hyperplane between  $X_0$  and  $X_1$  exists, the [PT Algorithm](#) will converge. If no such hyperplane exists, convergence cannot be guaranteed.
- ❑ A separating hyperplane can be found in polynomial time with linear programming. The *PT* Algorithm, however, may require an exponential number of iterations.



# Perceptron Learning

## Perceptron Convergence Theorem: Discussion [\[Theorem\]](#)

- If a separating hyperplane between  $X_0$  and  $X_1$  exists, the [PT Algorithm](#) will converge. If no such hyperplane exists, convergence cannot be guaranteed.
- A separating hyperplane can be found in polynomial time with linear programming. The *PT Algorithm*, however, may require an exponential number of iterations.
- Classification problems with noise (right-hand side) are problematic:



# Gradient Descent

## Classification Error

**Gradient descent** considers the true error (better: the hyperplane distance) and will converge even if  $X_1$  and  $X_0$  cannot be separated by a hyperplane. However, this convergence process is of an asymptotic nature and no finite iteration bound can be stated.

Gradient descent applies the so-called **delta rule**, which will be derived in the following. The delta rule forms the basis of the backpropagation algorithm.

# Gradient Descent

## Classification Error

**Gradient descent** considers the true error (better: the hyperplane distance) and will converge even if  $X_1$  and  $X_0$  cannot be separated by a hyperplane. However, this convergence process is of an asymptotic nature and no finite iteration bound can be stated.

Gradient descent applies the so-called **delta rule**, which will be derived in the following. The delta rule forms the basis of the backpropagation algorithm.

Consider the linear perceptron *without* a threshold function:

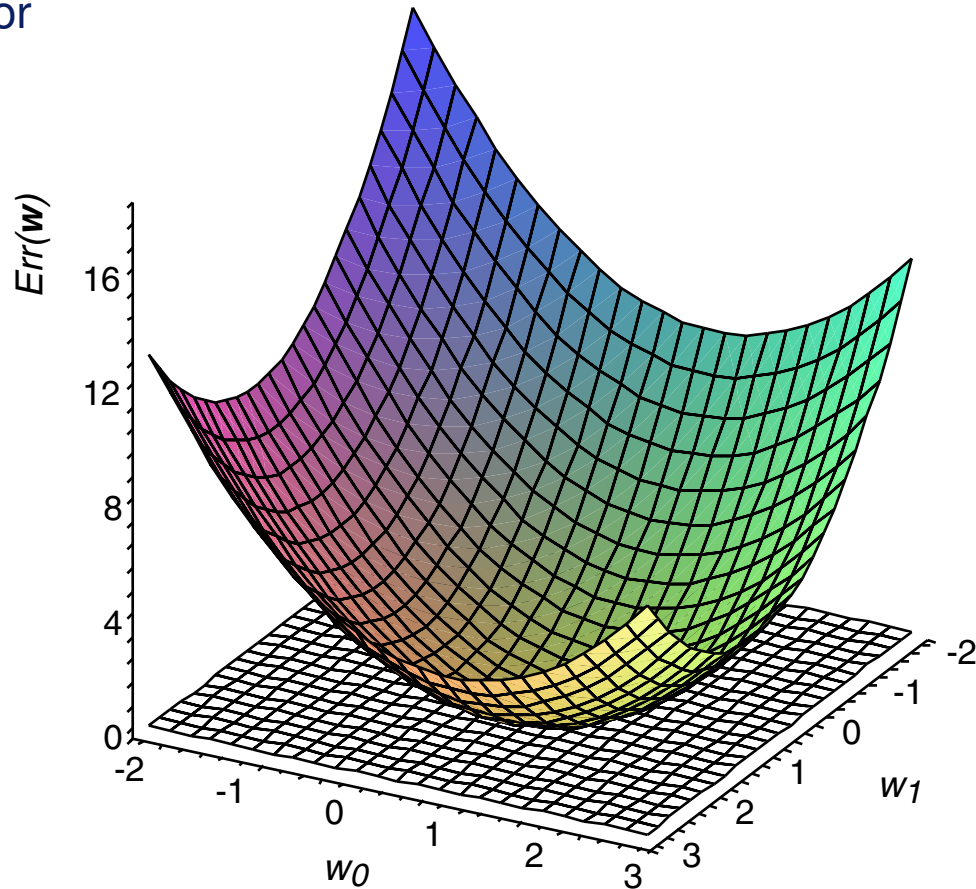
$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} = \sum_{j=0}^p w_j x_j \quad [\text{Heaviside}]$$

The classification error  $Err(\mathbf{w})$  of a weight vector (= hypothesis)  $\mathbf{w}$  with regard to  $D$  can be defined as follows:

$$Err(\mathbf{w}) = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 \quad [\text{Singleton error}]$$

# Gradient Descent

## Classification Error



The gradient of  $Err(\mathbf{w})$ ,  $\nabla Err(\mathbf{w})$ , defines the steepest ascent or descent:

$$\nabla Err(\mathbf{w}) = \left( \frac{\partial Err(\mathbf{w})}{\partial w_0}, \frac{\partial Err(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial Err(\mathbf{w})}{\partial w_p} \right)$$

# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w})$$

# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w})$$

$$\frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) = \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))^2$$

# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w})$$

$$\begin{aligned} \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))^2 \\ &= \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} 2(c(\mathbf{x}) - y(\mathbf{x})) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x})) \end{aligned}$$

# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w})$$

$$\begin{aligned} \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))^2 \\ &= \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} 2(c(\mathbf{x}) - y(\mathbf{x})) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x})) \\ &= \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \end{aligned}$$



# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w})$$

$$\begin{aligned} \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))^2 \\ &= \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} 2(c(\mathbf{x}) - y(\mathbf{x})) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x})) \\ &= \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \\ &= \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) (-x_j) \end{aligned}$$

# Gradient Descent

## Weight Adaptation

$$\mathbf{w} \leftarrow \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla \text{Err}(\mathbf{w}) \quad [\text{PT Algorithm}]$$

Componentwise ( $j = 0, \dots, p$ ) weight adaptation:

$$w_j \leftarrow w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) = \eta \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \cdot x_j$$

$$\frac{\partial}{\partial w_j} \text{Err}(\mathbf{w}) = \frac{\partial}{\partial w_j} \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - y(\mathbf{x}))^2 = \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))^2$$

$$= \frac{1}{2} \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} 2(c(\mathbf{x}) - y(\mathbf{x})) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - y(\mathbf{x}))$$

$$= \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \cdot \frac{\partial}{\partial w_j} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x})$$

$$= \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x})(-x_j)$$

# Gradient Descent

## Weight Adaptation: Batch Gradient Descent [Algorithms: IGD PT]

Algorithm: *BGD* Batch Gradient Descent

Input:  $D$  Training examples  $(\mathbf{x}, c(\mathbf{x}))$  with  $|\mathbf{x}| = p + 1$ ,  $c(\mathbf{x}) \in \{0, 1\}$ . ( $c(\mathbf{x}) \in \{-1, 1\}$ )  
 $\eta$  Learning rate, a small positive constant.

Internal:  $y(D)$  Set of  $y(\mathbf{x})$ -values computed from the elements  $\mathbf{x}$  in  $D$  given some  $\mathbf{w}$ .

Output:  $\mathbf{w}$  Weight vector.

$BGD(D, \eta)$

1. *initialize\_random\_weights*( $\mathbf{w}$ ),  $t = 0$
2. **REPEAT**
3.    $t = t + 1$
4.    $\Delta \mathbf{w} = 0$
5.   **FOREACH**  $(\mathbf{x}, c(\mathbf{x})) \in D$  **DO**
6.      $error = c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}$
7.      $\Delta \mathbf{w} = \Delta \mathbf{w} + \eta \cdot error \cdot \mathbf{x}$
8.   **ENDDO**
9.    $\mathbf{w} = \mathbf{w} + \Delta \mathbf{w}$
10. **UNTIL**(*convergence*( $D, y(D)$ )) **OR**  $t > t_{\max}$
11. *return*( $\mathbf{w}$ )

## Remarks:

- ❑  $\Delta \mathbf{w} \sim -\nabla \text{Err}(\mathbf{w})$ ; i.e., proportional to “−” and not to “+” to descend to the minimum.
- ❑ Each BGD iteration “REPEAT ... UNTIL” corresponds to finding the direction of steepest error descent as  $-\nabla \text{Err}(\mathbf{w}_t) = \sum_{(\mathbf{x}, c(\mathbf{x})) \in D} (c(\mathbf{x}) - \mathbf{w}_t^T \mathbf{x}) \cdot \mathbf{x}$  and updating  $\mathbf{w}_t$  by taking a step of length  $\eta$  in this direction.
- ❑ Using a constant step size  $\eta$  can severely impair the speed of convergence.  
When taking the optimal step size  $\eta_t := \operatorname{argmin}_{\eta} \text{Err}(\mathbf{w}_t - \eta \cdot \nabla \text{Err}(\mathbf{w}_t))$  at each iteration  $t$ , it can be shown that gradient descent has a linear rate of convergence, merely. [\[Meza 2010\]](#)
- ❑ The *convergence* function may compute the global error, either quantified as the sum of the squared residuals,  $\text{Err}(\mathbf{w}_t)$ , or as the norm of the error gradient,  $\|\nabla \text{Err}(\mathbf{w}_t)\|$ , and compare it to some small positive bound  $\varepsilon$ .

# Gradient Descent

## Weight Adaptation: Delta Rule

The weight adaptation in the BGD Algorithm is set-based: before modifying a weight component in  $\mathbf{w}$ , the total error of *all* examples (the “batch”) is computed.

Weight adaptation with regard to a *single* example  $(\mathbf{x}, c(\mathbf{x})) \in D$ :

$$\Delta \mathbf{w} = \eta \cdot (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}) \cdot \mathbf{x}$$

This adaptation rule is known under different names:

- ❑ delta rule
- ❑ Widrow-Hoff rule
- ❑ adaline rule
- ❑ least mean squares (LMS) rule

The classification error  $Err_d(\mathbf{w})$  of a weight vector (= hypothesis)  $\mathbf{w}$  with regard to a *single* example  $d \in D$ ,  $d = (\mathbf{x}, c(\mathbf{x}))$ , is given as:

$$Err_d(\mathbf{w}) = \frac{1}{2} (c(\mathbf{x}) - \mathbf{w}^T \mathbf{x})^2 \quad [\text{Batch error}]$$

# Gradient Descent

## Weight Adaptation: Incremental Gradient Descent [Algorithms: BGD PT LMS]

Algorithm: *IGD* Incremental Gradient Descent

Input:  $D$  Training examples  $(\mathbf{x}, c(\mathbf{x}))$  with  $|\mathbf{x}| = p + 1$ ,  $c(\mathbf{x}) \in \{0, 1\}$ . ( $c(\mathbf{x}) \in \{-1, 1\}$ )  
 $\eta$  Learning rate, a small positive constant.

Internal:  $y(D)$  Set of  $y(\mathbf{x})$ -values computed from the elements  $\mathbf{x}$  in  $D$  given some  $\mathbf{w}$ .

Output:  $\mathbf{w}$  Weight vector.

$IGD(D, \eta)$

```
1. initialize_random_weights( $\mathbf{w}$ ),  $t = 0$ 
2. REPEAT
3.    $t = t + 1$ 
4.   FOREACH  $(\mathbf{x}, c(\mathbf{x})) \in D$  DO
5.      $error = c(\mathbf{x}) - \mathbf{w}^T \mathbf{x}$ 
6.      $\Delta \mathbf{w} = \eta \cdot error \cdot \mathbf{x}$ 
7.      $\mathbf{w} = \mathbf{w} + \Delta \mathbf{w}$ 
8.   ENDDO
9. UNTIL ( $convergence(D, y(D))$  OR  $t > t_{max}$ )
10. return( $\mathbf{w}$ )
```

## Remarks:

- ❑ The sequence of incremental weight adaptations approximates the gradient descent of the batch approach. If  $\eta$  is chosen sufficiently small, this approximation can happen at arbitrary accuracy.
- ❑ The computation of the total error of batch gradient descent enables larger weight adaptation increments.
- ❑ Compared to batch gradient descent, the example-based weight adaptation of incremental gradient descent can better avoid getting stuck in a local minimum of the error function.
- ❑ Incremental gradient descent is also called *stochastic* gradient descent.

## Remarks (continued) :

- ❑ When, as is done here, the residual sum of squares, RSS, is chosen as error (loss) function, the incremental gradient descent algorithm [IGD] corresponds to the least mean squares algorithm [LMS].
- ❑ The incremental gradient descent algorithm [IGD] looks similar to the perceptron training algorithm [PT], since these algorithms differ only in the error computation (Line 5) where the latter applies the Heaviside function. However, this subtle syntactic difference is a significant conceptual difference, entailing a number of consequences:
  - Gradient descent is a regression approach and exploits the residuals, which are provided by an error function of choice, and whose differential is evaluated to control the hyperplane movement.
  - The PT algorithm is not based on residuals (in the  $(p + 1)$ -dimensional input-output-space) but refers to the input space only, where it simply evaluates the side of the hyperplane as a binary feature (correct side or not).
  - Provided linear separability, the PT algorithm will converge within a finite number of iterations, which, however, cannot be guaranteed for gradient descent.
  - Gradient descent will converge even if the data is not linearly separable.
  - Data sets can be constructed whose classes are linearly separable, but where gradient descent will not determine a hyperplane that classifies all examples correctly (whereas the PT Algorithm of course does).