

Chapter S:VI

VI. Relaxed Models

- ❑ Motivation
- ❑ ε -Admissible Speedup Versions of A^*
- ❑ Using Information about Uncertainty of h
- ❑ Risk Measures
- ❑ Nonadditive Evaluation Functions
- ❑ Heuristics Provided by Simplified Models
- ❑ Mechanical Generation of Admissible Heuristics
- ❑ Probability-Based Heuristics

Using Information about Uncertainty of h

Using a non-admissible heuristic function

Idea [Harris]:

The heuristic function h estimates the cheapest remaining cost h^* mostly quite well, but sometimes overestimates h^* by no more than ε .

→ A^* using such a heuristic function h is ε -admissible.

The condition for ε -admissibility of A^* is satisfied because at termination it holds

$$h(n) - h^*(n) \leq \varepsilon \quad \text{für alle } n \in \text{OPEN}.$$

Also the weakened form of admissibility of h is often too restrictive.

Often it is easier to find a heuristic estimate for h^* that mostly estimates precisely but sometimes overestimates h^* (by much more than any reasonable ε).




→ The error in the estimate is not limited, but a large error is unlikely.

Remarks:

- ❑ Heuristic functions h with $h \leq (1 + \varepsilon)h^*$ are called ε -admissible.
- ❑ Analogously to Lemma C^* -Bounded OPEN Node, it can be proven that, at any point in time before termination, there exists some node n in OPEN with $f(n) \leq (1 + \varepsilon)C^*$.
- ❑ The condition " $h(n) - h^*(n) \leq \varepsilon$ for all $n \in \text{OPEN}$ " is sufficient, but not necessary, for A* being ε -admissible.

Using Information about Uncertainty of h

Illustration of Underestimating and Overestimating Estimation Functions



Using Information about Uncertainty of h

Example: Search in “Random” Graphs

Given is a graph with randomly drawn edge costs. The minimum number of edges to a target node is known in each node.

- ❑ Edge costs $c(n, n')$ are known to be drawn independently from a common distribution function, uniform in interval $[0; 1]$.
- ❑ For long paths with N edges from a node n to a goal node in Γ it is known that $h^*(n)$ is most likely to be near $\frac{N}{2}$.
- ❑ The only *admissible* heuristic estimate for h^* is $h_1(n) = 0$.
- ❑ The most reasonable heuristic estimate for h^* is $h_2(n) = \frac{N}{2}$.

The heuristic estimate h_2 leads to a worst-case cost overestimation of $\frac{N}{2}$ and is therefore not (ε -)admissible. But the likelihood of this event is extremely small.

Using Information about Uncertainty of h

Example: Search in “Random” Graphs

Given is a graph with randomly drawn edge costs. The minimum number of edges to a target node is known in each node.

- ❑ Edge costs $c(n, n')$ are known to be drawn independently from a common distribution function, uniform in interval $[0; 1]$.
- ❑ For long paths with N edges from a node n to a goal node in Γ it is known that $h^*(n)$ is most likely to be near $\frac{N}{2}$.
- ❑ The only *admissible* heuristic estimate for h^* is $h_1(n) = 0$.
- ❑ The most reasonable heuristic estimate for h^* is $h_2(n) = \frac{N}{2}$.

The heuristic estimate h_2 leads to a worst-case cost overestimation of $\frac{N}{2}$ and is therefore not (ε -)admissible. But the likelihood of this event is extremely small.

→ Algorithm R^*_δ :

- Besides an estimation function h for h^* there is also knowledge about the uncertainty of the estimation process.
- Knowledge about the uncertainty of the estimation process is expressed in the form of a probability density function $\rho_{h^*}(x)$.

Using Information about Uncertainty of h

Describing the Estimation Uncertainty using Density Functions

Viewing cost functions as a random variables:

cost function	random variable
$h^*(n)$	h_n^*
$f^*(n) = g^*(n) + h^*(n)$	f_n^*
$f^+(n) = g(n) + h^*(n)$	f_n^+

Let $\rho_{h_n^*}$ be a density function for the random variable h_n^* .

Semantics:

On the basis of $\rho_{h_n^*}$ one can define the probability with which $h^*(n)$ can be found in a neighborhood of x costs.

$$P(h_n^* = x) = \rho_{h_n^*}(x)$$

Using Information about Uncertainty of h

Describing the Estimation Uncertainty using Density Functions (continued)

Let $\rho_{h_n^*}$ be a density function for the random variable h_n^* .

Further applies:

1. From $\rho_{h_n^*}(x)$ a density function $\rho_{f_n^*}(y)$ can be derived for the random variable f_n^* , if g^* is known (e.g., when searching a tree):

$$\rho_{f_n^*}(y) := \rho_{h_n^*}(y - g^*)$$

2. Let P_{s-n} be the cheapest known path from s to an OPEN node n .

From $\rho_{h_n^*}(x)$ a density function $\rho_{f_n^+}(y)$ can be derived for the random variable f_n^+ , which specifies the cost of an optimal solution path that continues P_{s-n} :

$$\rho_{f_n^+}(y) := \rho_{h_n^*}(y - g)$$


Remarks:

- ❑ The random variable f_n^+ with associated density function $\rho_{f_n^+}$ is given for each node n .
- ❑ The random variable f_n^+ describes the possible costs of an optimal solution path that contains the pointer path P_{s-n} as a subpath.
- ❑ If goal nodes can be reached from s , the OPEN list always contains a node n , to which $f^+(n) = f^*(n)$ applies. [Corollary Shallowest OPEN Node on Optimum Path]



Using Information about Uncertainty of h

Describing the Estimation Uncertainty using Density Functions (continued)


Uncertainty area:



Density functions ρ :



Related distribution function:

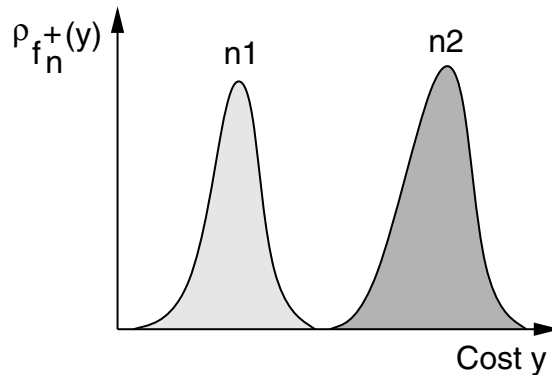


Using Information about Uncertainty of h


Describing the Estimation Uncertainty using Density Functions (continued)

How should an evaluation order be calculated from the density functions $\rho_{f_n^+}$ for the nodes in the OPEN list?

Possible shapes of two density functions:



Case (a)



Case (b)

Remarks:

- (a) If the density functions do not overlap, the node for which the corresponding density function ρ_{f^+} has the lowest density value f_a^+ with respect to all other nodes would be selected.
 - (b) $f_{n_1}^+$ has the lower expected value; n_2 *has the possibility* that the cost $f_{n_2}^+$ may be lower than n_1 .
An admissible algorithm would expand n_2 .
It would make more sense to expand n_1 because the “ $f^+(n_2) < f^+(n_1)$ ” event is unlikely.
- Due to uncertainty, costs can be overestimated or underestimated. I.e., not expanding a node in OPEN and terminating it too expensively as a result, represents a *risk*.
- Quantification of the risk of terminating with too high costs (= terminating too early).

Risk Measures

Defining the Order of Node Evaluations

Idea:

Estimate the risk of terminating too early using a *risk measure* R for each node in the OPEN list.

- ❑ For a given cost value C (of a goal node), the risk measure evaluates for each node n in the OPEN list to what extent C can be improved by expanding n .
- ❑ $R = R(C)$. The risk measure is a nondecreasing function of the C cost. The greater the $R(C)$ value of a node n , the greater the risk of missing an improvement of C if terminating with C without expanding n .
- ❑ $R(C)$ should use knowledge about the cost distribution for the node n , so it should be based on $\rho_{f_n^+}$.

Risk Measures

Principle of the Algorithm R^*_δ

Search continues until the risk value $R(C)$ of each node in the OPEN list is below a user accepted risk threshold δ .

- If a high risk is acceptable, nodes with a high risk value of $R(C)$ (= high cost reduction potential) remain unexpanded. As a result, cost underestimation becomes less likely.

If only a small risk is acceptable, even nodes with a low risk value of $R(C)$ (= low cost reduction potential) are expanded. As a result, cost underestimation becomes more likely.

- I.e., depending on a risk threshold δ the probability of a cost underestimation (= probability of admissibility or optimality) can be controlled.


Remarks:

- ❑ The observance of this principle by the algorithm R^*_δ is ensured — as shown later — by the use of certain risk functions $R(C)$.
- ❑ When a goal node with cost C is selected from OPEN, its risk function must guarantee the property $C \leq C_\delta$. Otherwise, one of the remaining nodes in OPEN can have a risk for C that is higher than δ .


Risk Measures

Potential for Improvement to a Current Solution

Let n_1, n_2 be nodes of the OPEN list.



Example of risk functions $R(C)$ for the nodes n_1, n_2 :



The nodes have different random variables $f_{n_1}^+$ and $f_{n_2}^+$ for the cost.


Remarks:

- ❑ The potential for improvement is a statistical quantity defined for a node n using f_n^+ .
- ❑ The *evaluation* of the potential for improvement regarding given costs C is done with the help of a risk-measure $R(C)$.

Risk Measures

Risk Threshold and Cost Threshold

The risk threshold $\delta \geq 0$ defines for each node n in OPEN its cost threshold $C_\delta(n)$:



Let n_1, n_2 be nodes in OPEN. If the search was terminated with node n_2 and cost $C' = C_\delta(n_2)$, the risk $R(C')$ for n_1 would be above the risk threshold δ .

- R_δ^* chooses the node n with the lowest cost threshold $C_\delta(n)$ in the OPEN list. In the above example, node n_1 would be preferred to node n_2 .

Risk Measures

Definition 75 (Risk Measure)

Let M be the ordered set of cost values. A risk measure $R(C)$ for a node is a nondecreasing function $R : M \rightarrow [0, \infty]$ measuring the risk associated with leaving that node unexplored when terminating with a solution with cost C .

Definition 76 (Cost Threshold)

Let δ be a nonnegative real number and let $R(C)$ be the risk measure for a node n . The solution $C_\delta(n)$ to the equation $R(C) = \delta$ is called the cost threshold.

Assuming the cost of a solution path found is C , then for each node n in OPEN with $C > C_\delta(n)$ the risk of missing a better solution path is higher than risk threshold δ . These nodes should be expanded before termination.

Remarks:

- ❑ Risk measures and risk thresholds must be seen in context: not every risk threshold makes sense for a risk measure.
- ❑ Depending on the f_n^+ cost random variable of a n node, the δ risk threshold can lead to different sequences in the OPEN list.
- ❑ The cost-threshold $C_\delta(n)$ indicates how high the cost of a solution may be without exceeding the δ -risk-threshold for the node n .

Risk Measures

Definition 77 (δ -Risk-Admissibility)

An algorithm is said to be δ -risk-admissible if it always terminates with a solution cost C such that $R(C) \leq \delta$ for each node left on OPEN.

The above version of the δ -risk-admissible condition is equivalent to stating that at termination, the cost of the solution found is not greater than $C_\delta(n)$ for each n on OPEN.

Definition 78 (Algorithm R^*_δ)

R^*_δ is a search algorithm which is identical to A^* except that it chooses for expansion that node n from OPEN with the lowest cost-threshold $C_\delta(n)$.

Note that with $\delta = 0$, R^*_δ is identical to A^* since it is guided by the (admissible) lowest tail of the density of f , namely by $g + h_a$.

For $\delta > 0$, R^*_δ may prefer a node with high f_a and narrow distribution over a node with low f_a but highly diffused density.

Remarks:



- ❑ The first definition of δ -risk-admissibility is from the perspective of risk, the second is from the perspective of cost.
- ❑ With $\delta = 0$, R^*_δ is identical to A^* . Justification:
 1. Computing the cost-threshold $C_\delta(n)$ for nodes in OPEN is solving the equation $R(C) = 0$ for $\delta = 0$.
 2. $R(C) = 0$ holds for the lowest point f_a on the tail of the density of f^+ .
 3. Hence, $f_a = g + h_a \leq g(n) + h^*(n)$
 4. R^*_δ is guided by an admissible heuristic function and, therefore, R^*_δ is admissible.
- ❑ As the δ increases, R^*_δ tends to abandon admissibility.

Risk Measures

Risk Measures of Type $R(C) = \varrho[C - f^+]$

Starting point are density functions for the random variables f_n^+ of nodes n in the OPEN list.

Examples:



f_a (resp. h_a) is the smallest positive preimage of the density function $\rho_{f_n^+}$ (resp. $\rho_{h_n^*}$).

Risk Measures

Risk Measures of Type $R(C) = \varrho[C - f^+]$ (continued)

1. Worst Case Risk R_1 :

$$R_1(C) = \sup_{\{y | \rho_{f_n^+}(y) > 0\}} (C - y) = C - f_a = C - g - h_a$$

2. Probability of Suboptimal Termination R_2 :

$$R_2(C) = P(C > f_n^+) = P(C - f_n^+ > 0) = \int_{y=-\infty}^C \rho_{f_n^+}(y) dy$$

3. Expected Risk R_3 :

$$R_3(C) = E(\max\{C - f_n^+; 0\}) = \int_{y=-\infty}^C (C - y) \rho_{f_n^+}(y) dy$$

Remarks:

- The risk measures R_1 and R_3 describe costs, the risk measure R_2 describes a probability.

- R_1 :

For the costs represented by the f_n^+ random variable, the smallest possible value is assumed. R_1 quantifies the maximum possible loss if a solution is satisfied with C costs.

The lowest costs are the worst case because they represent the extreme case of a missed cost reduction. The probability that the remaining costs are lower than h_a is 0.

- R_2 :

The probability for the occurrence of the event " $C > f_n^+$ " (i.e., event is a loss) is calculated if you are satisfied with a solution with C cost.

- R_3 :


For the costs represented by the random variable f_n^+ , the expected loss $E(\max\{C - f_n^+; 0\})$ is calculated if one is satisfied with a solution with costs C .

R_3 weights the probability of the loss (R_2) with the amount of the occurring loss.

Risk Measures

Example

Let f_n^+ be uniformly distributed between an optimistic estimate f_a and a pessimistic estimate f_b (The f_a and f_b estimates depend on n , where n is in OPEN.):



Density function:

$$\rho_{f_n^+}(y) = \begin{cases} \frac{1}{f_b - f_a} & f_a \leq y \leq f_b \\ 0 & \text{else} \end{cases}$$


Risk measure:

$$R_1(C) = \sup_{\{y | \rho_{f_n^+}(y) > 0\}} (C - y) = C - f_a$$

Risk Measures

Example (continued)

Let f_n^+ be uniformly distributed between an optimistic estimate f_a and a pessimistic estimate f_b (The f_a and f_b estimates depend on n , where n is in OPEN.):



Density function:

$$\rho_{f_n^+}(y) = \begin{cases} \frac{1}{f_b - f_a} & f_a \leq y \leq f_b \\ 0 & \text{else} \end{cases}$$


Risk measure:

$$R_2(C) = \int_{y=-\infty}^C \rho_{f_n^+}(y) dy = \begin{cases} 0 & C < f_a \\ \frac{(C - f_a)}{(f_b - f_a)} & f_a \leq C \leq f_b \\ 1 & f_b < C \end{cases}$$

Risk Measures

Example (continued)

Let f_n^+ be uniformly distributed between an optimistic estimate f_a and a pessimistic estimate f_b (The f_a and f_b estimates depend on n , where n is in OPEN.):



Density function:

$$\rho_{f_n^+}(y) = \begin{cases} \frac{1}{f_b - f_a} & f_a \leq y \leq f_b \\ 0 & \text{else} \end{cases}$$


Risk measure:

$$R_3(C) = \int_{y=-\infty}^C (C - y) \rho_{f_n^+}(y) dy = \begin{cases} 0 & C < f_a \\ \frac{(C - f_a)^2}{2(f_b - f_a)} & f_a \leq C \leq f_b \\ C - \frac{f_a + f_b}{2} & f_b < C \end{cases}$$

Risk Measures

Example (continued)

Shape of risk measures R_1 , R_2 , and R_3 (f_n^+ uniformly distributed):




The vertical axis represents the functions $R(C)$ for the three risk measures considered.

Risk Measures

Example (continued)

Computing the cost threshold C_δ for R_3 (f_n^+ uniformly distributed):

Let δ be the user's risk tolerance of the user. For each node n in OPEN, it defines its cost threshold $C_\delta(n)$ using the equation $R(C) = \delta$.



For a node n on OPEN, $C_\delta(n)$ is computed by transforming $R_3(C_\delta) = \delta$:


$$C_\delta(n) = \begin{cases} f_a & \delta = 0 \\ f_a + \sqrt{2 \cdot (f_b - f_a) \cdot \delta} & 0 < \delta \leq \frac{f_b - f_a}{2} \\ \delta + \frac{f_a + f_b}{2} & \frac{f_b - f_a}{2} < \delta \end{cases}$$

Risk Measures

Example (continued)

For a goal node γ in OPEN, we have $h(n) = h^*(n) = 0$. Therefore, there remains no uncertainty regarding f_γ^+ .

Graphs of random variable f_γ^+ for solution cost and risk measures R_1 , R_2 , and R_3 :



Cost threshold:

$$C_\delta(\gamma) = \begin{cases} g(\gamma) + \delta & \text{for risk measure } R_1 \\ g(\gamma) & \text{for risk measure } R_2 \text{ and } \delta < 1 \\ g(\gamma) + \delta & \text{for risk measure } R_3 \end{cases}$$

Risk Measures

Theorem 79 (δ -Risk-Admissibility of R^*_δ)

R^*_δ is δ -risk-admissible with respect to risk measures R_1 , R_2 , and R_3 when G is a search space graph with $Prop(G)$ and $E(h^*) < \infty$ on solution paths.

Risk Measures

Theorem 79 (δ -Risk-Admissibility of R^*_δ)

R^*_δ is δ -risk-admissible with respect to risk measures R_1 , R_2 , and R_3 when G is a search space graph with $Prop(G)$ and $E(h^*) < \infty$ on solution paths.

Proof (sketch)

1. δ -Risk-Admissibility:

- (a) According to the previous example, it holds for the cost C of a solution path found by R^*_δ :

$$C = g(\gamma) \leq C_\delta(\gamma) \quad \text{for the risk measures } R_1, R_2, R_3.$$

- (b) Since R^*_δ chooses for expansion that node n from OPEN with the lowest cost-threshold $C_\delta(n)$, δ -risk-admissibility of R^*_δ follows for risk measures R_1 , R_2 , and R_3 .

2. Completeness:

- (a) At all times OPEN contains a node n on a solution path for which $C_\delta(n)$ is finite. Obviously, $R_1(C) = \delta$ and $R_2(C) = \delta$ have a finite solution. If density $\rho_{h^*}(x)$ possesses a finite expectation $E(h^*) < \infty$ for any node on a solution path, for R_3 we have

$$R_3(C) \geq C \cdot (1 - P(f^+ > C)) - E(f^+) \geq C - 2E(f^+) = C - 2g - 2E(h^*)$$

- (b) $C_\delta(n) \geq g(n)$ holds for each node n in OPEN since there is no risk in abandoning n after finding a solution path with cost $\leq g(n)$. A positive lower bound of the edge cost values guarantees that R^*_δ can neglect nodes on solution paths only for a limited number of node expansions.

Remarks:

- ❑ Expectations can have the value ∞ , e.g., for a random variable that returns values 2^n with probability 2^{-n} .
- ❑ In step 2(a) we use the fact $R_3(0) = 0$ for graphs G with nonnegative edge cost values. As the lower bound for $R_3(C)$ increases with C , there is a finite value C with $R_3(C) > \delta$. Hence, $C_\delta(n) < \infty$.
- ❑ The exact form of $\rho_{h_n^*}$ is generally unknown. For this the edge costs must have been generated by a given probabilistic model.
- ❑ Generating a good estimate for $C_\delta(n)$ is often possible. For this, the knowledge of upper and lower bounds of h_n^* together with the often reasonable assumption of a standardized distribution between them, such as an uniform distribution, an exponential distribution or a normal distribution, is sufficient.
- ❑ The principle of the ε -admissible acceleration in A_ϵ^* for A^* can also be applied to R_δ^* and leads to the algorithm $R_{\delta,\epsilon}^*$. The special version $R_{\delta,\delta}^*$ is δ -risk-admissible with respect to risk measures R_1 , R_2 , and R_3 under the preconditions of the previous theorem.