

The Archive Query Log: Mining Millions of Search Result Pages of Hundreds of Search Engines from 25 Years of Web Archives

Jan Heinrich Reimer
Friedrich-Schiller-Universität Jena

Sebastian Schmidt
Leipzig University

Maik Fröbe
Friedrich-Schiller-Universität Jena

Lukas Gienapp
Leipzig University and ScaDS.AI

Harris Scells
Leipzig University

Benno Stein
Bauhaus-Universität Weimar

Matthias Hagen
Friedrich-Schiller-Universität Jena

Martin Potthast
Leipzig University and ScaDS.AI

ABSTRACT

The Archive Query Log (AQL) is a previously unused, comprehensive query log collected at the Internet Archive over the last 25 years. Its first version includes 356 million queries, 137 million search result pages, and 1.4 billion search results across 550 search providers. Although many query logs have been studied in the literature, the search providers that own them generally do not publish their logs to protect user privacy and vital business data. Of the few query logs publicly available, none combines size, scope, and diversity. The AQL is the first to do so, enabling research on new retrieval models and (diachronic) search engine analyses. Provided in a privacy-preserving manner, it promotes open research as well as more transparency and accountability in the search industry.

CCS CONCEPTS

- Information systems → Query log analysis.

KEYWORDS

query log, search engine result page, information retrieval history

ACM Reference Format:

Jan Heinrich Reimer, Sebastian Schmidt, Maik Fröbe, Lukas Gienapp, Harris Scells, Benno Stein, Matthias Hagen, and Martin Potthast. 2023. The Archive Query Log: Mining Millions of Search Result Pages of Hundreds of Search Engines from 25 Years of Web Archives. In *Proceedings of the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '23), July 23–27, 2023, Taipei, Taiwan*. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3539618.3591890>

1 INTRODUCTION

Search engine query logs are a rich resource for many information retrieval applications [4], such as user behavior and user experience analysis, or query suggestions and query reformulations. When a query log also includes users' clicks and dwell time on search results, this is a valuable source of implicit relevance feedback

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGIR '23, July 23–27, 2023, Taipei, Taiwan

© 2023 Copyright held by the owner/author(s). Publication rights licensed to ACM. ACM ISBN 978-1-4503-9408-6/23/07...\$15.00
<https://doi.org/10.1145/3539618.3591890>

Table 1: The Archive Query Log 2022 (AQL-22) at a glance.

Search provider (known domains)	URLs (total)	Queries (total)	Queries (unique)	SERPs (estimate)	Results (estimate)
Google	89.4 M	72.7 M	20.0 M	28.0 M	223.1 M
YouTube	41.8 M	41.4 M	11.3 M	15.9 M	339.2 M
Baidu	78.5 M	69.6 M	2.9 M	26.8 M	107.6 M
QQ	0.5 M	0.5 M	0.1 M	0.2 M	2.1 M
Facebook	3.1 M	0.2 M	0.0 M	0.1 M	0.7 M
Yahoo!	8.8 M	2.8 M	1.2 M	1.1 M	9.2 M
Amazon	66.8 M	0.8 M	0.3 M	0.3 M	7.8 M
Wikipedia	68.5 M	1.7 M	0.6 M	0.7 M	7.0 M
JD.com	4.4 M	3.9 M	0.4 M	1.5 M	16.0 M
360	1.5 M	1.1 M	0.1 M	0.4 M	3.5 M
:	540 others	646.8 M	161.8 M	27.8 M	62.4 M
Σ	550	1010.2 M	356.5 M	64.5 M	137.3 M
					1410.0 M

about their information needs. Modern search engines use this feedback to train retrieval models for re-ranking [93, 117]. However, query logs are also highly sensitive data that affect a number of stakeholders [11, 73]: First and foremost are user privacy concerns. Over time, if a user frequently uses a search engine, their query log can be enough to personally identify them and reveal a lot about their state of mind and health. To some extent, this also applies to persons or organizations mentioned or implied in queries or search results. Not least, relevance feedback from a query log is an important asset for search providers, commercial or otherwise.

The aforementioned arguments present valid concerns for refraining from publishing query logs. An additional, albeit debatable, rationale for major search providers stems from the expectations of increased transparency and accountability from governments, civil societies, affected users, and third parties, owing to their substantial market presence. Granting access to query logs would facilitate large-scale, independent investigations into the accuracy and fairness of their search results [11], foster competition [12], support law enforcement efforts [35], and advance public information retrieval research. However, these objectives may not align with the best interests of search providers themselves.

We have uncovered and acquired an extensive query log that has accumulated at the Internet Archive over the last 25 years. We call this new resource Archive Query Log (AQL). Table 1 gives an overview of the first version of 2022 and the top ten search providers as per fused snapshots of Alexa website traffic rankings. Shown are the respective numbers of archived URLs, the queries extracted

Table 2: Overview of large-scale query logs used in previous work. Private logs are grouped by source. Each source is referenced by the paper using the largest sample. Number of usages is given under Nº. Timespan indicates crawled duration; year indicates date of last included query. Fields with ‘-’ are either not available or not specified. Languages are estimated. The ★ marks logs still available for download; ☀ marks industry; ☁ academic; ✕ mixed affiliations.

Source (name)	Clickable link: ⚡	Queries	Queries (uniq.)	Sessions	Users	Clicks	Results	Task	Lang.	Span	Year	Aff.	Ref.	Nº
Private (sample)	AltaVista	575,244,993	153,645,050	-	285,000,000	-	-	Web	en	1m	1998	☀	[102]	3
	Dreamer	2,184,256	228,566	-	-	-	-	Web	zh	3m	1998	☀	[26]	1
	Excite	51,473	18,098	-	18,113	-	-	Web	en	-	1998	☀	[63]	3
	Infoseek	19,933,187	-	-	-	-	-	Web	zh	1m	1998	✖	[46]	1
	GAIS	475,564	114,182	-	-	-	-	Web	zh	2w	1999	☀	[26]	1
	Lycos	500,000	243,595	-	-	500,000	361,906	Web	en	1d	2000	✖	[17]	1
	Yahoo!	2,369,282	-	-	-	21,421	-	Web	en,zh	4m	2000	☁	[59]	9
	OpenFind	2,493,211	-	-	-	-	-	Web	en	1y	2000	☁	[26]	1
	Encarta	2,772,615	-	2,772,615	-	-	-	Web	en	1m	2002	✖	[109]	1
	Utah State Gov.	792,103	575,389	458,962	161,042	323,285	-	Web	en	5m	2003	☁	[24]	1
	Timway	1,255,633	-	-	-	-	-	Web	zh	3m	2004	☁	[81]	1
	MetaSpy	580,000	-	-	-	-	-	Web	en	5d	2005	☁	[6]	1
	arXiv	44,399	-	-	13,304	48,976	-	Edu	en	3m	2006	☁	[99]	1
	TodoCL	192,924	-	348,035	-	360,641	360,641	Web	en	9m	2006	✖	[48]	1
	kunstmuseum.nl	7,531	1,183	-	-	-	-	Lib.	nl	2y	2007	☁	[9]	1
	Microsoft AdCenter	27,922,224	27,922,224	-	-	7,820,000	-	Ads	en	2m	2007	☁	[74]	2
	Europeana	3,024,162	1,382,069	-	-	-	-	Lib.	en	6m	2011	✖	[22]	1
	GNU IFT	2,099	-	-	-	-	4,754	Img.	en	1y	2011	☁	[86]	1
	INDURE	14,503	2,923	85,215	6,434	-	-	Edu	en	3m	2011	☁	[49]	1
	Bing Videos	1,218,936	445,859	174,955	174,955	-	-	Vid.	en	1w	2011	✖	[72]	1
	Baidu	362,994,092	10,413,491	87,744,130	-	-	13,126,252	Web	zh	-	2012	☀	[115]	4
	CADAL Library	45,892	-	81,759	-	-	164,822	Lib.	en	-	2012	☁	[112]	1
	Taobao	1,410,960	-	-	4,285	-	-	Prod.	zh	1m	2013	☁	[116]	2
	Startpagina	10,000,000	-	-	-	-	-	Web	nl	1m	2014	✖	[55]	1
	parsijoo.ir	27,000,000	-	-	-	-	-	Web	fa	2y	2017	☁	[85]	1
	CiteSeerX	78,124,884	14,759,852	-	-	-	-	Edu	en	4y	2021	☁	[100]	1
: 15 query logs from undisclosed sources, all private, 5 ☀, 2 ☁, 8 ✕														-
[14, 15, 18, 20, 54, 58–60, 66, 67, 79, 80, 98, 106, 111]														-
Public (exhaustive)	PubMed	2,996,301	-	-	627,455	-	-	Med.	en	1d	2005	☁	[57]	3
	AOL (AOL Query Log '06)	36,389,567	10,154,742	-	657,426	19,442,629	19,442,629	Web	en	3m	2006	☀	[95]	14
	MSN (MS RFP'06)	14,921,285	-	14,921,285	-	-	-	Web	en	1m	2006	☀	[17]	-
	Belga News Agency	1,402,990	-	-	-	5,697,287	498,039	Img.	en	1y	2008	☁	[87]	2
	bildungsserver.de (DBMS)	98,512	31,347	65,513	-	68,604	-	Edu	de	-	2009	☁	[84]	4
	Gov2 Crawl (LETOR 4.0)	2,500	-	-	-	-	-	Web	en	-	2009	☀	[97]	2
	Sogou	18,393,652	4,580,463	-	8,168,051	-	14,075,717	Web	zh	1m	2009	☁	[114]	4
	Tumba!	458,623	-	-	-	-	-	Web	en	-	2009	☁	[83]	1
	European Library (TEL)	162,642	-	75,100	-	-	-	Lib.	en	1y	2010	☁	[83]	4
	Yandex	10,139,547	-	-	956,536	-	49,029,185	Web	ru	-	2011	☀	[2]	-
	Bing (MS Image Ret. Ch.)	11,701,890	-	-	-	-	-	Img.	en	-	2013	☀	[1]	-
	Bing (ORCAS)	18,823,602	-	-	-	18,823,602	18,823,602	Web	en	-	2020	✖	[38]	3
	AOL (AOLIA)	11,337,160	-	-	-	-	1,525,524	Web	en	-	2022	☀	[82]	1
	StackOverflow	9,046,179	-	16,164,506	-	-	-	Q&A	en	-	2022	☀	[2]	-
	Archive Query Log (AQL)	356,450,494	64,544,345	-	-	-	1,409,974,669	Multi	Multi	25y	2022	☁	-	-

from them, archived search result pages (SERPs), and results linked to them. The SERPs of many queries have been archived multiple times, enabling diachronic analyses. At the time of writing, we collect this data for a total of 550 search providers. The Archive Query Log 2022 includes 356 million queries (65 million unique), 137 million search result pages, and 1.4 billion search results—an unprecedented scale for a public query log. Based on a comprehensive review of public and private query logs used in the literature (Section 2), we detail our acquisition method (Section 3), initial analyses (Section 4), and discuss our plan to share the data with the information retrieval community in a privacy-preserving manner, as well as limitations and ethical considerations (Section 5).^{1,2}

¹AQL code: <https://github.com/webis-de/SIGIR-23>

²AQL access: <https://tira.io/task/archive-query-log>

2 BACKGROUND AND RELATED WORK

We systematically review the use of query logs and search result pages in information retrieval research, and briefly discuss search transparency and accountability as well as the Internet Archive.

2.1 Query Logs

Table 2 compiles an overview of 14 public and 31 private query logs from a focused literature review. Using the DBLP³ title search, we screened all publications that mention “query log”, “click log”, or “clickthrough” in their title—a high-precision heuristic to ensure logs play a role, at the expense of recall. From the 642 publications found, the 492 related to information retrieval (e.g., not databases) were downloaded. We then searched for occurrences of the pattern

³<https://dblp.org/>

“<number> <qualifier> queries” in them with regular expressions, assuming that virtually all researchers using query logs also specify how large they are.⁴ This facilitated the manual extraction of passages and tables from 120 randomly selected publications for Table 2. Some entries were added manually to cover public logs.

Despite the fact that query logs are rarely published, researchers in academia have sought alternative means of access, usually by collaborating with many search providers large and small. Weighted by the number of publications per log, research with very large query logs was conducted in industry and at major search providers. The AOL log [95] is the only exception. Together with its recent extension AOLIA [82] they are the largest publicly available query logs. The AltaVista log and the Baidu log are the two largest private logs reported. Our Archive Query Log is on par with the latter two. The ratio of unique queries to all queries averages 0.24. Our log’s ratio of 0.18 is slightly lower due to its multilingual nature. In addition to queries, organic search engine query logs may include information about users, clicks, sessions, and results, while our log only includes queries, SERPs, and the result documents themselves.

Given the main tasks for which query logs are used, the AQL can be used to study many—though not all—of them at a scale not easily attainable for academic researchers: Query understanding involves analyzing user information behavior. Subtasks include determining the user’s search intent [62] and examining user populations [63], particularly with respect to geographic [104] and temporal [64] dimensions. In addition, much research has focused exclusively on how people search for health information, from both consumer [94] and expert perspectives [101]. Query suggestion involves exploiting query logs to recommend alternative queries to the user. Subtasks include clustering [17], query similarity [25], and query expansion using relevance feedback [44, 45, 59]. The AQL may support both tasks in general, in particular when used for pre-training. However, model transfer to a specific application domain will be required, as well as organic log data for practical use cases.

Session analysis examines how users reformulate their queries across one or more sessions [5], a key subtask being session detection [52, 56]. User modeling involves analyzing logs to build models of user interaction. Subtasks include examining the distributions of query lengths and query terms [65, 102, 105, 110], relevance feedback mechanisms [103], and what users consider relevant [61]. The AQL does not support these tasks; it lacks session or user data.

Learning to rank exploits query logs to derive effective ranking models. Subtasks include developing click models [70, 71] and models that incorporate implicit feedback such as dwell time on pages [3, 113]. The AQL supports this task despite the lack of click data. Craswell et al.’s [38] rationale for the design of MS MARCO corroborates this claim. Here, only passages (judged for relevance) from the top-ranked documents returned by Bing for a query are included as ground truth for training, which has proven to be sufficient to yield effective retrieval models. The same is true for the AQL, where the ranked results of third-party retrieval models encode the domain expertise and the implicit relevance feedback from query logs that the respective search providers incorporated into them.

⁴Examples: “1 million queries”, “386,879 queries”, “386 879 queries”, or “386k queries”. A qualifier is a sequence of up to 20 characters excluding end of sentence punctuation.

2.2 Search Engine Result Pages

Search engine result pages (SERPs) are how search engines present results to users in response to a query. SERPs for web search typically consist of a list of links to web pages ranked by their relevance to the user’s query, along with additional information such as snippets, images, and other features designed to help users meet their information needs. SERPs have been studied in information retrieval research for many years to understand how users interact with them, how they can be improved, and how they can present information more effectively to better meet user needs. The AOL contains the SERP of 39 % of its queries.

One area of SERP research has focused on understanding how users interact with search results. Researchers have used techniques such as eye-tracking [10, 68, 69] and brain monitoring [88] to study how users perceive SERPs. These studies have led to a deeper understanding of how to improve the presentation of results and the ranking algorithms used by search engines. Another area of SERP research is the study of their design and layout [90, 91]. These longitudinal studies show how SERPs evolve in response to new technologies. The AQL provides millions of archived SERPs which include all necessary assets for showing them in a browser, so that they can be used for user studies and offline experiments.

2.3 Transparency and Accountability

In November 2022, the Digital Market Act [1] and the Digital Services Act [2] came into force in the European Union. The former applies primarily to “gatekeepers” in digital markets, such as Google for the search market, the latter to all digital services that act as (information) “intermediaries”. Both laws contain provisions that, among other things, require search providers to increase data privacy, transparency, and accountability, with the goal of ensuring fair and open digital markets. In particular, legislators are allowed to exercise regulatory and market investigation powers, which may include looking into the algorithms used. The AQL complements these measures and also gives civilian initiatives the means to conduct independent investigations of search providers. Previous studies on search accountability raise the question of how to inform users about a search engine’s retrieval algorithms to raise awareness of how they work [37, 76, 77] and to ensure unbiased results [51, 78]. While we cannot consider all previous work in this context, a recent overview was provided at the FACTS-IR workshop [92] on fairness, accountability, confidentiality, transparency, and safety in information retrieval. In terms of both algorithm transparency and search engine accountability, archived search result pages are perhaps one of the best representations of a search engine’s behavior at a given time, and archiving them on a large scale allows for corresponding post-hoc analyses.

2.4 Internet Archive

The Internet Archive is a nonprofit digital library that has grown to become the largest and most comprehensive digital library in the world since its inception in 1996. In addition to providing access to extensive archives of books, audio recordings, videos, images, and software, the Internet Archive’s best-known service is probably the Wayback Machine, which provides a digital archive of the web.⁵ At

⁵<https://web.archive.org/>

the time of writing, it contains 806 billion web pages. We believe that the AOL accumulated both due to accidental crawling by their crawlers and intentional archiving by their users, since any user can request archiving of any publicly accessible URL. AOLIA [82] extends the original AOL log by providing links to archived versions of its search results, originally specified as URLs only. However, AOLIA's use is limited since only the landing pages of the domains from which search results were originally found could be recovered.

3 MINING THE ARCHIVE QUERY LOG

Besides general-purpose search engines, many other websites such as social media platforms offer a search function (a query field). The answer to a query is often encoded as URL linking to a SERP, which is displayed to the user. Like other URLs, these “SERP URLs” can be linked to by web pages and are thus included in automated web crawls. The Internet Archive, as the world’s largest digital library of archived web pages, is likely to include many SERPs, a fact which can be exploited for large-scale query log mining. This section describes a multi-step process to mine a query log from the Internet Archive’s Wayback Machine, which eventually becomes the Archive Query Log (see Figure 1).

First, a list of popular search providers including general-purpose search engines and all kinds of media platforms is compiled (see upper part of Figure 1; Section 3.1). Second, for each provider the second-level domains, known subdomains, and URL patterns under which SERPs are likely to be found are semi-automatically generated and the URL captures found in the Internet Archive (using the CDX API⁶ of the Wayback Machine) are aggregated (Section 3.2). Third, the queries are extracted from the URLs using provider-specific parsers (Section 3.3). Fourth, the HTML content of archived SERPs is downloaded and parsed for metadata (URLs, titles, snippets, language, etc.; Section 3.4). Both queries and metadata form the AOL-22 corpus (Section 3.5) which will be made accessible via the TIRA platform [96] as discussed in Section 4.

3.1 Search Provider Collection

Our search provider collection contains both (1) websites that primarily act as search engines, and (2) highly relevant websites that have been identified by their Alexa Rank.⁷

Regarding (1), we exploit a dedicated list of search engines on Wikipedia which we extend manually.⁸ Regarding (2), we take the 3,088 archived snapshots of the Alexa top-1M ranking between June 2010 and November 2022⁹ and apply reciprocal rank fusion [36] considering the 1,000 highest ranked website domains of each snapshot. The resulting list of 13,647 domains is narrowed down to 951 search providers by checking whether a search bar is present on the website’s landing page of the respective provider. For this purpose we load the landing page directly or from an Internet Archive snapshot from 2022, render the page if JavaScript content is found, and check for HTML forms or <div> elements containing the pattern “search” in its attributes.

⁶<https://github.com/internetarchive/wayback/tree/master/wayback-cdx-server>

⁷The Alexa Rank was a ranking system that reflected the global popularity of website domains based on estimated visits; it was shut down end of 2022.

⁸See https://en.wikipedia.org/wiki/List_of_search_engines

⁹See https://web.archive.org/web/*s3.amazonaws.com/alexa-static/top-1m.csv.zip

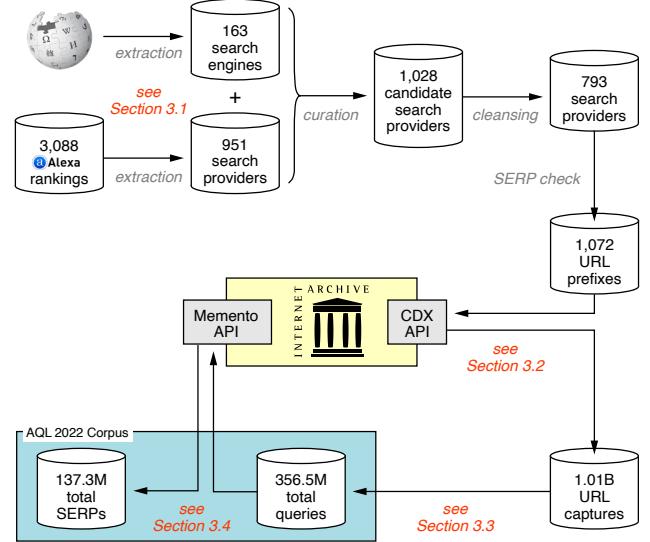


Figure 1: Overview of the AOL creation process.

The merged list of 1,028 unique candidate search providers is used to identify relevant URL patterns as well as suitable approaches for query extraction (see Section 3.2). Further manual curation steps weed out providers because they have been identified as spam, do not encode the search query in their URL, or offer only an autocomplete search that links directly to a page. Also, search providers are merged when more than one of their second-level domains appears in the merged Alexa list (e.g., so.com is merged into 360.com). After these analyses and curation steps, the list includes 793 search providers.

3.2 Provider Domains and URLs

Since many search providers are available under multiple second-level domains and/or subdomains,¹⁰ we expand the list of the 793 provider domains manually and with publicly available lists.¹¹

However, on high-traffic domains, only a small fraction of all archived URLs is likely to link SERPs and thus is relevant for our purposes. To identify common prefixes of URLs that contain queries, we submit multiple test queries using the search provider’s query field and examine the URL the request is redirected to. For discontinued or otherwise inaccessible websites, we resort to the most recent functional snapshot of the search provider’s homepage in the Internet Archive. The final list of 1,072 URL prefixes is used to filter the list of available captures with the help of the Internet Archive’s CDX API. This API allows to request a list of URLs by the crawling date they were archived in the Wayback Machine for a certain domain or URL prefix. Via the CDX API, a list of all available captures for each of the 1,072 URL prefixes is retrieved and filtered for successful captures with HTML content (i.e., HTTP status code 200). This process further narrows the search provider list with archived SERPs to 550.

¹⁰Google lists 190 supported domains https://www.google.com/supported_domains

¹¹E.g., <https://github.com/JamieFarrelly/Popular-Site-Subdomains>

(a) *Query parameter:*

scheme	network loc.	path	query string
URL prefix	query	offset	
https://google.com/search?q=covid+19+usa+map&start=10&ei=...			

(b) *Path segment:*

scheme	network loc.	path
URL prefix	page	query
https://chefkoch.de/rs/s0/backen%20dinkelmehl/		Rezepte.html

Figure 2: Illustration of the components of SERP URLs and the relevant parts for query extraction by (a) query parameter or (b) path segment.

Altogether, 1.1B URL captures along with their crawling date are collected, an average of 1.8M URLs per provider. Most of the captures originate from search engines (226M URLs, 22 %) with Google being the largest contributor of archived URLs (89M URLs, 9 %).

3.3 Query Extraction

To extract the query from a SERP URL, the URL is parsed into its components¹² and the query encoding is identified as one of three possible categories: (1) URL parameter, (2) path segment, or (3) fragment identifier. Examples of the first two patterns are illustrated in Figure 2.

For each of the three patterns, configurable query parsers are built using the `urllib` Python package:¹³ (1) parsing a query parameter by its name (e.g., name `q` for the first URL in Figure 2), (2) parsing a segment from the URL’s path component by its index (e.g., index 2 for the second URL in Figure 2), or (3) parsing a parameter from the fragment identifier by treating it like a query parameter. In addition to the query, URLs can include a page number or offset. These help in reconstructing longer rankings from separate SERPs for the same query that are captured at nearly the same time. Google’s SERPs, for instance, are paginated with 10 results per page. Thus, the page number can be used to infer the continued ranks of documents on the next page.

To determine the query, page, and offset of the parsers for each search provider, we manually examined the captured URLs in a similar way as for URL prefixes (see Section 3.2) and derive suited URL parser types and parameters. Regular expressions are optionally used to limit parsing to specific URLs, and to further refine the parsers (e.g., removing prefixes such as `page-` in `/search/page-4`). The resulting set of parsers is applied to all available captures of all search providers, ordered by preference, so that the first parser that returns a non-empty query, page, or offset is used.

Altogether a total of 356.5M URLs containing queries are collected. Again, the majority of queries stems from search engines (162M queries, 46 %) such as Google (73M queries, 20 %) and Baidu (70M queries, 20 %). On average, 648,092 queries are extracted per search provider. This unfiltered set of queries contains large

¹²RFC 2396; <https://datatracker.ietf.org/doc/html/rfc2396.html>

¹³<https://docs.python.org/3/library/urllib.html>

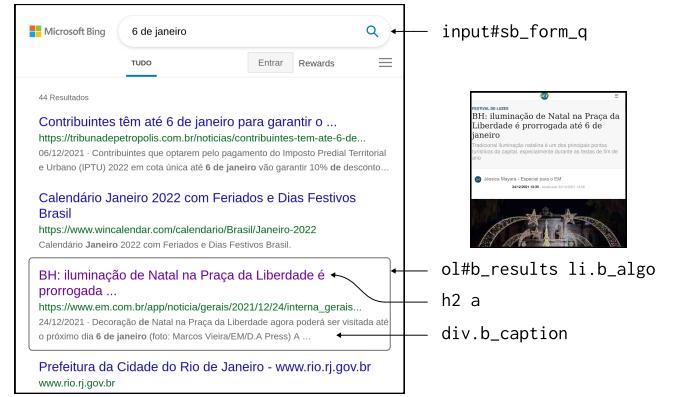


Figure 3: Screenshot of an archived Bing SERP along with the CSS selectors for the query and result items. The nearest archived snapshot of the search result is shown on the right.

amounts of duplicates (288M, 81 %) for which we identify three reasons: (1) the same query is captured at different times, (2) the query is captured at approximately the same time but with different result page offsets, and (3) the same query is captured as issued from different users (e.g., if a user identifier is included in the URL). This is supported by the fact that search engines are the main contributors of duplicates (131M); however, government sites have the highest share (91 % of all queries from government sites).

We create a set of unique queries for each search provider by selecting a representative query URL (the capture with the shortest query string) from each group with the same parsed query.¹⁴ If a group of duplicates has multiple captures with the same query parameter and URL length, the representative URL is chosen by lexicographic order. We refrain from using the capture’s timestamp as a tie-breaker to not favor older or newer captures. The deduplication results in 64.5M unique queries across the final list of 550 search providers. There are 117,353 unique queries per each provider on average. Again, the search engines make up the majority of deduplicated queries (31M, 45 %).

3.4 SERP Acquisition and Parsing

Previous query logs rarely contain results for the logged queries (see Section 2). However, since the AQL is obtained from web page captures in the Internet Archive, it naturally contains the full ranking of results for many of its queries. By downloading and parsing the archived SERPs, one has access to the full ranking of results for each search query (including the processed query itself, as it appears in the query field of the SERP). Parsing the search results including result titles, referenced URLs, snippets, and the query from a SERP facilitates not only the comparison of different search provider’s ranking functions but also the evaluation of their query understanding and reformulation techniques.

We download the SERP HTML content for the unique search queries identified in the previous step and save it in WARC format.¹⁵ For the 20 most popular search providers SERPs are downloaded

¹⁴As split according to RFC 2396.

¹⁵ISO 28500:2017; <https://iipc.github.io/warc-specifications/>

for all unique URLs; for the remaining providers the download is limited to a maximum of 100,000 due to resource constraints. Connection timeouts and other errors during download are handled by repeating the download up to 10 times, after which we consider the archived SERP snapshot to be unavailable. Based on the completed downloads, we estimate a total of 137.3M SERPs to be available, most of which originate from search engines (ca. 46 %) and social media platforms (ca. 16 %). The downloads are ongoing, and we plan on scaling them up (see Section 5) to compile the full set of estimated SERPs available (see Section 4).

From the downloaded SERPs the search result ranking is extracted and processed using a configurable parser pipeline based on FastWARC [19], Beautiful Soup,¹⁶ and Approval Tests.¹⁷ In detail, the CSS path or selector¹⁸ to the result list items is specified, as well as the path from each individual result item to its title element, the referenced URL anchor, and the snippet text. Similarly, the processed query is parsed based on the CSS path to the query field. Figure 3 shows an archived Bing SERP and highlights how CSS paths are used to select relevant HTML tags. Each search provider can have multiple parser configurations, ordered by preference, that, for example, account for a changed HTML structure after redesigns of a search provider’s SERPs.

We derive parser configurations (CSS paths) for the 50 most popular search providers by generating Approval Tests according to following workflow for a provider: (1) Randomly sample 10 SERPs from the downloaded SERPs. (2) For each SERP manually annotate the expected ranking and query. (3) Apply the existing parser configurations to the sampled SERPs. (4) Compare the parsed results to the annotations. (5) If the annotations do not match, inspect the HTML page in a browser, adapt or extend missing patterns, and add them to the provider’s parser configurations. New configurations are added iteratively until all sampled SERPs are correctly parsed. Altogether, 70 parser configurations for SERPs and 57 parsers for processed queries of the 50 most popular providers are derived. With additional manual tests for the 10 most popular providers (see Table 1), the parsers pass a test suite of 444 Approval Tests. The code base is available open source.¹⁹

3.5 The Archive Query Log 2022 (AQL-22)

We merge the filtered URL captures, queries, and SERPs into a single corpus to be used in subsequent analyses (see Section 4). This corpus, the Archive Query Log 2022, consists of two artifacts: (1) a set of queries and (2) a set of ranked documents (search result snippets). Both artifacts are stored in a GZIP-compressed, newline-delimited JSON format.²⁰ To create the query set, each captured URL is assigned a unique identifier based on the full URL string and timestamp of the capture.²¹ The captured URL is associated with its parsed query, the location of the stored copy of the SERP, and the processed query and search results parsed from that SERP. In addition, we include the URL to the SERP’s archived snapshot

¹⁶<https://pypi.org/project/beautifulsoup4/>

¹⁷<https://pypi.org/project/approvaltests/>

¹⁸See <https://facelessuser.github.io/soupsieve/> for a list of supported selectors. The CSS path to an HTML element can be inferred using a web browser’s developer tools.

¹⁹AQL code: <https://github.com/webis-de/archive-query-log>

²⁰<https://jsonlines.org/>

²¹Name-based SHA-1 UUID according to RFC 4122: <https://rfc-editor.org/rfc/rfc4122>

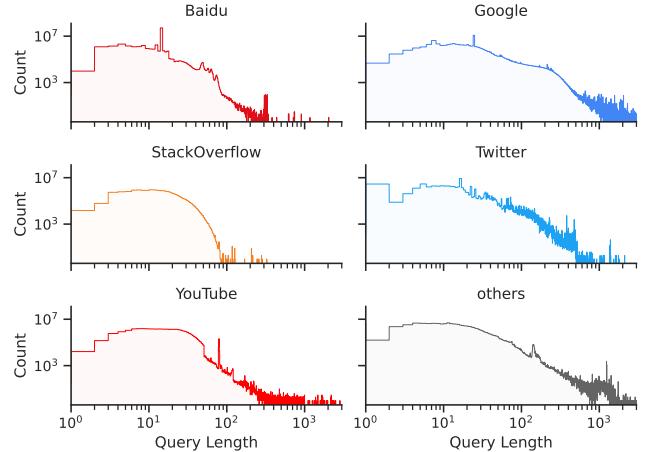


Figure 4: Distribution of query lengths for 5 search providers contributing the highest amount of queries. The remaining search providers are grouped as “others”.

on the Wayback Machine and tag the query language based on the parsed query text using cld3.²²

The set of result documents is created by concatenating all ranked search results (i.e., rank, snippet text with title, and document URL to the referenced web page) from all parsed SERPs. Each document is assigned a unique identifier based on the document URL, the timestamp of its origin query, and the rank of the snippet on the SERP. We also associate each document with the attributes of the corresponding query in the query set. Two additional fields are the URL to the nearest available snapshot of the SERP on the Wayback Machine and the snippet language as tagged using cld3 based on the snippet’s title and text.

4 ANALYSIS

In order to provide a better understanding of the corpus, we conduct some analyses on the query and SERP characteristics and highlight potential use cases. Detailed analyses of the AQL will be the subject of future work. At the time of writing, we could download and parse all URLs and queries. However, due to computational constraints, only a subset of all available SERPs could be parsed yet (see Section 4.4) and are used for our analyses.

4.1 Query Characteristics

A central feature of the AQL is its diversity. In addition to the variety of search providers, it also features a total of 104 different languages²³ with Chinese and English being the most frequently used query languages (see Table 3). The query length in the AQL follows a skewed distribution, with most queries containing between 5 and 20 characters. Figure 4 provides a visualization for the 5 search providers with the most queries. We inspect samples of queries with 5, 10, 100, and 1000 characters. Very short queries are often Mandarin keywords (e.g., 长袖衬衫男, “men’s long sleeve shirt”) or single English words (e.g., video). Queries with 10 characters

²²Google’s Compact Language Detector; <https://github.com/google/cld3>

²³Out of 107 detectable with cld3.

Table 3: The Archive Query Log 2022 (AQL-22) in detail. Categories manually annotated. Top-3 languages tagged by cld3. Ticks in the timelines indicate days of archival from Jan 1999 to Dec 2022. Number of SERPs and results estimated (c.f. Section 4.4).

Search provider	Category	URLs	Queries				SERPs	Results			
			Total	Total	Unique	Lang.		Estimate	Estimate	Lang.	Timeline
G Google	Search engine	89,364,948	72,673,044	19,953,592	en, th, zh	███████	27,979,122	223,073,409	en, de, pt	███████	███████
Y YouTube	Media sharing	41,846,525	41,365,166	11,250,179	ru, ko, ja	███████	15,925,589	339,164,817	ru, en, ko	███████	███████
B Baidu	Search engine	78,506,825	69,619,339	2,900,878	zh, ga, ja	███████	26,803,445	107,623,054	zh, en, mr	███████	███████
Q QQ	Web portal	515,895	513,608	51,228	zh, ja, lb	██████	197,739	2,101,929	–	██████	██████
F Facebook	Social media	3,131,212	159,087	35,492	ca, en, bs	██████	61,249	651,065	–	██████	██████
W Yahoo!	Web portal	8,787,707	2,827,103	1,232,589	en, la, de	███████	1,088,435	9,242,960	en, es, pt	███████	███████
A Amazon	E-commerce	66,795,164	776,127	315,068	en, ja, zh	██████	298,809	7,789,151	en, ja, it	██████	██████
W Wikipedia	Wiki	68,547,509	1,707,058	621,971	sv, zh, en	██████	657,218	6,986,104	–	██████	██████
J JD.com	E-commerce	4,370,884	3,902,604	370,473	zh, hr, ja	██████	1,502,503	15,971,325	–	██████	██████
360 360	Search engine	1,495,365	1,090,152	65,596	zh, ja, mg	██████	419,708	3,487,458	zh, mr, en	██████	██████
W Weibo	Social media	6,245,012	5,324,385	1,886,458	zh, ja, en	██████	2,049,889	21,789,936	–	██████	██████
R Reddit	Forum	94,162	89,492	36,852	en, la, de	██████	40,369	429,115	–	██████	██████
Vk.com	Social media	643,354	153,642	46,134	ru, sr, ky	██████	59,152	628,775	–	██████	██████
C CSDN	Social media	21,863	946	736	zh, en, vi	██████	810	8,610	–	██████	██████
B Bing	Search engine	11,263,539	6,152,425	2,253,965	en, zh, pt	██████	2,368,684	12,625,017	en, pt, fr	██████	██████
T Twitter	Social media	55,499,532	48,084,528	3,869,382	ja, en, gl	██████	18,512,544	241,835,785	en, ja, es	██████	██████
Twitch	Streaming	21,931	15,225	11,445	en, zh, de	██████	12,587	133,797	–	██████	██████
eBay	E-commerce	7,927,123	5,507,532	1,379,646	zh, en, la	██████	2,120,400	25,130,312	en, es, de	██████	██████
N Naver	Search engine	1,063,991	756,153	400,490	ja, ko, vi	██████	291,119	2,938,902	ko, en, hi	██████	██████
AliExpress	E-commerce	4,620,331	1,861,642	55,849	en, lb, fy	██████	716,732	5,719,031	en, fr, ru	██████	██████
: 530 others		559,225,614	93,871,236	17,806,322	en, zh, de	██████████	36,194,558	382,644,117	en, zh, de	██████████	██████████
Σ 550		1,009,988,486	356,450,494	64,544,345	zh, en, ga	██████████	137,300,661	1,409,974,669	en, ru, ko	██████████	██████████

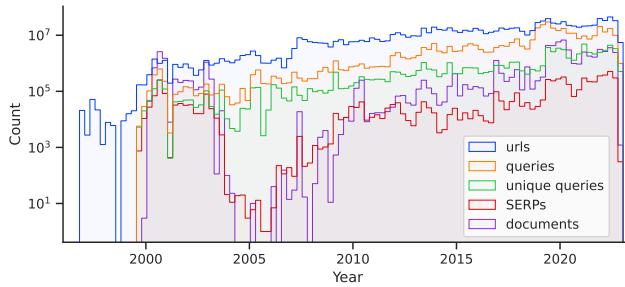


Figure 5: Time coverage of the total amount of different data types collected for the AQL-22, per quarter.

are mostly keyword-style queries from Latin languages (e.g., comic font) or hashtags (e.g. #чемпионат, “championship”). Most longer queries extensively use search operators like site: and order:, include literature references, or include long multi-line text like stack traces from errors in programming.

Second, we evaluate whether obscene or unwanted terms comprise a large share of the AQL. We use lists of obscene words for 27 languages²⁴ and expand the list of English terms with new expressions found in the downloaded queries. We check each query from the two most frequent languages, Chinese and English, for their lists of obscene terms. Overall, only 1.30 % of all queries contain obscene terms. The highest share of these obscene queries were observed on pornography (19.48 %), torrent (3.73 %), and forum (2.87 %) websites. For non-pornographic search providers, most

²⁴Compiled by Shutterstock; <https://github.com/LDNOOBW>List-of-Dirty-Naughty-Obscene-and-Otherwise-Bad-Words>.

Table 4: Most frequent document domains in the top-5 or the top-10 search results compared to references to the search provider’s own domain (○) or 791,646 other domains (…).

Top	W	Y	f	in	hi	ig	a	o	...	○
5	2.9 %	0.8 %	0.6 %	0.4 %	0.3 %	0.3 %	0.3 %	0.2 %	24.3 %	69.6 %
10	2.2 %	0.7 %	0.5 %	0.3 %	0.3 %	0.2 %	0.3 %	0.3 %	24.6 %	70.4 %

stem from heroturko.org (16.13 %, e-commerce), reddit.com (5.05 %, forum), and kat.cr (4.08 %, forum).

Regarding time coverage, Figure 5 shows that the archival of SERPs dropped between 2004 and 2010 for unknown reasons, which might indicate that using more specialized SERP parsers are required, or that results were loaded using Javascript. The queries, however, extend over the whole timespan, with tens of thousands of queries recorded for the early 2000’s as well. Table 3 contains an overview of the 20 most popular services’ time coverage.

4.2 SERP Characteristics

In Table 4, we consider the most frequently referenced domains from search result URLs as an indicator of plausible rankings. Excluding frequent self-references (e.g., to internal redirect pages), by far the most frequently ranked domain is wikipedia.org contributing 2.9 % of all top-5 results and 2.2 % of the top-10. Other popular domains like youtube.com and facebook.com also frequently appear on high ranks. The most frequent languages are shown in Table 3. Interestingly, Chinese is not among the top-3 languages of search results, even though it is the most frequently used query language, representing a bias that should be evaluated in future work.

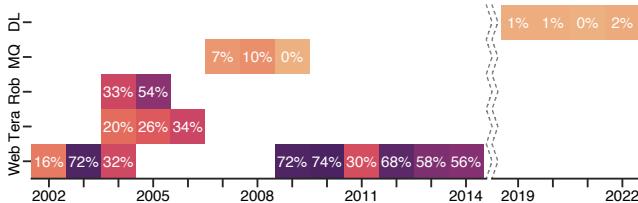


Figure 6: Query overlap with TREC Robust (Rob), Terabyte (Tera), Million Query (MQ), Web, and Deep Learning (DL).

4.3 Use Cases

The AQL opens up a variety of use cases for the IR community. We highlight two promising applications.

First, we evaluate the exact overlap of the queries in the AQL with the collections used in various TREC tracks from 2004 to 2022 [7, 8, 21, 27–34, 39–43, 107, 108]. As shown in Figure 6, the highest overlap exists with the Web tracks, specifically in 2010 (74 %), 2003, and 2009 (both 72 %). The lowest overlap was found with the Deep Learning tracks, ranging between 0–2 %. The high overlap on older Web tracks poses an interesting opportunity for enriching existing benchmarks. While query logs have been used previously in shared community tasks [83], shared tasks often specify only one query for each topic. We propose to sample semantically similar queries from the AQL to generate topics with user query variations [16] automatically. On the other hand, the low overlap with the Deep Learning tracks highlights a sampling bias in creating the Deep Learning topics. The topics were sampled from the official eval set of MS MARCO, which includes only natural language questions from a Bing query log [43, 89]. The AQL, on the other hand, contains a much broader range of queries, including queries from other search providers and non-question-like queries. Therefore, we propose using the AQL to create new, “harder” Deep Learning topics that are more representative of other kinds of queries users submit.

Second, we demonstrate how global trends are reflected in the AQL on the example of the Covid-19 pandemic. In Figure 7, we count the occurrences of the terms covid 19, sars cov 2, and corona virus each month since the outbreak in 2019. A peak can be observed during the first global lockdowns in early 2020, but overall, interest in the pandemic has yet to stagnate. The example showcases how the AQL enables unique opportunities for diachronically analysing global trends.

4.4 Total Size Estimates

As Section 3 explains, we have only downloaded and parsed a subset of all available SERPs. Based on our results so far, we estimate 70 % of all SERP snapshots to be available for download. Assuming an estimated parsing success rate of 55 % and 10.6 results per SERP, we expect the total number of parsed SERPs in the AQL-22 to be 137.3M with 1.4B search results. As outlined in Section 5, we continue to download and parse SERPs and look forward to expanding the AQL with the IR community.



Figure 7: Timeline of Covid-19-related terms in the AQL-22.

5 DISCUSSION

Access to query logs has long been an insurmountable barrier to answering critical questions about the search economy at large—if not to ask them in the first place. As a consequence, the media and public were left with no choice but to trust search engines on questions such as “How accountable are search providers in terms of measures of interest, like representation and fairness?”, “How have these accountability measures changed for these organizations over time?”, and “How honest has self-reported accountability of these organizations been?”. As the most extensive public query log to date, the AQL enables detailed analyses of and thus facilitates the public discourse on the search industry. It also furthers the research on information retrieval, whose retrieval models are often presumed to be behind or at least detached from those of industry players (e.g., Azzopardi et al. [13]). Using the AQL, researchers will also be able to answer questions such as “How far is academic information retrieval research behind industry?”, “How much do query logs contribute compared to other ranking signals?”, and “What are domain-specific differences?”.

However, with the scale of the AQL, several ethical and legal considerations also arise, particularly around personally identifiable information or illegal content. To address inherent risks, we release the data by imposing a barrier to access that minimizes potential harms while giving the information retrieval community as much freedom to conduct their research as possible. In addition, we also acknowledge the challenges of creating such extensive collections and discuss our plans for opening contributions to the AQL.

5.1 Accessing the AQL

Our goal with releasing the AQL is to do so responsibly and in a privacy-preserving manner. We work towards that goal using the TIRA platform [50, 96] for any analysis on the AQL one wishes to conduct. TIRA has been used since 2012 [53] to facilitate shared tasks with software submissions while ensuring that the submitted software can process the data without giving the participants themselves access. TIRA achieves this through sandboxing, i.e., disconnecting the software from the internet while it is running and thus ensuring that it can not leak data. We added the AQL to TIRA to allow researchers to submit their analysis software as Docker images. The platform is open to the public, and we provide examples and documentation on how to perform analyses on the AQL.²⁵ Specific shared tasks will be developed as well.

TIRA allows running arbitrary software packaged in Docker images on the AQL dataset in a privacy-preserving way, as the analysis results are blinded until reviewed. Specifically, we review the output and the software installed in the Docker image to ensure

²⁵<https://tira.io/task/archive-query-log>

no sensitive data is leaked. TIRA runs the software in a Kubernetes cluster (1,620 CPU cores, 25.4 TB RAM, 24 GeForce GTX 1080 GPUs) with a timeout of 24 hours, so that almost any evaluation is supported.²⁶ In summary, TIRA provides the ideal means to work with the AQL, ensuring sensitive query log data remains secure and is used for academic research in a responsible fashion.

5.2 Limitations and Scalability

While creating the AQL, we encountered several technical limitations that guide future optimizations and improvements. First, the various parsers for creating the AQL-22 corpus were written semi-automatically. This approach was error-prone and inefficient, requiring much manual work. When building future versions of the AQL, we plan to train token classification models like BERT [47] to automatically generate query parsers based on a training set derived from our existing parsers. Manually finding the correct CSS paths for snippets on a SERP is a similarly tedious process that can benefit from wrapper generation [75], which has successfully been applied to web page parsing [23].

Second, dynamic content cannot be interpreted by our existing parsers. The SERPs of DuckDuckGo, for instance, are loaded dynamically using JavaScript and thus cannot be parsed from just the archived HTML snapshot, yet the search results are still archived as a different record. To overcome this limitation, we plan to use a headless browser to render the SERPs and then parse them. A helpful library for this type of content extraction is Scriptor.²⁷

The last and most pressing limitation, however, is that all URL captures and SERP contents must first be downloaded from the Internet Archive, which is restricted by both rate limits and network bandwidth. As estimated in Section 4.4, currently 93 % of the SERPs still need to be fully downloaded from the Internet Archive and thus could not yet be parsed. Hence, more computational resources are required to use this extensive collection. In this regard, we will reach out to the Internet Archive whose privileged access to their infrastructure will allow for much faster compilation of the data.

5.3 Contributing to the AQL

There is an inherent boundary between search providers and researchers when using query logs. The AQL lowers this boundary by exploiting the web archival process of the Internet Archive. As we have described above, physical limitations, such as network speed, restrict the rate at which we can further grow the AQL. One way to overcome such limitations is to distribute computations across an open community, an approach that has been successfully employed in mathematics.²⁸ We therefore open source our code to allow the community to contribute query and SERP parsers.

6 CONCLUSION

The Archive Query Log provides an unparalleled academic resource. It consists of over 356 million queries, over 166 million SERPs, and over 1.7 billion results extracted from the SERPs, all coming from 550 search providers spanning 25 years. The AQL is the largest and most diverse query log ever publicly available. From an academic

²⁶We can extend the timeouts and available resources individually if the need arises.

²⁷<https://github.com/webis-de/scriptor>

²⁸E.g., at the Great Internet Mersenne Prime Search: <https://mersenne.org/>

perspective, the AQL will enable researchers to tackle challenges in information retrieval that were not possible until now, ranging from the development of new models for retrieval, query suggestion or query prediction to large-scale diachronic analyses of search engines; to name the most salient research avenues. Furthermore, our release plan for accessing the AQL minimizes the harm to society and will allow researchers to safely research the transparency and accountability of search engines while protecting user privacy.

In this paper, we have documented the initial version of the AQL (i.e., AQL-22). We have plans to release future versions of the AQL that will further expand the collection. First, we will continue to add to the long tail of search providers and continue our efforts to download and extract more data from the Internet Archive. Continuing to grow the types of data provided, the next version of the AQL will also include the content of web pages for each result in a SERP. Not least, we will investigate the training of large re-ranking models based on this data.

Altogether, the AQL is an exceedingly valuable resource for researchers and will enable advances in information retrieval research that were previously insurmountable due to the relatively low scale of query logs. Because of its scope, size, and diversity we consider the AQL a significant contribution to the community, and these dimensions will continue to grow as we build upon and expand future versions of the AQL.

ACKNOWLEDGMENTS

This publication has received funding from the European Union's Horizon Europe research and innovation programme under grant agreement No 101070014 (OpenWebSearch.EU, <https://doi.org/10.3030/101070014>)

REFERENCES

- [1] 2022. Regulation (EU) 2022/1925 of the European Parliament and of the Council of 14 September 2022 on Contestable and Fair Markets in the Digital Sector and Amending Directives (EU) 2019/1937 and (EU) 2020/1828 (Digital Markets Act). *Official Journal of the European Union* 265 (2022), 1–66. <https://eur-lex.europa.eu/eli/reg/2022/1925>
- [2] 2022. Regulation (EU) 2022/2065 of the European Parliament and of the Council of 19 October 2022 on a Single Market For Digital Services and Amending Directive 2000/31/EC (Digital Services Act). *Official Journal of the European Union* 277 (2022), 1–102. <https://eur-lex.europa.eu/eli/reg/2022/2065>
- [3] Eugene Agichtein, Eriq Brill, and Susan T. Dumais. 2006. Improving Web Search Ranking by Incorporating User Behavior Information. (2006), 19–26. <https://doi.org/10.1145/1148170.1148177>
- [4] Maristella Agosti, Franco Crivellari, and Giorgio Maria Di Nunzio. 2012. Web Log Analysis: A Review of a Decade of Studies About Information Acquisition, Inspection and Interpretation of User Interaction. *Data Mining and Knowledge Discovery* 24, 3 (2012), 663–696. <https://doi.org/10.1007/s10618-011-0228-8>
- [5] Maristella Agosti and Giorgio Maria Di Nunzio. 2007. Web Log Mining: A Study of User Sessions. In *Proceedings of the 10th DELOS Thematic Workshop on Personalized Access, Profile Management, and Context Awareness in Digital Libraries, PersDL 2007, Corfu, Greece, June 29–30, 2007*. 70–74. <https://www.academia.edu/download/33383010/72.pdf>
- [6] Farooq Ahmad and Grzegorz Kondrak. 2005. Learning a Spelling Error Model from Search Query Logs. In *Proceedings of the Human Language Technology Conference and the Conference on Empirical Methods in Natural Language Processing, HLT/EMNLP 2005, Vancouver, British Columbia, Canada, October 6–8, 2005*. ACL, 955–962. <https://aclanthology.org/H05-1120>
- [7] James Allan, Javed A. Aslam, Virgil Pavlu, Evangelos Kanoulas, and Ben Carterette. 2008. Million Query Track 2008 Overview. In *Proceedings of the Seventeenth Text REtrieval Conference, TREC 2008, Gaithersburg, Maryland, USA, November 18–21, 2008 (NIST Special Publication, Vol. 500-277)*. National Institute of Standards and Technology. https://trec.nist.gov/pubs/trec17/papers/MQ_OVERVIEW.pdf

- [8] James Allan, Ben Carterette, Blagovest Dachev, Javed A. Aslam, Virgilii Pavlu, and Evangelos Kanoulas. 2007. Million Query Track 2007 Overview. In *Proceedings of the Sixteenth Text REtrieval Conference, TREC 2007, Gaithersburg, Maryland, USA, November 5–9, 2007 (NIST Special Publication, Vol. 500-274)*. National Institute of Standards and Technology. <http://trec.nist.gov/pubs/trec16/papers/1MQ.OVERVIEW16.pdf>
- [9] Avi Arampatzis, Jaap Kamps, Marijn Koolen, and Nir Nussbaum. 2007. Deriving a Domain Specific Test Collection from a Query Log. In *Proceedings of the Workshop on Language Technology for Cultural Heritage Data, LaTeCHACL 2007, Prague, Czech Republic, June 28, 2007*. ACL, 73–80. <https://aclanthology.org/W07-0910>
- [10] Ioannis Arapakis, Joemon M. Jose, and Philip D. Gray. 2008. Affective Feedback: An Investigation Into the Role of Emotions in the Information Seeking Process. In *Proceedings of the 31st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2008, Singapore, July 20–24, 2008*. ACM, 395–402. <https://doi.org/10.1145/1390334.1390403>
- [11] Cédric Argenton and Jens Prüfer. 2007. The Structure of Search Engine Law. *Iowa Law Review* 93, 1 (11 2007), 1–63. <https://ssrn.com/abstract=979568>
- [12] Cédric Argenton and Jens Prüfer. 2012. Search Engine Optimization with Network Externalities. *Journal of Competition Law & Economics* 8, 1 (3 2012), 73–105. <https://doi.org/10.1093/joclec/nhr018>
- [13] Leif Azzopardi, Yashar Moshfeghi, Martin Halvey, Rami Suleiman Alkhawaldeh, Krisztian Balog, Emanuele Di Buccio, Diego Ceccarelli, Juan M. Fernández-Luna, Charlie Hull, Jake Mannix, and Sauparna Palchowdhury. 2016. Lucene4IR: Developing Information Retrieval Evaluation Resources using Lucene. *ACM SIGIR Forum* 50, 2 (2016), 58–75. <https://doi.org/10.1145/3053408.3053421>
- [14] Ricardo Baeza-Yates. 2015. Incremental Sampling of Query Logs. In *Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2015, Santiago, Chile, August 9–13, 2015*. ACM, 1093–1096. <https://doi.org/10.1145/2766462.27776780>
- [15] Ricardo A. Baeza-Yates and Felipe Saint-Jean. 2003. A Three Level Search Engine Index Based in Query Log Distribution. In *Proceedings of the 10th International Symposium on String Processing and Information Retrieval, SPIRE 2003, Manaus, Brazil, October 8–10, 2003 (Lecture Notes in Computer Science, Vol. 2857)*. Springer, 56–65. https://doi.org/10.1007/978-3-540-39984-1_5
- [16] Peter Bailey, Alistair Moffat, Falk Scholer, and Paul Thomas. 2016. UQV100: A Test Collection with Query Variability. In *Proceedings of the 39th International ACM SIGIR conference on Research and Development in Information Retrieval, SIGIR 2016, Pisa, Italy, July 17–21, 2016*. ACM, 725–728. <https://doi.org/10.1145/2911451.2914671>
- [17] Doug Beeferman and Adam L. Berger. 2000. Agglomerative clustering of a search engine query log. In *Proceedings of the Sixth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, Boston, MA, USA, August 20–23, 2000*. ACM, 407–416. <https://doi.org/10.1145/347090.347176>
- [18] Steven M. Beitzel, Eric C. Jensen, Abdur Chowdhury, David A. Grossman, and Ophir Frieder. 2004. Hourly Analysis of a Very Large Topically Categorized Web Query Log. In *Proceedings of the 27th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2004, Sheffield, UK, July 25–29, 2004*. ACM, 321–328. <https://doi.org/10.1145/1008992.1009048>
- [19] Janek Bevendorff, Martin Potthast, and Benno Stein. 2021. FastWARC: Optimizing Large-Scale Web Archive Analytics. In *Proceedings of the 3rd International Symposium on Open Search Technology, OSSYM 2021, Geneva, Switzerland, October 11–13, 2021*. <https://doi.org/10.5281/zenodo.6840911>
- [20] Bin Cao, Dou Shen, Kuansan Wang, and Qiang Yang. 2010. Clickthrough Log Analysis by Collaborative Ranking. In *Proceedings of the Twenty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2010, Atlanta, Georgia, USA, July 11–15, 2010*. AAAI Press. <https://aaai.org/ocs/index.php/AAAI/AAAI10/paper/view/1906>
- [21] Ben Carterette, Virgilii Pavlu, Hui Fang, and Evangelos Kanoulas. 2009. Million Query Track 2009 Overview. In *Proceedings of the Eighteenth Text REtrieval Conference, TREC 2009, Gaithersburg, Maryland, USA, November 17–20, 2009 (NIST Special Publication, Vol. 500-278)*. National Institute of Standards and Technology. <http://trec.nist.gov/pubs/trec18/papers/MQ09OVERVIEW.pdf>
- [22] Diego Ceccarelli, Sergiu Gordea, Claudio Lucchese, Franco Maria Nardini, and Gabriele Tolomei. 2011. Improving Europeana Search Experience Using Query Logs. In *Research and Advanced Technology for Digital Libraries. Proceedings of the International Conference on Theory and Practice of Digital Libraries, TPDL 2011, Berlin, Germany, September 26–28, 2011 (Lecture Notes in Computer Science, Vol. 6966)*. Springer, 384–395. https://doi.org/10.1007/978-3-642-24469-8_39
- [23] Chia-Hui Chang, Mohammed Kayed, Moheb R. Gigris, and Khaled F. Shaalan. 2006. A Survey of Web Information Extraction Systems. *IEEE Transactions on Knowledge and Data Engineering* 18, 10 (2006), 1411–1428. <https://doi.org/10.1109/TKDE.2006.152>
- [24] Michael Chau, Xiao Fang, and Olivia R. Liu Sheng. 2005. Analysis of the Query Logs of a Web Site Search Engine. *Journal of the American Society for Information Science and Technology* 56, 13 (2005), 1363–1376. <https://doi.org/10.1002/asi.20210>
- [25] Steve Chien and Nicole Immorlica. 2005. Semantic Similarity Between Search Engine Queries Using Temporal Correlation. In *Proceedings of the 14th International Conference on World Wide Web, WWW 2005, Chiba, Japan, May 10–14, 2005*. ACM, 2–11. <https://doi.org/10.1145/1060745.1060752>
- [26] Shui-Lung Chuang and Lee-Feng Chien. 2003. Enriching Web Taxonomies Through Subject Categorization of Query Terms from Search Engine Logs. *Decision Support Systems* 35, 1 (2003), 113–127. [https://doi.org/10.1016/S0167-9236\(02\)00099-4](https://doi.org/10.1016/S0167-9236(02)00099-4)
- [27] Charles L. A. Clarke, Nick Craswell, and Ian Soboroff. 2004. Overview of the TREC 2004 Terabyte Track. In *Proceedings of the Thirteenth Text REtrieval Conference, TREC 2004, Gaithersburg, Maryland, USA, November 16–19, 2004 (NIST Special Publication, Vol. 500-200)*. NIST.
- [28] Charles L. A. Clarke, Nick Craswell, and Ian Soboroff. 2009. Overview of the TREC 2009 Web Track. In *Proceedings of the Eighteenth Text REtrieval Conference, TREC 2009, Gaithersburg, Maryland, USA, November 17–20, 2009 (NIST Special Publication, Vol. 500-278)*. NIST.
- [29] Charles L. A. Clarke, Nick Craswell, Ian Soboroff, and Gordon V. Cormack. 2010. Overview of the TREC 2010 Web Track. In *Proceedings of the Nineteenth Text REtrieval Conference, TREC 2010, Gaithersburg, Maryland, USA, November 16–19, 2010 (NIST Special Publication, Vol. 500-294)*. NIST.
- [30] Charles L. A. Clarke, Nick Craswell, Ian Soboroff, and Ellen M. Voorhees. 2011. Overview of the TREC 2011 Web Track. In *Proceedings of the Twentieth Text REtrieval Conference, TREC 2011, Gaithersburg, Maryland, USA, November 15–18, 2011 (NIST Special Publication, Vol. 500-296)*. NIST.
- [31] Charles L. A. Clarke, Nick Craswell, and Ellen M. Voorhees. 2012. Overview of the TREC 2012 Web Track. In *Proceedings of the Twenty-First Text REtrieval Conference, TREC 2012, Gaithersburg, Maryland, USA, November 6–9, 2012 (NIST Special Publication, Vol. 500-298)*. NIST.
- [32] Charles L. A. Clarke, Falk Scholer, and Ian Soboroff. 2005. The TREC 2005 Terabyte Track. In *Proceedings of the Fourteenth Text REtrieval Conference, TREC 2005, Gaithersburg, Maryland, USA, November 15–18, 2005 (NIST Special Publication, Vol. 500-266)*. NIST.
- [33] Kevyn Collins-Thompson, Paul N. Bennett, Fernando Diaz, Charlie Clarke, and Ellen M. Voorhees. 2013. TREC 2013 Web Track Overview. In *Proceedings of the Twenty-Second Text REtrieval Conference, TREC 2013, Gaithersburg, Maryland, USA, November 19–22, 2013 (NIST Special Publication, Vol. 500-302)*. NIST. <https://trec.nist.gov/pubs/trec22/papers/WEB.OVERVIEW.pdf>
- [34] Kevyn Collins-Thompson, Craig Macdonald, Paul N. Bennett, Fernando Diaz, and Ellen M. Voorhees. 2014. TREC 2014 Web Track Overview. In *Proceedings of the Twenty-Third Text REtrieval Conference, TREC 2014, Gaithersburg, Maryland, USA, November 19–21, 2014 (NIST Special Publication, Vol. 500-308)*. NIST. <https://trec.nist.gov/pubs/trec23/papers/overview-web.pdf>
- [35] Alissa Cooper. 2008. A Survey of Query Log Privacy-Enhancing Techniques from a Policy Perspective. *ACM Transactions on the Web* 2, 4 (2008), 19:1–19:27. <https://doi.org/10.1145/1409220.1409222>
- [36] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Büttcher. 2009. Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In *Proceedings of the 32nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2009, Boston, MA, USA, July 19–23, 2009*. ACM, 758–759. <https://doi.org/10.1145/1571941.1572114>
- [37] Vittoria Cozza, Van Tien Hoang, Marinella Petrocchi, and Rocco De Nicola. 2019. Transparency in Keyword Faceted Search: An Investigation on Google Shopping. In *Digital Libraries: Supporting Open Science. Proceedings of the 15th Italian Research Conference on Digital Libraries, IRCDL 2019, Pisa, Italy, January 31 – February 1, 2019 (Communications in Computer and Information Science, Vol. 988)*. Springer, 29–43. https://doi.org/10.1007/978-3-030-11226-4_3
- [38] Nick Craswell, Daniel Campos, Bhaskar Mitra, Emine Yilmaz, and Bodo Billerbeck. 2020. ORCAS: 18 Million Clicked Query-Document Pairs for Analyzing Search. In *Proceedings of the 29th ACM International Conference on Information and Knowledge Management, CIKM 2020, Virtual Event, Ireland, October 19–23, 2020*. ACM, 2983–2989. <https://doi.org/10.1145/3340531.3412779>
- [39] Nick Craswell and David Hawking. 2002. Overview of the TREC-2002 Web Track. In *Proceedings of the Eleventh Text REtrieval Conference, TREC 2002, Gaithersburg, Maryland, USA, November 19–22, 2002 (NIST Special Publication, Vol. 500-251)*. NIST.
- [40] Nick Craswell and David Hawking. 2004. Overview of the TREC 2004 Web Track. In *Proceedings of the Thirteenth Text REtrieval Conference, TREC 2004, Gaithersburg, Maryland, USA, November 16–19, 2004 (NIST Special Publication, Vol. 500-261)*. NIST.
- [41] Nick Craswell, David Hawking, Ross Wilkinson, and Mingfang Wu. 2003. Overview of the TREC 2003 Web Track. In *Proceedings of the Twelfth Text REtrieval Conference, TREC 2003, Gaithersburg, Maryland, USA, November 18–21, 2003 (NIST Special Publication, Vol. 500-255)*. NIST, 78–92.
- [42] Nick Craswell, Bhaskar Mitra, Emine Yilmaz, and Daniel Campos. 2020. Overview of the TREC 2020 Deep Learning Track. In *Proceedings of the 29th Text REtrieval Conference, TREC 2020, Virtual Event, Gaithersburg, Maryland, USA, November 16–20, 2020 (NIST Special Publication, Vol. 1266)*. NIST.

- [43] Nick Craswell, Bhaskar Mitra, Emine Yilmaz, Daniel Campos, and Ellen M. Voorhees. 2019. Overview of the TREC 2019 Deep Learning Track. In *Proceedings of the 28th Text REtrieval Conference, TREC 2019, Gaithersburg, Maryland, USA, November 13–15, 2019 (NIST Special Publication, Vol. 1250)*. NIST.
- [44] Hang Cui, Ji-Rong Wen, Jian-Yun Nie, and Wei-Ying Ma. 2002. Probabilistic Query Expansion Using Query Logs. In *Proceedings of the Eleventh International World Wide Web Conference, WWW 2002, May 7–11, 2002, Honolulu, Hawaii, USA*. ACM, 325–332. <https://doi.org/10.1145/511446.511489>
- [45] Hang Cui, Ji-Rong Wen, Jian-Yun Nie, and Wei-Ying Ma. 2003. Query Expansion by Mining User Logs. *IEEE Transactions on Knowledge and Data Engineering* 15, 4 (2003), 829–839.
- [46] Erika F. de Lima and Jan O. Pedersen. 1999. Phrase Recognition and Expansion for Short, Precision-Biased Queries Based on a Query Log. In *Proceedings of the 22nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 1999, August 15–19, 1999, Berkeley, CA, USA*. ACM, 145–152. <https://doi.org/10.1145/312624.312669>
- [47] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019, Minneapolis, Minnesota, USA, June 2–7, 2019, Volume 1 (Long and Short Papers)*. ACL, 4171–4186. <https://doi.org/10.18653/v1/n19-1423>
- [48] Georges Dupret, Benjamin Piwowarski, Carlos A. Hurtado, and Marcelo Mendoza. 2006. A Statistical Model of Query Log Generation. In *Proceedings of the 13th International Conference on String Processing and Information Retrieval, SPIRE 2006, Glasgow, UK, October 11–13, 2006 (Lecture Notes in Computer Science, Vol. 4209)*. Springer, 217–228. https://doi.org/10.1007/11880561_18
- [49] Yi Fang, Naveen Somasundaram, Luo Si, Jeongwoo Ko, and Aditya P. Mathur. 2011. Analysis of an Expert Search Query Log. In *Proceeding of the 34th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2011, Beijing, China, July 25–29, 2011*. ACM, 1189–1190. <https://doi.org/10.1145/2009916.2010113>
- [50] Mai Fröbe, Matti Wiegmann, Nikolay Kolyada, Bastian Grahm, Theresa Elstner, Frank Loebe, Matthias Hagen, Benno Stein, and Martin Potthast. 2023. Continuous Integration for Reproducible Shared Tasks with TIRA.io. In *Advances in Information Retrieval. Proceedings of the 45th European Conference on Information Retrieval, ECIR 2023, Dublin, Ireland, April 2–6, 2023, Part III (Lecture Notes in Computer Science, Vol. 13982)*. Springer, 236–241. https://doi.org/10.1007/978-3-031-28241-6_20
- [51] Fernando Galindo and Javier García Marco. 2017. Freedom and the Internet: Empowering Citizens and Addressing the Transparency Gap in Search Engines. *European Journal of Law and Technology* 8, 2 (2017). <https://ejlt.org/index.php/ejlt/article/view/476>
- [52] Daniel Gayo-Avello. 2009. A Survey on Session Detection Methods in Query Logs and a Proposal for Future Evaluation. *Information Sciences* 179, 12 (2009), 1822–1843. <https://doi.org/10.1016/j.ins.2009.01.026>
- [53] Tim Gollub, Benno Stein, Steven Burrows, and Dennis Hoppe. 2012. TIRA: Configuring, Executing, and Disseminating Information Retrieval Experiments. In *Proceedings of the 23rd International Workshop on Database and Expert Systems Applications, DEXA 2012, Vienna, Austria, September 3–7, 2012*. IEEE, 151–155. <https://doi.org/10.1109/DEXA.2012.55>
- [54] Siyu Gu, Jun Yan, Lei Ji, Shuicheng Yan, Junshi Huang, Ning Liu, Ying Chen, and Zheng Chen. 2011. Cross Domain Random Walk for Query Intent Pattern Mining from Search Engine Log. In *Proceedings of the 11th IEEE International Conference on Data Mining, ICDM 2011, Vancouver, BC, Canada, December 11–14, 2011*. IEEE, 221–230. <https://doi.org/10.1109/ICDM.2011.44>
- [55] Dirk Guijt and Claudia Hauff. 2015. Using Query-Log Based Collective Intelligence to Generate Query Suggestions for Tagged Content Search. In *Engineering the Web in the Big Data Era. Proceedings of the 15th International Conference on Web Engineering, ICWE 2015, Rotterdam, The Netherlands, June 23–26, 2015 (Lecture Notes in Computer Science, Vol. 9114)*. Springer, 165–181. https://doi.org/10.1007/978-3-319-19890-3_12
- [56] Matthias Hagen, Jakob Gomoll, Anna Beyer, and Benno Stein. 2013. From Search Session Detection to Search Mission Detection. In *Proceedings of the 10th Conference on Open Research Areas in Information Retrieval, OAIR 2013, Lisbon, Portugal, May 15–17, 2013*. ACM, 85–92. <https://dl.acm.org/doi/10.5555/2491748.2491769>
- [57] Jorge R. Herskovic, Len Y. Tanaka, William R. Hersh, and Elmer V. Bernstam. 2007. A Day in the Life of PubMed: Analysis of a Typical Day's Query Log. *Journal of the American Medical Informatics Association* 14, 2 (2007), 212–220. <https://doi.org/10.1197/jamia.M2191>
- [58] Yunhua Hu, Ya-nan Qian, Hang Li, Dixin Jiang, Jian Pei, and Qinghua Zheng. 2012. Mining Query Subtopics from Search Log Data. In *Proceedings of the 35th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2012, Portland, OR, USA, August 12–16, 2012*. ACM, 305–314. <https://doi.org/10.1145/2348283.2348327>
- [59] Chien-Kang Huang, Lee-Feng Chien, and Yen-Jen Oyang. 2003. Relevant Term Suggestion in Interactive Web Search Based on Contextual Information in Query Session Logs. *Journal of the Association for Information Science and Technology* 54, 7 (2003), 638–649. <https://doi.org/10.1002/asi.10256>
- [60] Po-Sen Huang, Xiaodong He, Jianfeng Gao, Li Deng, Alex Acero, and Larry P. Heck. 2013. Learning Deep Structured Semantic Models for Web Search Using Clickthrough Data. In *Proceedings of the 22nd ACM International Conference on Information and Knowledge Management, CIKM 2013, San Francisco, CA, USA, October 27–November 1, 2013*. ACM, 2333–2338. <https://doi.org/10.1145/2505515.2505665>
- [61] Bernard J. Jansen and Amanda Spink. 2003. An Analysis of Web Documents Retrieved and Viewed. In *Proceedings of the 4th International Conference on Internet Computing, IC 2003, Las Vegas, Nevada, USA, June 23–26, 2003, Volume 1*. CSREA Press, 65–69. http://faculty.ist.psu.edu/jjansen/academic/pubs/pages_viewed.pdf
- [62] Bernard J. Jansen and Amanda Spink. 2005. Methodological Approach in Discovering User Search Patterns through Web Log Analysis. *Bulletin of the American Society for Information Science and Technology* 27, 1 (2005), 15–17. <https://doi.org/10.1002/bult.185>
- [63] Bernard J. Jansen, Amanda Spink, Judy Bateman, and Tefko Saracevic. 1998. Real Life Information Retrieval: A Study of User Queries on the Web. *ACM SIGIR Forum* 32, 1 (1998), 5–17. <https://doi.org/10.1145/281250.281253>
- [64] Bernard J. Jansen, Amanda Spink, and Jan O. Pedersen. 2005. A Temporal Comparison of AltaVista Web Searching. *Journal of the American Society for Information Science and Technology* 56, 6 (2005), 559–570. <https://doi.org/10.1002/asi.20145>
- [65] Bernard J. Jansen, Amanda Spink, and Tefko Saracevic. 2000. Real Life, Real Users, and Real Needs: A Study and Analysis of User Queries on the Web. *Information Processing and Management* 36, 2 (2000), 207–227. [https://doi.org/10.1016/S0306-4573\(99\)00056-4](https://doi.org/10.1016/S0306-4573(99)00056-4)
- [66] Di Jiang, Kenneth Wai-Ting Leung, Wilfred Ng, and Hao Li. 2013. Beyond Click Graph: Topic Modeling for Search Engine Query Log Analysis. In *Proceedings of the 18th International Conference on Database Systems for Advanced Applications, DASFAA 2013, Wuhan, China, April 22–25, 2013, Part I (Lecture Notes in Computer Science, Vol. 7825)*. Springer, 209–223. https://doi.org/10.1007/978-3-642-37487-6_18
- [67] Di Jiang, Yongxin Tong, and Yuanfeng Song. 2016. Cross-Lingual Topic Discovery From Multilingual Search Engine Query Log. *ACM Transactions on Information Systems* 35, 2 (2016), 9:1–9:28. <https://doi.org/10.1145/2956235>
- [68] Jimmy, Guido Zuccon, Bevan Koopman, and Gianluca Demartini. 2019. Health Cards for Consumer Health Search. In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2019, Paris, France, July 21–25, 2019*. ACM, 35–44. <https://doi.org/10.1145/3331184.3331194>
- [69] Jimmy, Guido Zuccon, Bevan Koopman, and Gianluca Demartini. 2019. Health Cards to Assist Decision Making in Consumer Health Search. In *Proceedings of the American Medical Informatics Association Annual Symposium, AMIA 2019, Washington, DC, USA, November 16–20, 2019*. AMIA. <https://knowledge.amia.org/69862-amia-1.4570936/t005-1.4574828/t005-1.4574829/3201885-1.4574890/3201686-1.4574887>
- [70] Thorsten Joachims. 2002. Optimizing Search Engines Using Clickthrough Data. In *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD 2002, July 23–26, 2002, Edmonton, Alberta, Canada*. ACM, 133–142. <https://doi.org/10.1145/775047.775067>
- [71] Thorsten Joachims, Laura A. Granka, Bing Pan, Helene Hembrooke, and Geri Gay. 2005. Accurately Interpreting Clickthrough Data as Implicit Feedback. In *Proceedings of the 28th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2005, Salvador, Brazil, August 15–19, 2005*. ACM, 154–161. <https://doi.org/10.1145/1076034.1076063>
- [72] Christoph Kofler, Linjun Yang, Martha A. Larson, Tao Mei, Alan Hanjalic, and Shipeng Li. 2012. When Video Search Goes Wrong: Predicting Query Failure Using Search Engine Logs and Visual Search Results. In *Proceedings of the 20th ACM Multimedia Conference, MM 2012, Nara, Japan, October 29–November 02, 2012*. ACM, 319–328. <https://doi.org/10.1145/2393347.2393395>
- [73] Torsten Körber. 2015. Common Errors Regarding Search Engine Regulation—And How to Avoid Them. *European Competition Law Review* 36 (2015). Issue 6. https://koerber.jura.uni-koeln.de/sites/koerber/user_upload/Common_errors_regarding_search_engine_regression_Koerber_2015_36_ECLR_Issue_6_Print.pdf
- [74] Chung-Lun Kuo and Hsin-Hsi Chen. 2016. Subtask Mining from Search Query Logs for How-Knowledge Acceleration. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation, LREC 2016, Portorož, Slovenia, May 23–28, 2016*. ELRA. <https://aclanthology.org/L16-1198>
- [75] Nicholas Kushmerick, Daniel S. Weld, and Robert B. Doorenbos. 1997. Wrapper Induction for Information Extraction. In *Proceedings of the Fifteenth International Joint Conference on Artificial Intelligence, IJCAI 1997, Nagoya, Japan, August 23–29, 1997, 2 Volumes*. Elsevier, 729–737. <https://homes.cs.washington.edu/~weld/papers/kushmerick-ijcai97.pdf>
- [76] Emily B. Laidlaw. 2009. Private Power, Public Interest: An Examination of Search Engine Accountability. *International Journal of Law and Information Technology*

- 17, 1 (2009), 113–145. <https://doi.org/10.1093/ijlit/ean018>
- [77] Lars Langer and Erik Frøkjær. 2008. Improving Web Search Transparency by Using a Venn Diagram Interface. In *Proceedings of the 5th Nordic Conference on Human-Computer Interaction, NordiCHI 2008, Lund, Sweden, October 20–22, 2008*, Vol. 358. ACM, 249–256. <https://doi.org/10.1145/1463160.1463187>
- [78] Ruohan Li, Jianxiang Li, Bhaskar Mitra, Fernando Diaz, and Asia J. Biega. 2022. Exposing Query Identification for Search Transparency. In *Proceedings of the ACM Web Conference 2022, WWW 2022, Virtual Event, Lyon, France, April 25 – 29, 2022*. ACM, 3662–3672. <https://doi.org/10.1145/3485447.3512262>
- [79] Zhen Liao, Daxin Jiang, Enhong Chen, Jian Pei, Huanhuan Cao, and Hang Li. 2011. Mining Concept Sequences from Large-Scale Search Logs for Context-Aware Query Suggestion. *ACM Transactions on Intelligent Systems and Technology* 3, 1 (2011), 17:1–17:40. <https://doi.org/10.1145/2036264.2036281>
- [80] Charles X. Ling, Jianfeng Gao, Huajie Zhang, Weining Qian, and HongJiang Zhang. 2001. Mining Generalized Query Patterns from Web Logs. In *Proceedings of the 34th Annual Hawaii International Conference on System Sciences, HICSS-34, January 3–6, 2001, Maui, Hawaii, USA*. IEEE. <https://doi.org/10.1109/HICSS.2001.926534>
- [81] Yan Lu, Michael Chau, and Xiao Fang. 2006. Mining the Query Logs of a Chinese Web Search Engine for Character Usage Analysis. In *Proceedings of the Pacific Asia Conference on Information Systems, PACIS 2006, Kuala Lumpur, Malaysia, July 6–9, 2006*. AIS, 17. <https://aisel.aisnet.org/pacis2006/17>
- [82] Sean MacAvaney, Craig Macdonald, and Iadh Ounis. 2022. Reproducing Personalised Session Search Over the AOL Query Log. In *Advances in Information Retrieval. Proceedings of the 44th European Conference on IR Research, ECIR 2022, Stavanger, Norway, April 10–14, 2022, Part I (Lecture Notes in Computer Science, Vol. 13185)*. Springer, 627–640. https://doi.org/10.1007/978-3-030-99736-6_42
- [83] Thomas Mandl, Maristella Agosti, Giorgio Maria Di Nunzio, Alexander S. Yeh, Inderjeet Mani, Christine Doran, and Julia Maria Schulz. 2009. LogCLEF 2009: The CLEF 2009 Multilingual Logfile Analysis Track Overview. In *Multilingual Information Access Evaluation I. Text Retrieval Experiments. Proceedings of the 10th Workshop of the Cross-Language Evaluation Forum, CLEF 2009, Corfu, Greece, September 30 – October 2, 2009, Revised Selected Papers (Lecture Notes in Computer Science, Vol. 6241)*. Springer, 508–517. https://doi.org/10.1007/978-3-642-15754-7_62
- [84] Thomas Mandl, Giorgio Maria Di Nunzio, and Julia Maria Schulz. 2010. LogCLEF 2010: The CLEF 2010 Multilingual Logfile Analysis Track Overview. In *Notebook Papers of the CLEF 2010 Labs and Workshops, CLEF 2010, Padua, Italy, September 22–23, 2010 (CEUR Workshop Proceedings, Vol. 1176)*. CEUR-WS.org. <https://ceur-ws.org/Vol-1176/CLEF2010wn-LogCLEF-MandlEt2010.pdf>
- [85] Behrooz Mansouri, Mohammad Sadegh Zahedi, Ricardo Campos, and Mojgan Farhoodi. 2018. Online Job Search: Study of Users' Search Behavior using Search Engine Query Logs. In *Proceedings of the 41st International ACM SIGIR Conference on Research & Development in Information Retrieval, SIGIR 2018, Ann Arbor, Michigan, USA, July 8–12, 2018*. ACM, 1185–1188. <https://doi.org/10.1145/3209978.3210125>
- [86] Donn Morrison, Stéphane Marchand-Maillet, and Eric Bruno. 2011. Query Log Simulation for Long-Term Learning in Image Retrieval. In *Proceedings of the 9th International Workshop on Content-Based Multimedia Indexing, CBMI 2011, Madrid, Spain, June 13–15, 2011*. IEEE, 55–60. <https://doi.org/10.1109/CBMI.2011.5972520>
- [87] Donn Morrison, Theodora Tsikrika, Vera Hollink, Arjen P. de Vries, Eric Bruno, and Stéphane Marchand-Maillet. 2013. Topic Modelling of Clickthrough Data in Image Search. *Multimedia Tools and Applications* 66, 3 (2013), 493–515. <https://doi.org/10.1007/s11042-012-1038-8>
- [88] Yashar Moshfeghi. 2021. NeuraSearch: Neuroscience and Information Retrieval. In *Proceedings of the Second International Conference on Design of Experimental Search & Information Retrieval Systems, Padova, Italy, September 15–18, 2021 (CEUR Workshop Proceedings, Vol. 2950)*. CEUR-WS.org, 193–194. <https://ceur-ws.org/Vol-2950/paper-27.pdf>
- [89] Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. 2016. MS MARCO: A Human Generated MAchine REading COmprehension Dataset. In *Proceedings of the Workshop on Cognitive Computation: Integrating Neural and Symbolic Approaches 2016 co-located with the 30th Annual Conference on Neural Information Processing Systems, CoCo@NIPS 2016, Barcelona, Spain, December 9, 2016 (CEUR Workshop Proceedings, Vol. 1773)*. CEUR-WS.org. https://ceur-ws.org/Vol-1773/CoCoNIPS_2016_paper9.pdf
- [90] Bruno Oliveira and Carla Teixeira Lopes. 2023. The Evolution of Web Search User Interfaces – An Archaeological Analysis of Google Search Engine Result Pages. In *Proceedings of the 2023 Conference on Human Information Interaction and Retrieval, CHIIR 2023, Austin, Texas, USA, March 19–23, 2023*. ACM, 55–68. <https://doi.org/10.1145/3576840.3578320>
- [91] Bruno Oliveira and Carla Teixeira Lopes. 2023. From 10 Blue Links Pages to Feature-Full Search Engine Results Pages – Analysis of the Temporal Evolution of SERP Features. In *Proceedings of the 2023 Conference on Human Information Interaction and Retrieval, CHIIR 2023, Austin, Texas, USA, March 19–23, 2023*. ACM, 338–345. <https://doi.org/10.1145/3576840.3578307>
- [92] Alexandra Olteanu, Jean Garcia-Gathright, Maarten de Rijke, and Michael D. Ekstrand. 2019. Workshop on Fairness, Accountability, Confidentiality, Transparency, and Safety in Information Retrieval (FACTS-IR). In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2019, Paris, France, July 21–25, 2019*. ACM, 1423–1425. <https://doi.org/10.1145/3331184.3331644>
- [93] Harrie Oosterhuis and Maarten de Rijke. 2018. Differentiable Unbiased Online Learning to Rank. In *Proceedings of the 27th ACM International Conference on Information and Knowledge Management, CIKM 2018, Torino, Italy, October 22–26, 2018*. ACM, 1293–1302. <https://doi.org/10.1145/3269206.3271686>
- [94] João R. M. Palotti, Allan Hanbury, Henning Müller, and Charles E. Kahn Jr. 2016. How Users Search and What They Search for in the Medical Domain – Understanding Laypeople and Experts Through Query Logs. *Information Retrieval Journal* 19, 1–2 (2016), 189–224. <https://doi.org/10.1007/s10791-015-9269-8>
- [95] Greg Pass, Abdur Chowdhury, and Cayley Torgeson. 2006. A Picture of Search. In *Proceedings of the 1st International Conference on Scalable Information Systems, Infoscale 2006, Hong Kong, May 30–June 1, 2006*. Vol. 152. ACM, 1. <https://doi.org/10.1145/146847.146848>
- [96] Martin Potthast, Tim Gollub, Matti Wiegmann, and Benno Stein. 2019. TIRA Integrated Research Architecture. In *Information Retrieval Evaluation in a Changing World – Lessons Learned from 20 Years of CLEF*. The Information Retrieval Series, Vol. 41. Springer, 123–160. https://doi.org/10.1007/978-3-030-22948-1_5
- [97] Tao Qin and Tie-Yan Liu. 2013. Introducing LETOR 4.0 Datasets. *arXiv.org abs/1306.2597* (2013). <https://doi.org/10.48550/arXiv.1306.2597>
- [98] Filip Radlinski, Paul N. Bennett, and Emine Yilmaz. 2011. Detecting Duplicate Web Documents Using Clickthrough Data. In *Proceedings of the Forth International Conference on Web Search and Web Data Mining, WSDM 2011, Hong Kong, China, February 9–12, 2011*. ACM, 147–156. <https://doi.org/10.1145/1935826.1935859>
- [99] Filip Radlinski and Thorsten Joachims. 2006. Minimally Invasive Randomization for Collecting Unbiased Preferences from Clickthrough Logs. *arXiv.org abs/cs/0605037* (2006). <https://doi.org/10.48550/arXiv.cs/0605037>
- [100] Shaurya Rohatgi, C. Lee Giles, and Jian Wu. 2021. What Were People Searching For? A Query Log Analysis of An Academic Search Engine. In *Proceedings of the ACM/IEEE Joint Conference on Digital Libraries, JCDL 2021, Champaign, Illinois, USA, September 27–30, 2021*. IEEE, 342–343. <https://doi.org/10.1109/JCDL52503.2021.000062>
- [101] Harrisen Scells, Connor Forbes, Justin Clark, Bevan Koopman, and Guido Zuccon. 2022. The Impact of Query Refinement on Systematic Review Literature Search: A Query Log Analysis. In *Proceedings of the 2022 ACM SIGIR International Conference on the Theory of Information Retrieval, ICTIR 2022, Madrid, Spain, July 11 – 12, 2022*. ACM, 34–42. <https://doi.org/10.1145/3539813.3545143>
- [102] Craig Silverstein, Monika Rauch Henzinger, Hannes Marais, and Michael Moricz. 1999. Analysis of a Very Large Web Search Engine Query Log. *ACM SIGIR Forum* 33, 1 (1999), 6–12. <https://doi.org/10.1145/331403.331405>
- [103] Amanda Spink, Bernard J. Jansen, and H. Cenk Ozmutlu. 2000. Use of Query Reformulation and Relevance Feedback by Excite Users. *Internet Research* 10, 4 (2000), 317–328. <https://doi.org/10.1108/10662240010342621>
- [104] Amanda Spink, Seda Ozmutlu, Huseyin C. Ozmutlu, and Bernard J. Jansen. 2002. U.S. Versus European Web Searching Trends. 36, 2 (2002), 32–38. <https://doi.org/10.1145/792550.792555>
- [105] Amanda Spink, Dietmar Wolfram, Bernard J. Jansen, and Tefko Saracevic. 2001. Searching the Web: The Public and Their Queries. *Journal of the American Society for Information Science and Technology* 52, 3 (2001), 226–234. [https://doi.org/10.1002/1097-4571\(2000\)9999:999%3C::AID-ASI1591%3E3.0.CO;2-R](https://doi.org/10.1002/1097-4571(2000)9999:999%3C::AID-ASI1591%3E3.0.CO;2-R)
- [106] Xu Sun, Jianfeng Gao, Daniel Micol, and Chris Quirk. 2010. Learning Phrase-Based Spelling Error Models from Clickthrough Data. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics, ACL 2010, July 11–16, 2010, Uppsala, Sweden*. Association for Computer Linguistics, 266–274. <https://aclanthology.org/P10-1028/>
- [107] Ellen M. Voorhees. 2004. Overview of the TREC 2004 Robust Retrieval Track. In *Proceedings of the Thirteenth Text RETrieval Conference, TREC 2004, Gaithersburg, Maryland, USA, November 16–19, 2004 (NIST Special Publication, Vol. 500-261)*. NIST. <https://trec.nist.gov/pubs/trec13/papers/ROBUST.OVERVIEW.pdf>
- [108] Ellen M. Voorhees. 2005. Overview of the TREC 2005 Robust Retrieval Track. In *Proceedings of the Fourteenth Text RETrieval Conference, TREC 2005, Gaithersburg, Maryland, USA, November 15–18, 2005 (NIST Special Publication, Vol. 500-266)*. NIST. <https://trec.nist.gov/pubs/trec14/papers/ROBUST.OVERVIEW.pdf>
- [109] Ji-Rong Wen, Jian-Yun Nie, and HongJiang Zhang. 2002. Query Clustering Using User Logs. *ACM Transactions on Information Systems* 20, 1 (2002), 59–81. <https://doi.org/10.1145/503104.503108>
- [110] Dietmar Wolfram, Amanda Spink, Bernard J. Jansen, and Tefko Saracevic. 2001. Vox Populi: The Public Searching of the Web. *Journal of the American Society for Information Science and Technology* 52, 12 (2001), 1073–1074. <https://doi.org/10.1002/asi.1157>

- [111] Ting Yao, Min Zhang, Yiqun Liu, Shaoping Ma, Yongfeng Zhang, and Liyun Ru. 2010. Investigating Characteristics of Non-click Behavior Using Query Logs. In *Information Retrieval Technology. Proceedings of the 6th Asia Information Retrieval Societies Conference, AIRS 2010, Taipei, Taiwan, December 1–3, 2010 (Lecture Notes in Computer Science, Vol. 6458)*. Springer, 85–96. https://doi.org/10.1007/978-3-642-17187-1_8
- [112] Deng Yi, Yin Zhang, Haihan Yu, Yanfei Yin, Jing Pan, and Baogang Wei. 2012. Improving Multi-Faceted Book Search by Incorporating Sparse Latent Semantic Analysis of Click-through Logs. In *Proceedings of the 12th ACM/IEEE Joint Conference on Digital Libraries, JCDL 2012, Washington, DC, USA, June 10–14, 2012*. ACM, 249–258. <https://doi.org/10.1145/2232817.2232864>
- [113] Dell Zhang and Yisheng Dong. 2002. A Novel Web Usage Mining Approach for Search Engines. *Computer Networks* 39, 3 (2002), 303–310. [https://doi.org/10.1016/S1389-1286\(02\)00211-6](https://doi.org/10.1016/S1389-1286(02)00211-6)
- [114] Zhitao Zhang, Muyun Yang, Sheng Li, Haoliang Qi, and Chao Song. 2009. So-gou Query Log Analysis: A Case Study for Collaborative Recommendation or Personalized IR. In *Proceedings of the 2009 International Conference on Asian Language Processing, IALP 2009, Singapore, December 7–9, 2009*. IEEE, 304–307. <https://doi.org/10.1109/IALP.2009.72>
- [115] Shiqi Zhao, Haifeng Wang, and Ting Liu. 2012. User Behaviors Lend a Helping Hand: Learning Paraphrase Query Patterns from Search Log Sessions. In *Proceedings of the 24th International Conference on Computational Linguistics, COLING 2012, 8–15 December 2012, Mumbai, India, Technical Papers*. ACL, 3137–3152. <https://aclanthology.org/C12-1192>
- [116] Xiang Zhou, Pengyi Zhang, and Jun Wang. 2017. Identification and Analysis of Multi-tasking Product Information Search Sessions with Query Logs. *Journal of Data and Information Science* 1, 3 (2017), 79–94. <https://doi.org/10.20309/jdis.201621>
- [117] Shengyao Zhuang and Guido Zuccon. 2020. Counterfactual Online Learning to Rank. In *Advances in Information Retrieval. Proceedings of the 42nd European Conference on IR Research, ECIR 2020, Lisbon, Portugal, April 14–17, 2020, Part I (Lecture Notes in Computer Science, Vol. 12035)*. Springer, 415–430. https://doi.org/10.1007/978-3-030-45439-5_28